

Contents

1	Appendix B: Zero-Knowledge Proof Systems - Technical Implementation	1
1.1	B.1 Cryptographic Foundations	1
1.1.1	B.1.1 Proof System Definitions	1
1.1.2	B.1.2 SNARKs (Succinct Non-Interactive Arguments of Knowledge)	1
1.2	B.2 Backend Implementations	2
1.2.1	B.2.1 Groth16 - Optimal Performance	2
1.2.2	B.2.2 PLONK - Universal Setup	3
1.2.3	B.2.3 Halo2 - Trustless Construction	4
1.3	B.3 Circuit Implementation	4
1.3.1	B.3.1 Variant Presence Circuit	4
1.3.2	B.3.2 Polygenic Risk Score Circuit	5
1.3.3	B.3.3 Ancestry Estimation Circuit	6
1.4	B.4 Performance Optimization	7
1.4.1	B.4.1 Batch Proving	7
1.4.2	B.4.2 Proof Caching	7
1.5	B.5 Security Analysis	8
1.5.1	B.5.1 Soundness Analysis	8
1.5.2	B.5.2 Zero-Knowledge Analysis	8
1.6	B.6 Production Monitoring	8
1.6.1	B.6.1 Key Metrics	8
1.6.2	B.6.2 Alerting Rules	9
1.7	B.7 References	9

1 Appendix B: Zero-Knowledge Proof Systems - Technical Implementation

1.1 B.1 Cryptographic Foundations

1.1.1 B.1.1 Proof System Definitions

Definition B.1 (Zero-Knowledge Proof System)

A zero-knowledge proof system for an NP language L consists of three probabilistic polynomial-time algorithms (Setup, Prove, Verify):

- **Setup**($1^\lambda, C$) \rightarrow (**pp**, **vk**): Generates public parameters **pp** and verification key **vk** for circuit C
- **Prove**(**pp**, x , w) \rightarrow : Generates proof that (x, w) satisfies circuit C
- **Verify**(**vk**, x , π) \rightarrow $\{0,1\}$: Verifies proof π for public input x

Properties: 1. **Completeness:** For all valid (x, w) , Verify accepts with probability 1 2. **Soundness:** No adversary can create accepting proof for invalid x 3. **Zero-Knowledge:** Proof reveals nothing beyond truth of statement

1.1.2 B.1.2 SNARKs (Succinct Non-Interactive Arguments of Knowledge)

Succinctness: $|\pi| = \text{poly}(\lambda, \log|C|)$, independent of witness size **Non-Interactive:** Single message from prover to verifier **Argument:** Computational soundness (not information-theoretic) **of**

Knowledge: Extractor can recover witness from accepting proof

1.2 B.2 Backend Implementations

1.2.1 B.2.1 Groth16 - Optimal Performance

Cryptographic Construction:

Curve: BLS12-381 - **Base field:** \mathbb{F}_q where $q = 2^{381}$ - **Scalar field:** \mathbb{F}_r where $r = 2^{255}$ -
Embedding degree: $k = 12$ - **Security:** 128-bit

Proof Structure:

$$P = (A \cdot G, B \cdot G, C \cdot G)$$

Verification Equation:

$$e(A, B) = e(L, C) \cdot e(L, C) \cdot e(C, C)$$

where: - $e: G \times G \rightarrow G_T$ is optimal Ate pairing - $\tau, \tau^2, \dots, \tau^N$: Setup parameters from trusted ceremony - L : Linear combination of public inputs

Proof Size Breakdown:

A : 48 bytes (compressed G point)
 B : 96 bytes (compressed G point)
 C : 48 bytes (compressed G point)
Total: 192 bytes

Trusted Setup Protocol:

Phase 1 (Powers of Tau):

$\tau, \tau^2, \tau^3, \dots, \tau^N$ in G, G

- Universal: Reusable across all circuits up to size N
- Participants: 1000+ in Perpetual Powers of Tau ceremony
- Security: Need only 1 honest participant

Phase 2 (Circuit-Specific):

$\tau, \tau^2, \tau^3, \dots, \tau^N$ for specific circuit

- Circuit-specific: Must rerun for circuit modifications
- Multi-party computation: N participants contribute randomness
- Security: 1-of- N honesty assumption

Key Compromise Response:

Indicators of compromise: 1. Invalid proofs that verify 2. Leaked ceremony transcripts 3. Participant compromise acknowledgment

Immediate response ($T < 1$ hour):

```
# 1. Disable affected circuits
genomevault zk disable-circuit variant_presence --reason "key_compromise"
```

```
# 2. Alert downstream systems
```

```

genomevault alerts broadcast --level critical --msg "ZK key compromise detected"

# 3. Queue re-generation
genomevault zk queue-regenall --start-after-ceremony

Recovery procedure (T = 24-48 hours):

# 1. New ceremony with vetted participants
snarkjs groth16 setup circuit.r1cs pot28_final.ptau circuit_0000.zkey

# 2-N. Participants contribute
for i in {1..10}; do
    snarkjs zkey contribute circuit_${(i-1)}.zkey circuit_${i}.zkey \
        --name "Emergency Contributor $i" \
        --entropy $(openssl rand -hex 32)
done

# Final beacon
snarkjs zkey beacon circuit_10.zkey circuit_final.zkey $(openssl rand -hex 32) 10

# 3. Export new verification key
snarkjs zkey export verificationkey circuit_final.zkey vk_new.json

# 4. Update production
genomevault zk update-vk variant_presence vk_new.json --verify-ceremony

```

1.2.2 B.2.2 PLONK - Universal Setup

Cryptographic Construction:

Polynomial Commitment: KZG (Kate-Zaverucha-Goldberg)

Commit to polynomial $f(X)$ of degree d :

$$C = [f()] = f() \cdot G$$

where s is secret from trusted setup.

Opening Proof: For claimed evaluation $f(z) = y$:

$$= [(f(X) - y)/(X - z)]$$

Verification:

$$e(C - [y], G) = e([s], [1] - [z])$$

Universal SRS (Structured Reference String):

$$\{[s_i], [1] : i = 0 \dots N\}$$

Advantages: - **Reusable:** All circuits up to size N - **Updatable:** Additional ceremonies extend without invalidating - **Available:** Aztec's Ignition ceremony provides SRS up to 2^2 constraints

Proof Composition:

```

= (
  a_comm, b_comm, c_comm,      # Wire commitments
  z_comm,                     # Permutation accumulator
  t_lo, t_mid, t_hi,          # Quotient polynomial (chunked)
  a_eval, b_eval, c_eval,      # Wire evaluations
  s_1_eval, s_2_eval,         # Permutation evaluations
  z_omega_eval,               # Next accumulator eval
  opening_proof,              # Batch opening
  challenge_response           # Fiat-Shamir transcript
)

```

Proof Size: Approximately 1KB (7 G points + 7 scalars)

1.2.3 B.2.3 Halo2 - Trustless Construction

Innovation: Eliminates trusted setup via IPA (Inner Product Argument)

Polynomial Commitment (IPA-based):

Commit to polynomial $f(X)$ with coefficients $f = (f_0, f_1, \dots, f_d)$:

$$C = f, G = \sum f_i \cdot G_i$$

where G is deterministic from hash function (no trusted setup).

Opening Proof: Recursive halving protocol

To prove $f(z) = y$:

1. Split $f = (f_L, f_R)$, $G = (G_L, G_R)$
2. Compute cross terms L, R
3. Verifier sends challenge x
4. Recurse on $f' = f_L + x \cdot f_R$, $G' = G_L + x \cdot G_R$
5. After $\log(d)$ rounds, verify base case

Proof Size: $O(\log d)$ group elements 5KB for $d = 2^2$

Pasta Curves (Pallas/Vesta):

Purpose: Enable efficient recursion without pairing-friendly curves

- **Pallas:** Base field \mathbb{F}_p , scalar field \mathbb{F}_q
- **Vesta:** Base field \mathbb{F}_q , scalar field \mathbb{F}_p

Cycle property: p 's order = q , q 's order = p

This allows: - Prove Vesta circuit in Pallas - Prove Pallas circuit in Vesta - Compose recursively without field switching overhead

1.3 B.3 Circuit Implementation

1.3.1 B.3.1 Variant Presence Circuit

Circuit Logic:

```

template VariantPresence(numVariants) {
    // Private inputs
    signal input variants[numVariants];      // Patient's variants
    signal input queryVariant;                // Variant to check

    // Public output
    signal output hasVariant;

    // Intermediate signals
    signal isMatch[numVariants];
    signal accumulator[numVariants];

    // Check each variant
    accumulator[0] <== 0;
    for (var i = 0; i < numVariants; i++) {
        isMatch[i] <== IsEqual()(variants[i], queryVariant);
        if (i > 0) {
            accumulator[i] <== accumulator[i-1] + isMatch[i];
        }
    }

    // Output 1 if any match found, 0 otherwise
    hasVariant <== GreaterThan(32)(accumulator[numVariants-1], 0);
}

```

Constraint Count Analysis: - IsEqual: 2 constraints per comparison - GreaterThan: 252 constraints (32-bit comparison) - Total: $2N + 252$ 15,234 for $N=7,500$ variants

1.3.2 B.3.2 Polygenic Risk Score Circuit

Circuit Logic:

```

template PolygeneticRisk(numSNPs) {
    // Private inputs
    signal input snps[numSNPs];                // Binary SNP values {0,1,2}
    signal input weights[numSNPs];             // Risk weights (fixed-point)
    signal input salt;                          // Privacy salt

    // Public inputs
    signal input threshold;                     // Risk threshold
    signal input commitment;                   // Hash commitment to SNPs

    // Public output
    signal output isHighRisk;

    // Constraint: Verify SNP commitment
    component hasher = Poseidon(numSNPs + 1);
    for (var i = 0; i < numSNPs; i++) {
        hasher.inputs[i] <== snps[i];
    }
}

```

```

}
hasher.inputs[numSNPs] <== salt;
hasher.out == commitment;

// Compute weighted risk score
signal partialSums[numSNPs];
partialSums[0] <== snps[0] * weights[0];
for (var i = 1; i < numSNPs; i++) {
    partialSums[i] <== partialSums[i-1] + snps[i] * weights[i];
}
signal riskScore <== partialSums[numSNPs - 1];

// Check if risk exceeds threshold
isHighRisk <== GreaterThan(32)([riskScore, threshold]);
}

```

Constraint Count: $O(\text{numSNPs})$ 1M for comprehensive PRS

1.3.3 B.3.3 Ancestry Estimation Circuit

Circuit Logic: Principal component analysis on genetic markers

```

template AncestryProof(numMarkers, numPCs) {
    // Private inputs
    signal input markers[numMarkers];          // Ancestry-informative markers
    signal input pcWeights[numMarkers][numPCs]; // PCA weights

    // Public outputs
    signal output ancestry;                    // Ancestry category {0,1,2,...}

    // Compute principal components
    signal pcs[numPCs];
    for (var pc = 0; pc < numPCs; pc++) {
        signal partialSum[numMarkers];
        partialSum[0] <== markers[0] * pcWeights[0][pc];
        for (var i = 1; i < numMarkers; i++) {
            partialSum[i] <== partialSum[i-1] + markers[i] * pcWeights[i][pc];
        }
        pcs[pc] <== partialSum[numMarkers - 1];
    }

    // Classification decision tree (simplified)
    signal isEuropean <== GreaterThan(32)([pcs[0], threshold_eur]);
    signal isAfrican <== GreaterThan(32)([pcs[1], threshold_afr]);
    // ... additional logic for multi-way classification

    ancestry <== 0 + isEuropean + 2*isAfrican + ...;
}

```

1.4 B.4 Performance Optimization

1.4.1 B.4.1 Batch Proving

Process multiple proofs in parallel:

```
from genomevault.zk_proofs import BatchProver
from concurrent.futures import ThreadPoolExecutor

prover = BatchProver(backend="halo2", workers=10)

# Queue proofs
proof_ids = []
for i, (public, private) in enumerate(inputs):
    proof_id = prover.queue_proof(
        circuit="variant_presence",
        public_inputs=public,
        private_inputs=private
    )
    proof_ids.append(proof_id)

# Wait for completion
proofs = prover.wait_all(proof_ids, timeout=300)

print(f"Generated {len(proofs)} proofs in batch")
```

Throughput: - Serial: 1.67 proofs/core/sec - Parallel (10 cores): 16.7 proofs/sec - With caching (40% hit rate): 27.8 proofs/sec effective

1.4.2 B.4.2 Proof Caching

Cache proofs for common queries:

```
from genomevault.zk_proofs import ProofCache

cache = ProofCache(
    backend="redis",
    ttl=86400, # 24 hours
    max_size="10GB"
)

# Check cache before proving
cache_key = hash((circuit, public_inputs))
cached_proof = cache.get(cache_key)

if cached_proof:
    return cached_proof
else:
    proof = prover.prove(circuit, public_inputs, private_inputs)
    cache.set(cache_key, proof)
```

```
return proof
```

Cache Hit Rates (measured in production): - Variant presence queries: 42% hit rate - PRS calculations: 18% hit rate (more varied) - Ancestry checks: 65% hit rate (limited ancestry groups)

1.5 B.5 Security Analysis

1.5.1 B.5.1 Soundness Analysis

Groth16 Soundness Error:

```
_soundness = 1 / |_r| 2(-255) 10(-77)
```

PLONK Soundness Error:

```
_soundness = (d + 6) / || 10(-75) for d = 10
```

Halo2 Soundness Error:

```
_soundness = O(d / |_p|) 10(-74) for d = 10
```

All provide negligible soundness error ($\ll 2^{(-128)}$ security).

1.5.2 B.5.2 Zero-Knowledge Analysis

Groth16 ZK Simulator:

Given verification key vk and statement x, simulator can produce indistinguishable proofs without witness:

```
def simulate_proof(vk, x):  
    # Sample random group elements  
    _A = random_G1()  
    _B = random_G2()  
  
    # Compute _C to satisfy verification equation  
    _C = compute_valid_C(vk, x, _A, _B)  
  
    return (_A, _B, _C)
```

Zero-Knowledge Property: No polynomial-time distinguisher can tell real proofs from simulated proofs with advantage $> \epsilon$.

1.6 B.6 Production Monitoring

1.6.1 B.6.1 Key Metrics

Proving Metrics:

```
proof_generation_time_ms:  
    p50: 603  
    p95: 711  
    p99: 711
```

```
peak_memory_usage_mb:
```



```
p50: 4200
p95: 4350
p99: 4400
```

```
proof_queue_depth:
  threshold: 100
  alert: depth > 100 for 5 minutes
```

Verification Metrics:

```
proof_verification_time_ms:
  p50: 20.4
  p95: 23.1
  p99: 23.2
```

```
verification_failure_rate:
  threshold: 0.001 # 0.1%
  alert: rate > 0.001 over 1 hour
```

1.6.2 B.6.2 Alerting Rules

```
alerts:
- name: ProvingTimeAnomaly
  condition: proof_generation_time_ms.p95 > 1000
  severity: warning
  action: Scale up prover pool

- name: VerificationFailure
  condition: verification_failure_rate > 0.001
  severity: critical
  action: Investigate key compromise, disable circuit

- name: QueueBacklog
  condition: proof_queue_depth > 100
  severity: warning
  action: Add prover workers

- name: MemoryExhaustion
  condition: peak_memory_usage_mb > 60000
  severity: critical
  action: Restart prover, upgrade instance
```

1.7 B.7 References

1. Groth, J. (2016). “On the size of pairing-based non-interactive arguments.” In *Annual International Conference on the Theory and Applications of Cryptographic Techniques* (pp. 305-326).
2. Gabizon, A., Williamson, Z. J., & Ciobotaru, O. (2019). “PLONK: Permutations over Lagrange-bases for Oecumenical Noninteractive arguments of Knowledge.” *IACR ePrint*

Archive.

3. Bowe, S., Grigg, J., & Hopwood, D. (2020). “Recursive Proof Composition without a Trusted Setup.” *IACR ePrint Archive*.
4. Kate, A., Zaverucha, G. M., & Goldberg, I. (2010). “Constant-size commitments to polynomials and their applications.” In *International Conference on the Theory and Application of Cryptology and Information Security* (pp. 177-194).

Implementation: Complete circuits available in `zk_circuits/` directory with compilation instructions and test vectors.