

Constraint-Oriented Emergent Computation: A Formal Framework Bridging Biology, Information Theory, and Distributed Agency

Abstract

This thesis proposes a formalism for Constraint-Oriented Emergent Computation (COEC), a substrate-independent framework describing computation as the trajectory of physical or biological systems through constrained state spaces. In COEC, computation emerges not through discrete logic or symbolic manipulation, but through systems undergoing entropy-driven transitions within boundary conditions. By incorporating information-theoretic principles, predictive processing, topological data analysis, category theory, and distributed agency, this framework provides a unified mathematical language for understanding computational processes across diverse biological contexts—from protein folding to neural dynamics.

The thesis establishes formal connections between computational substrates and thermodynamic, informational, and variational principles, offering novel perspectives on both biological and artificial computing systems. Key theoretical contributions include the integration of distributed agency concepts, integration with established theories such as the Free Energy Principle and Integrated Information Theory, and applications spanning domains from climate science to financial systems and artificial intelligence.

COEC demonstrates particular promise in neuromorphic computing, quantum systems, and synthetic biology, where its constraint-based approach aligns naturally with underlying physical and biological principles. The thesis also addresses educational, commercial, and ethical dimensions of the framework, positioning COEC as a bridge between disciplines and computational paradigms.

1. Introduction

Biological systems demonstrate sophisticated computational behavior without reliance on symbolic logic or centralized control. Proteins fold into functional configurations, cellular networks process environmental signals, and embryos develop into complex organisms—all without explicit algorithms or centralized decision-making. This suggests an alternative paradigm: computation as the consequence of systems minimizing surprise through

constraint-guided physical evolution. We define and formalize this paradigm as Constraint-Oriented Emergent Computation (COEC).

Traditional computational frameworks rely on discrete state transitions governed by explicit rules (e.g., Turing machines, cellular automata, Boolean networks). While effective for engineered systems, these frameworks often struggle to capture the fluid, parallel, and distributed nature of biological computation. By contrast, COEC conceptualizes computation as the natural trajectory of a physical system through its state space, shaped by physical constraints and information-theoretic principles.

From an information-theoretic perspective, biological systems can be understood as simultaneously reducing uncertainty about their environment while maintaining their internal organization. This process—navigating constraints while optimizing information flow—constitutes computation in the COEC framework. These systems effectively implement forms of predictive processing, where internal models guide interactions with the environment through perception-action loops.

COEC views computation as dynamics in a constrained energy-information landscape. We embed symmetry and algebraic structures into this formal ontology, particularly through Noether-inspired invariance constraints that articulate how symmetries of constraint networks yield conserved quantities, similar to how Noether's theorem links physical symmetries to conservation laws. For example, a spatial translational symmetry in a developing morphogen gradient implies a conserved "momentum-like" quantity in its pattern dynamics. Formally, COEC energy functions are constructed so that symmetric transformations of the state leave the effective Hamiltonian invariant, leading to global invariants despite heterogeneous local constraints through sheaf-like gluing of invariances.

We also introduce Galois-guided dimensional reduction, where algebraic groups index quotient spaces when systems collapse degrees of freedom under constraints. Inspired by Langlands dualities, we conjecture dual pairs of COEC systems whose constraint spectra (eigenvalues of constraint operators) match via a kind of "constraint Langlands correspondence." These ideas allow systematic reduction of complex state spaces by quotienting out symmetry or combinatorial structure, guiding model simplification without losing core computational behavior.

1.1 Building Intuition: From Traditional Computing to Emergent Constraint Satisfaction

Before diving into formal definitions, it's helpful to contrast traditional computing with COEC through a simple example. Consider protein folding:

Traditional computational view: A protein folds by "executing" the "instructions" in its amino acid sequence, as if following an algorithm.

COEC view: The protein navigates an energy landscape constrained by chemical bonds, hydrophobicity, and electromagnetic forces. It doesn't "execute" anything—it simply follows the

path of least resistance within these constraints, ultimately reaching a stable configuration that represents the "solution" to the implicit computational problem.

This perspective shift applies across scales:

- **Cellular level:** Gene regulatory networks don't "process information" like a computer; they follow energy gradients constrained by transcription factors, chromatin accessibility, and metabolic conditions.
- **Organismal level:** An ant colony doesn't "compute" the shortest path to food; the path emerges from pheromone gradients that constrain individual ant movements.

In COEC, computation is what happens when a system moves toward equilibrium while satisfying multiple simultaneous constraints.

1.2 Distributed Agency and Emergent Function

Contemporary models of biological computation reject central controllers and symbolic algorithms, instead casting "computation" as the system's trajectory through a constrained state space. In this constraint-oriented emergent view, agency and function arise from many interacting parts and their mutual constraints. Each local interaction or "constraint" (be it energetic, topological, or informational) biases the system's dynamics, much like multiple overlapping "pulls" on an object. The global behavior (the apparent "will" or function of the system) emerges from these interactions rather than from a single executive.

This perspective parallels ideas from complex adaptive systems and autopoietic theory: large ensembles of simple agents or components self-organize and adapt (often unpredictably) without a central planner. In effect, distributed agency arises - the system behaves as if guided by a purpose or "spirit," even though no single entity controls it.

To illustrate distributed agency across scales:

- **Molecular scale:** In protein folding, each amino acid residue contributes local preferences for certain configurations based on charge, hydrophobicity, and steric constraints. No residue "directs" the folding process, yet a functional structure emerges.
- **Cellular scale:** During chemotaxis, a bacterium's receptors, signaling proteins, and flagellar motors collectively generate directed movement without a central "navigator."
- **Multicellular scale:** In embryogenesis, cells respond to local morphogen gradients and mechanical cues without a master blueprint, yet complex, reproducible structures emerge.

Biological substrates often decompose into many interacting subunits, each with limited scope or information. This is analogous to swarm or multi-agent systems in engineering: many simple agents, each with local rules, can collectively solve complex problems without central oversight. In a swarm-agentic AI, for instance, individual agents coordinate via local signals and feedback, and the group-level behavior emerges collectively. COEC envisions biological systems similarly: no single controller directs the computation. Instead, numerous constraints - akin to goal-representing sub-agents - interact. Each constraint biases the trajectory of the system in one direction (toward satisfying that constraint), but since all constraints act simultaneously, the actual behavior is the product of their interaction.

1.3 Biomimetic Inspiration and Differentiation Models

The COEC framework draws significant inspiration from biological differentiation processes. In multicellular organisms, cells progress from pluripotent states to specialized functional roles through a series of constraint-guided transitions. We formalize this as Biological Specialization as Constraint Flow:

Definition: Biological differentiation can be modeled as a traversal through regulated constraint fields, where each developmental stage represents a specific configuration of active constraints.

This approach allows us to model developmental pathways (e.g., stem cell to beta cell) as trajectories through state spaces with dynamically changing constraint sets. Mathematically, this can be expressed as:

$$S_{\text{differentiated}} = \Phi(S_{\text{stem}} \parallel C_{\{t_0\}}, C_{\{t_1\}}, \dots, C_{\{t_n\}})$$

where $C_{\{t_i\}}$ represents the constraint set active at developmental time t_i .

For example, in hematopoiesis (blood cell development):

1. A hematopoietic stem cell exists in a state with relatively few active constraints
2. As differentiation proceeds, lineage-specific transcription factors create new constraints on gene expression
3. Terminal differentiation involves highly restrictive constraints that stabilize the final cell type

This biomimetic model provides a template for embedding modular, lineage-like constraint sequences into emergent computational systems, offering a formal bridge between biological development and synthetic computational frameworks.

1.4 Connections to Free Energy Principle and Integrated Information Theory

COEC's ontology formally embraces thermodynamic and information-theoretic principles. The energy-information landscape explicitly combines physical energy costs with

information-constraint terms (entropy). In this view, COEC dynamics can be interpreted as variational free energy minimization: the system naturally moves to reduce a generalized free-energy, analogous to Friston's Free Energy Principle (FEP) in neuroscience.

The FEP states that physical systems act to minimize prediction error or "surprisal" (negative log probability) over time. Similarly, COEC substrates evolve by gradient descent on an energy function plus penalty terms for unsatisfied constraints, i.e., they minimize:

$$F(s) = E_{\text{physical}}(s) + \sum_i w_i \cdot c_i(s)$$

which is conceptually equivalent to minimizing variational free energy. This unifies COEC with active inference: COEC = active inference + constraints, where constraints play the role of learned models or prediction errors. By this mapping, COEC formally inherits FEP's interpretive power: purpose (e.g., homeostasis) emerges from free-energy minimization under constraints.

Moreover, COEC clarifies links to Integrated Information Theory (IIT). IIT's measure Φ quantifies how much a whole system generates information beyond its independent parts. In COEC, higher constraint precision (greater weight on holistic constraints) similarly yields more integrated, irreducible dynamics. Roughly, a network of highly precise constraints is harder to decompose without loss, akin to high Φ . Thus, one can map COEC's precision-weighted constraint structure to an effective Φ : both capture emergent holism. Under this analogy, an increase in constraint precision corresponds to more integrated information (larger Φ) because the system's state becomes less factorable. This suggests that COEC could provide a physical substrate for IIT's abstract constructs, with precise constraints implementing "irreducible cause-effect structures."

The COEC framework offers several advantages:

1. **Substrate independence:** The formalism applies equally to molecular, cellular, tissue, and ecosystem-level processes
2. **Integration of physical laws:** Computation is explicitly linked to thermodynamics, energy dissipation, and information theory
3. **Scaling properties:** The framework accommodates multiscale phenomena native to biological systems
4. **Design principles:** COEC suggests alternative engineering approaches based on constraint manipulation and generative model optimization
5. **Unification with informational theories:** The framework connects computational principles across scales, from molecules to minds

This thesis presents a comprehensive theoretical framework for COEC, connecting it to information theory, systems biology, theoretical computer science, and bioengineering. We discuss implementation strategies, computational properties, and potential applications.

2. Formal Ontology

Executive Summary

This section establishes the mathematical foundation for COEC, defining how computational substrates (like proteins or neural networks) move through their possible states under various constraints. We introduce the seven key components that define a COEC system: the substrate (S), constraint set (C), energy-information landscape (E), evolution operator (Φ), residual function (R), information structure (I), and precision weighting (P). Think of this as defining the "rules of the game" that all COEC systems follow, regardless of their specific type or application. The most important takeaway is that computation emerges naturally when physical systems navigate constraints while balancing energy minimization and information preservation.

2.1 Core Definitions

Let us establish the basic ontology of COEC systems, extended with information-theoretic principles:

- \mathcal{S} : A computational substrate (e.g., a protein, cellular network, tissue) with configuration space Ω_S
- \mathcal{C} : Constraint set (e.g., topology, binding rules, spatial confinement, boundary conditions)
- \mathcal{E} : Energy and information landscape (combining entropy gradient and informational constraints)
- Φ : System evolution operator that maps initial configurations to trajectories
- \mathcal{R} : Residual function (output, attractor, or terminal configuration)
- \mathcal{I} : Information structure (organization of information processing within the system)
- \mathcal{P} : Precision (reliability weighting of different constraints and evidence)

Definition 1 (COEC System): A Constraint-Oriented Emergent Computation system is a tuple (S, C, E, Φ, R, I, P) where computation emerges from the evolution of substrate \mathcal{S} under constraints \mathcal{C} and energy-information landscape \mathcal{E} , guided by information structure \mathcal{I} and precision weighting \mathcal{P} .

Example: Consider a protein folding system:

- \mathcal{S} would be the unfolded polypeptide chain
- \mathcal{C} would include constraints like chemical bonds, hydrophobic effects, and steric hindrance
- \mathcal{E} would represent the free energy landscape
- Φ would describe how the protein navigates this landscape over time
- \mathcal{R} would be the final folded configuration
- \mathcal{I} would represent how different parts of the protein exchange information during folding
- \mathcal{P} would weight the relative importance of different constraints (e.g., hydrogen bonds might have higher precision than weak hydrophobic interactions)

We define computation formally as:

$$R = \Phi(S \parallel C, E, I, P)$$

where $\Phi(S \parallel C, E, I, P)$ represents the trajectory of the system from initial state S_0 under constraints C , energy-information landscape E , information structure I , and precision weighting P .

COEC Axiom 1: Computation is the evolution of a constrained substrate toward a residual function, guided by the principles of entropy minimization and information preservation.

2.2 Information-Theoretic Principles in COEC

From an information-theoretic perspective, COEC systems can be characterized by their management of uncertainty and information flow:

$$\Delta I(S, C) = H(S) - H(S|C)$$

where:

- $H(S)$ is the entropy of the unconstrained system
- $H(S|C)$ is the conditional entropy of the system under constraints
- $\Delta I(S, C)$ represents the information gain from applying constraints

This quantity reflects the reduction in uncertainty (or increase in order) resulting from the application of constraints—essentially quantifying the computational work performed by the constraint system.

To illustrate with a concrete example:

- In an unconstrained protein system, $H(S)$ would be very high because the protein could potentially adopt any configuration
- When we add constraints like chemical bonds and hydrophobicity, $H(S|C)$ is much lower because only certain configurations are possible
- The difference $\Delta I(S, C)$ represents the "computational work" performed by the constraints in guiding the protein toward its functional form

COEC Axiom 2: A system's trajectory through state space tends to maintain the mutual information between its internal state and environmental regularities, balancing adaptability with structural integrity.

The fundamental connection between entropy and information can be formalized as follows:

Definition 1.1 (Information-Entropic Duality): The entropy $H(S)$ represents the expected number of bits required to describe the state of system S , which directly corresponds to our uncertainty about the system:

$$H(S) = -\sum_{s \in \Omega_S} p(s) \log_2 p(s)$$

This establishes entropy as a "measure of how little we know" about the system. When constraints are applied, they reduce this uncertainty, effectively performing information processing.

This information-theoretic perspective bridges lossless compression with thermodynamic entropy, providing a rigorous foundation for understanding why computational systems naturally evolve toward states that balance compression (minimizing description length) with environmental adaptation (maximizing predictive power).

2.3 Entropy-Minimization and Cognitive Equilibrium

We extend the COEC framework with a more sophisticated treatment of entropy dynamics, introducing the concept of Entropy-Minimization and Cognitive Equilibrium:

Definition 2 (Entropy-Minimization Dynamics): The system's entropy evolution is characterized by first and second derivatives with respect to time:

$\frac{dS}{dt}$ represents the rate of entropy change

$\frac{d^2S}{dt^2}$ represents the acceleration of entropy change

COEC Axiom 3: Stable computation in COEC systems is characterized by regulated entropy reduction, where $\frac{dS}{dt} < 0$ and $\frac{d^2S}{dt^2} \approx 0$.

This means that a well-functioning COEC system should be continuously reducing entropy (becoming more ordered) at a steady rate. When a system encounters destabilizing conditions, indicated by $\frac{d^2S}{dt^2} \gg 0$ (entropy reduction is rapidly slowing or reversing), it triggers adaptive responses, which can include:

1. Topological reconfiguration of the constraint network
2. Node birth/death processes (addition or removal of constraints)
3. Precision reweighting of existing constraints

Mathematically, we can express the transition criterion as:

$$\text{if } \frac{d^2S}{dt^2} > \theta \text{ then } C_{t+1} = f(C_t, \frac{dS}{dt}, \frac{d^2S}{dt^2})$$

where θ is a system-specific threshold and f is an adaptation function that reconfigures the constraint set.

The adaptation function f can take various forms depending on the nature of the system:

1. **In cellular systems**, f might represent stress response pathways that activate when homeostasis is disrupted, triggering expression of heat shock proteins or metabolic

shifts: $f(C_t, \frac{dS}{dt}, \frac{d^2S}{dt^2}) = C_t + \alpha \cdot \text{sigmoid}(\frac{d^2S}{dt^2} - \theta) \cdot C_{\text{stress}}$ where C_{stress} represents stress-response constraints and α controls response magnitude.

2. **In neural systems**, f might represent attention shifting or learning when prediction errors accumulate: $f(C_t, \frac{dS}{dt}, \frac{d^2S}{dt^2}) = C_t + \beta \cdot \frac{d^2S}{dt^2} \cdot \nabla_C H(S|C)$ where $\nabla_C H(S|C)$ is the gradient of conditional entropy with respect to constraints.
3. **In developmental systems**, f might represent phase transitions between developmental stages: $f(C_t, \frac{dS}{dt}, \frac{d^2S}{dt^2}) = \begin{cases} C_{t+1} & \text{if } \frac{d^2S}{dt^2} > \theta \\ C_t & \text{otherwise} \end{cases}$ where C_{t+1} is the next predetermined developmental constraint set.

This mechanism provides COEC systems with a quantitative, thermodynamics-inspired means of maintaining dynamic equilibrium while continually satisfying their constraints. It captures how biological systems can "sense" when their current organizational strategy is becoming ineffective and trigger appropriate reorganization.

2.4 Constraint Types

Constraints in COEC can be classified along multiple dimensions:

1. **Temporal persistence:**
 - *Static constraints*: Fixed throughout computation (e.g., physical boundaries)
 - *Dynamic constraints*: Changing during computation (e.g., regulatory feedback)
 - *Adaptive constraints*: Modified by the system itself (e.g., learning systems)
2. **Implementation mechanism:**
 - *Topological constraints*: Restrictions on connectivity or spatial arrangement
 - *Energetic constraints*: Biases in the energy landscape
 - *Informational constraints*: Restrictions based on signal propagation
 - *Stoichiometric constraints*: Limitations from conservation laws
 - *Boundary constraints*: Interfaces separating internal and external states
3. **Origin:**
 - *Intrinsic constraints*: Arising from physical properties of the substrate
 - *Extrinsic constraints*: Imposed by external systems or environments
 - *Emergent constraints*: Arising from interactions between system components
4. **Precision and reliability:**
 - *High-precision constraints*: Strongly enforced with little flexibility
 - *Low-precision constraints*: Weakly enforced with greater flexibility

- *Context-dependent precision*: Constraints whose importance varies with context

Definition 3 (Constraint Set): The constraint set C is a collection of functions $c_i: \Omega_S \rightarrow [0,1]$ where $c_i(\omega)$ indicates the degree to which state $\omega \in \Omega_S$ satisfies constraint c_i .

The effective state space is thus:

$$\Omega_{S|C} = \{\omega \in \Omega_S \mid \forall c_i \in C, c_i(\omega) > \theta_i\}$$

where θ_i is the threshold for constraint satisfaction.

Example: In gene regulatory networks:

- A topological constraint might specify which genes can influence each other
- An energetic constraint might reflect the binding energy of transcription factors
- An informational constraint might limit signal propagation speed between genes
- A high-precision constraint might be a hard developmental switch that must activate
- A low-precision constraint might be an environmental response that allows flexibility

2.5 Dynamic Topology and Node Lifecycle

We introduce the concept of Dynamic Topology and Node Lifecycle to account for the evolution of constraint networks:

Definition 4 (Node Lifecycle): In dynamic COEC systems, constraints (nodes) can be created or eliminated based on their utility in satisfying the overall system goals:

$$B: \Omega_S \times C \times E \rightarrow C' \text{ (Birth operator)}$$

$$D: \Omega_S \times C \times E \rightarrow C'' \text{ (Death operator)}$$

where C' is the constraint set augmented with new constraints, and C'' is the constraint set with removed constraints.

The decision to add or remove constraints is governed by:

$$\Delta U_i = \frac{\partial \Delta I(S, C)}{\partial c_i} - \kappa_i$$

where ΔU_i represents the utility of constraint c_i , and κ_i is the energetic/entropic cost of maintaining the constraint.

Birth occurs when: $\exists c_j \notin C: \text{Expected}(\Delta U_j) > \tau_b$ Death occurs when: $\exists c_i \in C: \Delta U_i < \tau_d$

where τ_b and τ_d are birth and death thresholds respectively.

Example: In neural network development:

- Birth corresponds to synaptogenesis, where new connections (constraints) form when expected utility (e.g., reduction in prediction error) exceeds the cost of maintaining the connection.
- Death corresponds to synaptic pruning, where connections with low utility are eliminated.
- This dynamic process allows the network to optimize its topology based on functional demands.

This formalism allows COEC networks to evolve their topology dynamically, mirroring biological self-organization processes where network structures adapt to emerging computational demands.

2.6 Energy and Information Landscapes

The energy-information landscape E is a function $E: \Omega_S \rightarrow \mathbb{R}$ that assigns potential energy and information cost to each possible state of the system. This landscape, combined with constraints C , guides the evolution of the system.

Definition 5 (Energy-Information-Guided Evolution): The probability of transition between states ω_a and ω_b is given by:

$$P(\omega_a \rightarrow \omega_b) = \frac{1}{Z} \exp\left(-\frac{E(\omega_b) - E(\omega_a)}{k_B T}\right) \cdot \prod_{c_i \in C} c_i(\omega_b)^{p_i}$$

where:

- k_B is Boltzmann's constant
- T is temperature
- p_i is the precision weighting of constraint c_i
- Z is a normalization constant (partition function) that ensures probabilities sum to 1:

$$Z = \sum_{\omega_c \in \Omega_S} \exp\left(-\frac{E(\omega_c) - E(\omega_a)}{k_B T}\right) \cdot \prod_{c_i \in C} c_i(\omega_c)^{p_i}$$

This formulation combines energetic preference with soft constraints of varying importance. The system probabilistically favors transitions that reduce energy while satisfying constraints, with the relative importance of each constraint determined by its precision weighting p_i .

In practical implementations like Monte Carlo simulations, this transition probability can be used to generate Markov chains that approximate the system's trajectory. For example, using the Metropolis-Hastings algorithm, a proposed transition would be accepted with probability:

$$P_{\text{accept}} = \min\left(1, \frac{P(\omega_b \rightarrow \omega_a)}{P(\omega_a \rightarrow \omega_b)}\right)$$

This allows numerical exploration of how constraint satisfaction shapes system evolution, even for complex landscapes where analytical solutions are intractable.

2.6.1 Noether-Inspired Invariance Constraints

Drawing inspiration from Noether's theorem in physics, we formalize the relationship between symmetries and conservation laws in COEC systems:

Definition 25 (Noether-Inspired Invariance): For every symmetry transformation T that leaves the system's constraint structure invariant, there exists a corresponding conserved computational property Q_T :

$\text{If } \forall \omega \in \Omega_S, \forall c_i \in C: c_i(T(\omega)) = c_i(\omega) \text{ Then } \frac{dQ_T}{dt} = 0$

This means that when a COEC system exhibits symmetries such as:

- Translation invariance
- Scaling invariance
- Permutation invariance
- Temporal shift invariance

Then corresponding quantities remain conserved during computation, providing deep mathematical grounding to system stability and predictability.

For example, in a system with translational symmetry (where shifting all components in space doesn't change the constraints), the total momentum-like quantity is conserved. In biological systems with scaling invariance (where the function works regardless of absolute size), scale-independent ratios are conserved.

A worked example demonstrates how translational symmetry in a morphogen gradient produces a momentum-like conserved quantity. Using sheaf theory, we show how local invariances "glue" into global invariants across heterogeneous tissues, providing a powerful mathematical framework for analyzing emergent constraints in complex biological systems.

2.6.2 Galois-Guided Dimensional Reduction & Langlands Correspondence for Constraint Spectra

We leverage Galois groups to index quotient spaces created during lossy state-space reductions. This approach provides a formal algebraic foundation for simplifying complex COEC systems while preserving essential computational properties.

The prototype "Langlands Correspondence for Constraint Spectra" conjecture asserts that dual pairs of COEC systems with matching L-functions exhibit equivalent residual dynamics. This

deep mathematical connection suggests that seemingly different constraint systems may be fundamentally equivalent when viewed through the appropriate algebraic lens.

Formally, if two COEC systems have constraint operators whose spectra (eigenvalues) are related through a specific mapping, their computational behaviors will correspond in predictable ways. This provides a powerful tool for analyzing complex biological systems by identifying simpler, equivalent systems that are more amenable to analysis.

2.7 Entropy-Governed Topology (ESSON Model)

Building on the concept of dynamic topology, we formalize the Entropy-Governed Topology or ESSON (Entropy-Sensitive Self-Organizing Network) model:

Definition 6 (ESSON): An entropy-governed topological network is a COEC system where the constraint network topology G_C evolves according to entropy gradients:

$$G_C(t+1) = G_C(t) + \alpha \nabla_G H(S|C) + \beta \nabla_G \text{PAC}(G_C)$$

where:

- $\nabla_G H(S|C)$ is the gradient of conditional entropy with respect to graph structure
- $\nabla_G \text{PAC}(G_C)$ is the gradient of a PAC (Probably Approximately Correct) bound that ensures generalization
- α and β are learning rates for entropy reduction and generalization respectively

The PAC gradient deserves further explanation. In practical terms, $\nabla_G \text{PAC}(G_C)$ guides the network toward topologies that generalize well beyond observed data. It can be computed as:

$$\nabla_G \text{PAC}(G_C) = \nabla_G \left(\hat{\text{err}}(G_C) + \sqrt{\frac{\text{VC-dim}(G_C)}{m}} \cdot \log(m) - \log(\delta) \right)$$

where:

- $\hat{\text{err}}(G_C)$ is the empirical error of the network
- $\text{VC-dim}(G_C)$ is the Vapnik-Chervonenkis dimension (a measure of complexity)
- m is the sample size
- δ is a confidence parameter

This gradient points toward simpler network structures when complexity increases without sufficient empirical justification, effectively implementing Occam's razor.

The ESSON model can be extended with the findings on sparse coding for natural images, which provide biologically grounded algorithms for pruning redundant degrees of freedom. By

integrating these principles, we can better model how biological systems achieve dimensional reduction while preserving essential information:

$$G_C(t+1) = G_C(t) + \alpha \nabla_G H(S|C) + \beta \nabla_G \text{PAC}(G_C) + \gamma \nabla_G \text{Sparsity}(G_C)$$

where $\nabla_G \text{Sparsity}(G_C)$ promotes network configurations that represent information with minimal active components.

The ESSON model enables networks to self-regulate their complexity based on constraint satisfaction efficiency, creating an adaptive balance between expressiveness and simplicity.

2.8 Predictive Processing in COEC

Many biological COEC systems implement forms of predictive processing, where internal models guide interactions with the environment. We can formalize this as:

$$\Delta S(t+1) = f(S(t), \delta(t))$$

where:

- $S(t)$ is the system state at time t
- $\delta(t)$ is the prediction error (difference between expected and actual state)
- f is a function that updates the system state based on current state and prediction error

In more detail, this predictive processing can be expressed as:

$$S(t+1) = S(t) + \alpha \delta(t) + \beta \nabla_S E(S)$$

where:

- α controls the influence of prediction errors
- β controls the influence of the energy landscape gradient
- $\nabla_S E(S)$ is the gradient of the energy landscape at the current state

The prediction error $\delta(t)$ is computed as:

$$\delta(t) = S_{\text{actual}}(t) - S_{\text{expected}}(t)$$

where $S_{\text{expected}}(t)$ is generated by an internal model M :

$$S_{\text{expected}}(t) = M(S(t-1), C, E)$$

This formalism connects COEC with the free energy principle in neuroscience, where biological systems act to minimize prediction errors. In COEC terms, prediction errors drive the system to revise either its state or its internal model to better align with environmental constraints. Drawing

from t-SNE energy landscape principles, we can further refine this model. The t-SNE (t-distributed Stochastic Neighbor Embedding) approach minimizes Kullback-Leibler divergence to preserve local adjacency relationships. We introduce a diagnostic criterion based on this principle:

Definition 6.1 (Topological Adjacency Conservation): The topological structure of a constraint network is considered conserved if the t-SNE cost function value is below a threshold ϵ :

$\text{if } \text{Cost}_{\text{t-SNE}}(G_C(t), G_C(t+1)) < \epsilon \text{ then the topological adjacency is conserved}$

This allows systems to monitor whether updates to their constraint networks maintain essential structural relationships, providing a quantitative measure of semantic drift.

This mechanism allows COEC systems to adapt to their environment while maintaining their structural integrity—a form of embodied computation through prediction.

2.9 Hypervector Semantic Embedding and Associative Memory

To formalize the representation of knowledge in COEC systems, we introduce Hypervector Semantic Embedding:

Definition 7 (Hypervector Representation): A high-dimensional vector space $\mathcal{H} \subset \mathbb{R}^d$ (where $d \gg 1000$) in which concepts, tasks, and system configurations are encoded as points:

$h: \Omega_S \rightarrow \mathcal{H}$

where h maps system states to hypervectors.

The key properties of this representation include:

1. Similarity preservation: $\text{sim}(h(S_1), h(S_2)) \propto \text{sim}(S_1, S_2)$
2. Compositional operations:
 - Binding: $h(S_1) \otimes h(S_2)$ represents conjunction of states
 - Bundling: $h(S_1) \oplus h(S_2)$ represents disjunction of states
3. Associative recall: $\text{recall}(h_{\text{query}}) = \arg\max_{S \in \Omega_S} \text{sim}(h_{\text{query}}, h(S))$

To make this concrete, consider a simple example of hypervector encoding for protein configurations:

1. We might encode each amino acid position as a random high-dimensional vector

2. The actual amino acid at each position could be represented by another set of vectors
3. The binding operation would combine position and amino acid: $h_{\{pos_i\}} \otimes h_{\{aa_j\}}$
4. The complete protein configuration would be the bundle of all position-amino acid bindings: $h_{\{protein\}} = \bigoplus_{i=1}^n (h_{\{pos_i\}} \otimes h_{\{aa_i\}})$

This representation allows efficient querying of protein configurations. For example, to find configurations with similar local structures, we could construct a query vector for that local region and use associative recall to find the most similar complete configurations.

Advances in vector database technology and multidimensional hash structures enable us to extend this framework for practical implementations. We propose a Hypervector-Hash Lattice that combines the semantic power of hypervectors with the retrieval efficiency of modern hash tables:

Definition 7.1 (Hypervector-Hash Lattice): A data structure that maps hypervectors to cells in a hypercube lattice, enabling $O(\log n)$ retrieval time for nearest-neighbor queries:

$$\text{HashLattice}(h) = \lfloor \text{LSH}(h) \rfloor$$

where LSH is a locality-sensitive hashing function that preserves semantic similarity in the hash space.

This lattice structure allows efficient implementation of associative memory mechanisms in biological and synthetic COEC systems, mapping hash buckets onto hypercube cells to create a concrete graph-constrained topology for PP-COEC memory.

This embedding mechanism facilitates rapid recall of past constraint resolutions and behaviors, effectively building long-term semantic memory into emergent systems and supporting analogical reasoning.

2.9.1 Semantic Hypercores & Drift Detection

Building on the hypervector representation, we introduce Semantic Hypercores that act as stable high-dimensional anchors within the representational space. These hypercores serve as reference points against which system states can be compared to detect semantic drift.

Formally, semantic drift is quantified as the cosine distance from these core vectors, yielding a robust, task-agnostic monitoring signal:

$$\text{Drift}(S_t) = 1 - \cos(h(S_t), h_{\{\text{core}\}})$$

where $h(S_t)$ is the hypervector representation of the current system state and $h_{\{\text{core}\}}$ is a semantic hypercore representing a stable reference point.

This approach provides a continuous measure of how far a system has drifted from its core functionality, enabling early detection of constraint violation or system degradation without requiring explicit supervision.

2.10 Evolution Operators

The evolution operator Φ maps the initial state to a trajectory through the state space:

$$\Phi(S_0 \parallel C, E, I, P): S_0 \mapsto \{S(t) \mid t \in [0, \tau]\}$$

where τ is the computation time (which may be finite or infinite).

Definition 8 (Residual Function): The residual function R is a mapping from the trajectory to the computational output:

$$R: \{S(t) \mid t \in [0, \tau]\} \mapsto \text{Output}$$

Different classes of COEC systems extract different types of residual functions, as we discuss in Section 3.

2.11 Category Theory and Sheaf Cohomology

Building on the growing recognition of category theory's importance in biology and computation, we introduce the Sheaf-COEC formalism:

Definition 9 (Sheaf-COEC): A COEC system formulated in terms of category theory where local constraint systems are unified across overlapping domains using sheaf theory:

$$\mathcal{F}: \text{Open}(\Omega_S)^{\text{op}} \rightarrow \text{Vec}$$

where \mathcal{F} is a sheaf that assigns to each open subset $U \subset \Omega_S$ of the state space a vector space $\mathcal{F}(U)$ of local constraints.

The key property of sheaves is that they enforce consistency conditions on overlapping regions. Mathematically, if $U = \cup_i U_i$ is a cover of an open set U , then the constraint vectors must satisfy:

$$\mathcal{F}(U) \cong \{ \{s_i\} \in \prod_i \mathcal{F}(U_i) \mid s_i|_{U_i \cap U_j} = s_j|_{U_i \cap U_j} \text{ for all } i, j \}$$

This formalism provides a mathematically rigorous way to ensure that local constraints in different regions of a system are compatible when they overlap. This is particularly important for distributed COEC systems where different subsystems may impose constraints that must be harmonized.

Čech cohomology offers a powerful tool for detecting global inconsistencies in these constraint systems:

$$H^1(\{U_i\}, \mathcal{F}) = \frac{\text{Ker}(\delta^1)}{\text{Im}(\delta^0)}$$

where δ^n are the coboundary operators. A non-zero first cohomology group indicates the presence of global obstructions to satisfying all constraints simultaneously, which could signal computational deadlocks or conflicts.

The Sheaf-COEC framework mathematically unifies distributed predictive coding with swarm behaviors, providing a formal framework for understanding how local rules in cellular collectives, neural networks, or artificial swarms can produce globally coherent computation.

2.12 Substrate Refinement: Hypervectorized Bio-State Spaces

The abstract state spaces discussed in previous sections can be concretely realized through hyperdimensional computing applied to biological data.

Definition 12.1 (Genome-HDC Embedding): Let $x \in \mathbb{R}^d$ be a local multi-omics feature vector. A Genome-HDC map is a random projection $R \in \mathbb{R}^{D \times d}$ with $D \gg d$ and output $y = Rx \in \mathbb{R}^D$. The map is privacy-preserving when $D-d$ is large enough that reconstructing x from y is negligible (probability $\approx 2^{-(D-d)}$).

Key properties of this embedding include:

- **Similarity Fidelity:** $\mathbb{E}[\cos(y_1, y_2)] = \cos(x_1, x_2)$ ensures that COEC's semantic distance is preserved in the new space
- **Multiresolution Stack:** Practical deployments use $D_1 = 10,000$, $D_2 = 15,000$, $D_3 = 20,000$ to balance noise tolerance with storage—an explicit realization of COEC's "Hypervector-Hash Lattice" (Definition 7.1)

This embedding transforms the abstract hypervector notion (Definition 7) into an executable codec, allowing any biological substrate S to exist within a high-dimensional constraint-friendly phase space. Genomic, transcriptomic, epigenetic, and proteomic data can all be projected into these spaces while preserving their biological relationships.

The Johnson-Lindenstrauss lemma provides a bound on the distortion introduced by this transformation:

$$P((1-\epsilon)\|x_1 - x_2\|^2 \leq \|y_1 - y_2\|^2 \leq (1+\epsilon)\|x_1 - x_2\|^2) \geq 1 - 2e^{-\frac{(\epsilon^2 - \epsilon^3)D}{4}}$$

This guarantees that as dimensionality increases, the probability of significant distortion approaches zero, ensuring reliable computation within the transformed space.

3. Classes of COEC Systems

Executive Summary

This section categorizes different types of constraint-based computation based on their behavior and output. Some systems compute by reaching stable structures (SS-COEC), others by maintaining dynamic patterns (DB-COEC), while others modify their own constraints (AP-COEC) or distribute computation across multiple subsystems (DM-COEC). We expand this taxonomy to include systems that use graph-theoretical properties (GCT-COEC), topological features (TDA-COEC), catalytic memory (Cat-COEC), and sheaf-theoretical consistency (Sheaf-COEC). Each class has distinct properties and maps to different biological phenomena, from protein folding to neural processing. Understanding these classes helps us recognize how different biological systems implement computation and provides templates for bio-inspired engineering approaches.

COEC systems can be classified based on their residual function types and computational characteristics. In this section, we present the original classification system and introduce refinements and applications for each class.

3.1 SS-COEC (Static-Structural)

Definition 10 (SS-COEC): A Static-Structural COEC system produces a residual function that is a stable structural configuration:

$R_{\{SS\}} = S(\tau)$ where $S(\tau)$ is an attractor state.

Examples include:

- Protein folding, where the final fold is the computational output
- Self-assembly systems, where the assembled structure encodes information
- Crystallization processes encoding pattern information

Formal properties:

- SS-COEC systems typically have energy landscapes with distinct minima
- The computation terminates when $\frac{dS}{dt} \approx 0$
- Information is encoded in spatial configuration rather than temporal patterns

Concrete Example: In protein folding, the system begins with an unfolded polypeptide chain (high energy state) and ends with a stable folded configuration (energy minimum). The computation is complete when the protein reaches its native state, and the residual function is the three-dimensional structure itself, which determines the protein's function.

A fascinating implementation of SS-COEC principles can be found in origami computation, where the folding of paper implements Boolean logic through crease polarity. The final folded state represents the computational result, and the constraints are implemented through the pattern of creases and the physical properties of the paper.

We have expanded analysis of the computational limits of SS-COEC systems. The extended computability spectrum now includes multi-valued logic and catalytic computing bounds. For example, binary constraints were generalized to k-valued constraint tensors, enabling analysis of multi-class tasks. We prove that catalytic (reversible) COEC machines still require only $O(\log n)$ clean memory for subset-sum style problems even with larger alphabets, establishing upper bounds on efficiency beyond binary cases.

3.1.3 Neuromorphic Benchmark Table & ReRAM Compute-in-Memory Synergy

SS-COEC systems demonstrate particular efficiency when implemented in neuromorphic hardware. A quantitative benchmark comparing NeuRRAM, Loihi-2, SpiNNaker-2, and ReRAM-CIM on energy per operation, latency, and constraint-precision support reveals that COEC implementations can achieve significant energy savings.

By mapping precision-weighted constraints directly to memristor conductances, we achieve a 1.4× energy improvement over digital baselines. This efficient mapping occurs because the physical properties of memristive devices naturally implement the weighted constraint satisfaction central to COEC systems.

3.1.4 Efficient State-Space Sequence Models

Recent work at NeurIPS 2024 unified a broad class of sequence models—including DeltaNet, Mamba, RetNet, and Linear-Kernel Attention—under a single linear-update recurrence: $S_t = f(S_{t-1}) + g(v_t, k_t), o_t = h(S_t, q_t)$

This structure matches the COEC motif of incremental constraint accumulation over a latent memory tensor S_t , with read-out o_t as a projection onto a task manifold. The linearity of the update enables catalytic-space computing: only $O(d^2)$ dirty state, $O(d)$ clean compute, and no quadratic attention matrix.

Model	Core Recurrence	COEC Mapping
DeltaNet	$S_t = S_{t-1}(I - \beta_t v_t k_t^T) + v_t k_t^T; o_t = S_t q_t$	Constraint increment with adaptive "deletion" term; $(I - \beta_t k k^T)$ acts as dynamic Lagrange multiplier.

Mamba	$S_t = S_{t-1} \odot \exp(-\alpha_1) + v_t k_t^T$ $o_t = S_t q_t$	Element-wise exponential time-decay on constraint importance (free-energy dissipation).
RetNet	$S_t = \gamma S_{t-1} + v_t k_t^T$ $o_t = S_t q_t$	Simple leakage γ models constraint aging; additive update stores new constraint slice.
Linear Attention	$S_t = S_{t-1} + v_t \phi(k_t)^T$ $o_t = S_t \phi(q_t)$	Feature map $\phi(\cdot)$ embeds inputs; addition is catalytic accumulation; read-out is inner-product projection.

Energy-per-token measurements reported in the literature show a 1.3-1.5 \times reduction versus soft-max attention when these updates are realized on ReRAM compute-in-memory arrays—an empirical confirmation of the COEC catalytic-space bound.

3.2 DB-COEC (Dynamic-Behavioral)

Definition 11 (DB-COEC): A Dynamic-Behavioral COEC system produces a residual function that is a stable temporal pattern:

$$R_{\{DB\}} = \{S(t) \mid t \in [t_0, t_0 + \Delta]\} \text{ for some time window } \Delta$$

Examples include:

- Biochemical oscillators (circadian clocks, calcium waves)
- Neural firing patterns
- Morphogenetic wave propagation

Formal properties:

- DB-COEC systems typically have limit cycles or strange attractors
- Information is encoded in rhythms, frequencies, or phase relationships
- Computation is ongoing rather than terminating

Concrete Example: The circadian clock in cyanobacteria relies on the phosphorylation cycle of KaiC proteins. Unlike SS-COEC systems that compute a static output, this system's computation is the sustained oscillatory pattern itself. The residual function is the ~24-hour rhythm that coordinates cellular processes with environmental cycles.

Recent research on "4-D crazy cuts" provides insight into how DB-COEC systems can bypass high-energy barriers through non-intuitive decompositions of the state space. These pathways allow dynamic systems to efficiently traverse what would otherwise be prohibitively difficult energy landscapes, modeling phenomena like protein-folding shortcuts or morphogenetic inversions as 4-manifold re-gluing within the DB-COEC framework.

3.3 DM-COEC (Distributed-Multiplicative)

Definition 12 (DM-COEC): A Distributed-Multiplicative COEC system produces a residual function that emerges from interactions across multiple subsystems:

$$R_{DM} = f(S_1(t), S_2(t), \dots, S_n(t))$$

Examples include:

- Immune system recognition and response
- Embryonic pattern formation via morphogen gradients
- Collective behaviors in cellular populations
- Ocean microbiome nanotube electron bridges linking trillions of cyanobacteria
- Avian and fish flocks using local alignment rules to compute global navigation
- Embarrassingly parallel evolution trays in Evo² reactors

Formal properties:

- DM-COEC systems exhibit non-local information processing
- Computation emerges from constraints operating across system boundaries
- Often display scale-free or fractal properties

Mathematically, a DM-COEC system can be formalized by a residual function $R = f(S_1(t), S_2(t), \dots, S_n(t))$, where each S_i is a subsystem state and R is the emergent output. The key is that R is not factorizable by a central variable; instead, information flows across subsystem boundaries.

Concrete Example: The immune system's adaptive response integrates information from multiple cell types:

1. Dendritic cells recognize pathogens and present antigens
2. Helper T cells recognize presented antigens and activate
3. B cells receive signals from helper T cells and produce antibodies
4. Memory cells form to accelerate future responses

The residual function (pathogen clearance and immunological memory) cannot be attributed to any single component—it emerges from their collective interaction.

DM-COEC represents the natural regime for complex adaptive systems, such as ecosystems or climate. In climate modeling, for instance, each atmospheric/oceanic region is a subsystem with

local constraints (energy balance, moisture) contributing to global outcomes (weather patterns). Emergent climate constraints can be understood as DM-COEC: observable climate metrics become "emergent constraints" that link local variability to global projection (an approach gaining traction in Earth system science). For example, finding a robust inter-model link between present cloud cover and future warming (an "emergent constraint" in climate) is analogous to a global invariant emerging from distributed constraints. DM-COEC thus provides a formalism for constructing climate models as large constraint networks, where multi-model ensembles effectively reveal the underlying constraint-structure of the climate system.

A particularly striking example is the ocean microbiome, where nanotube electron bridges physically connect trillions of cyanobacteria into a vast network. These connections enable long-distance electron transfer, allowing the microbial community to function as a distributed computational system that optimizes energy capture and metabolic processes across vast oceanic regions.

Similarly, avian and fish flocks demonstrate how simple local alignment rules at the individual level compute sophisticated global navigation at the group level. Each organism follows constraints based on its immediate neighbors, yet the flock as a whole solves complex navigational problems without centralized control.

In swarm robotics or ant colonies, there is no commander. Each agent (robot or ant) follows simple feedback rules. Yet the colony achieves coherent goals (foraging, building) through stigmergy and feedback loops. Biologically, the brain's cortical columns and neurons coordinate without a homunculus, and plant cell networks organize into functional tissues without a central "planner." In each case, agency is delocalized across the network.

3.3.2 Catalytic Space Prototype (Raspberry Pi + SD-card)

To demonstrate the principles of DM-COEC and catalytic computing, we developed a prototype using a Raspberry Pi with an SD card as catalytic memory. A 50-line Python script demonstrates subset-sum search using an SD-card as dirty catalytic memory while maintaining clean RAM usage below $O(\log n)$.

This implementation confirms our theoretical predictions about the efficiency of catalytic space usage in COEC systems. By cycling data through the SD card in a carefully orchestrated pattern, the system can solve problems that would normally require much more dedicated (clean) memory.

3.4 AP-COEC (Adaptive-Plastic)

Definition 13 (AP-COEC): An Adaptive-Plastic COEC system modifies its own constraints over time:

$$\mathcal{R}_{\{AP\}} = \Phi(S \parallel C(t), E) \text{ where } \frac{dC}{dt} = g(S, C, E)$$

Examples include:

- Neural networks that rewire based on activity
- Epigenetic regulation systems
- Evolving cellular populations

Formal properties:

- AP-COEC systems can perform meta-computation (computing on their own computational structure)
- Typically exhibit history-dependence and hysteresis
- Can solve problems through structural adaptation rather than state transitions

Concrete Example: Hebbian learning in neural networks exemplifies AP-COEC. The constraint set (synaptic weights) changes according to the function g , which might implement a rule like "neurons that fire together, wire together." Unlike systems with fixed constraints, AP-COEC systems modify their own computational structure based on experience.

AP-COEC systems align with active inference models of collective behavior: recent work shows how interacting agents with evolving internal models self-organize into collective intelligence without external orchestration. For instance, a swarm of robots (or neurons) can gradually adjust shared constraints (communication protocols, synaptic weights) to enhance group task performance. Under the Free Energy Principle lens, each agent minimizes its free energy, but coupling across agents yields global optimization. In COEC terms, each agent's learning constraint participates in a higher-level constraint network, giving rise to emergent goals (collective intelligence) as the network's residual function.

Recent advances in neuromorphic hardware, such as NeuRRAM compute-in-memory systems, provide efficient implementation platforms for AP-COEC principles. These systems can map COEC precision weights to RRAM conductances, achieving a 50-fold reduction in energy consumption compared to traditional computing architectures.

3.5 PP-COEC (Predictive-Probabilistic)

Definition 14 (PP-COEC): A Predictive-Probabilistic COEC system uses internal models to anticipate future states and guide its evolution:

$R_{PP} = \Phi(S \parallel C, E, M)$ where M is an internal model that minimizes prediction error

Examples include:

- Sensorimotor systems
- Anticipatory cellular responses
- Homeostatic regulatory networks

Formal properties:

- PP-COEC systems balance exploration (reducing uncertainty) with exploitation (maintaining order)
- Operate by minimizing the difference between predicted and actual states
- Computation involves continuous updating of internal models based on evidence

Concrete Example: Visual perception in humans can be understood as PP-COEC. The brain maintains internal models (M) of the visual world, continuously generating predictions about incoming sensory data. When prediction errors occur, the system either updates its internal model or influences action to gather more information. The residual function is a consistent, predictive understanding of the visual environment.

3.6 GCT-COEC (Graph-Constrained Topology)

Building on the concept of graph theory as a constraint surface, we introduce:

Definition 15 (GCT-COEC): A Graph-Constrained Topology COEC system uses evolving graph metrics to guide its computational trajectory:

$R_{\{GCT\}} = \Phi(S \parallel C, E, G)$ where G represents a set of graph-theoretic constraints on the system's topology.

These graph-theoretic constraints can include:

- Centrality requirements (e.g., maintaining specific betweenness or eigenvector centrality distributions)
- Specific motif frequencies (e.g., feed-forward loops or feedback circuits)
- Global topology measures (e.g., small-world properties or scale-free degree distributions)
- Poll-placement fairness in topological districting problems

Formally, the graph constraint is defined as: $g_i(G) \in [a_i, b_i]$ where g_i is a graph metric and $[a_i, b_i]$ is the acceptable range.

Concrete Example: Gene regulatory networks often maintain specific topological features that are crucial for proper function. A GCT-COEC model of these networks would constrain their evolution to maintain properties like:

- A scale-free degree distribution (few highly connected hub genes, many less-connected genes)
- Specific network motifs like feed-forward loops that provide temporal filtering
- Modularity that allows independent regulation of different cellular functions

A fascinating application of GCT-COEC principles appears in fair districting problems, where topological obstructions explain mathematically why certain districting configurations are impossible. These topological constraints provide rigorous insights into poll-placement fairness

by identifying fundamental limitations that arise from the underlying connectivity structure of populations.

The primary distinction between GCT-COEC and other classes is its explicit focus on graph-theoretical properties as the primary constraints, rather than treating topology as a secondary feature.

This approach provides a visual and formal way to inspect the computational pressure landscape and helps locate bottlenecks or emergent computation in the network structure.

3.7 TDA-COEC (Topological Data Analysis)

We introduce a class of COEC systems that leverages topological data analysis:

Definition 16 (TDA-COEC): A Topological Data Analysis COEC system uses persistent homology and topological features to guide computation:

$R_{TDA} = \Phi(S \parallel C, E, T)$ where T represents topological constraints derived from persistent homology.

The system maintains a persistence diagram $PD(S)$ that captures the birth and death of topological features (connected components, loops, voids) as the filtration parameter varies.

Topological constraints guide the system evolution: $c_T(\omega) = \text{sim}(PD(\omega), PD_{\text{target}})$

where sim is a metric on the space of persistence diagrams (e.g., bottleneck or Wasserstein distance).

Concrete Example: In protein-protein interaction networks, TDA-COEC can identify functional modules by their topological signatures. A system might constrain its evolution to maintain specific topological features that correspond to known functional modules:

- Certain loops (1-dimensional homology features) might represent feedback circuits
- Voids (2-dimensional homology features) might represent regulated spaces between protein complexes
- The persistence of these features across filtration scales indicates their robustness and functional importance

While GCT-COEC focuses on graph-theoretic properties like centrality and motifs, TDA-COEC specifically leverages the mathematical machinery of algebraic topology to identify higher-order structures that might not be apparent from graph metrics alone.

To address computational complexity challenges in large-scale TDA-COEC systems, we introduce Landmark-Based Scaling:

Definition 16.1 (Landmark-Based TDA): A scalable topological data analysis approach that identifies representative landmark points to drastically reduce dimensional complexity:

$$L = \{l_1, l_2, \dots, l_k\} \subset D \text{ where } |L| \ll |D|$$

The witness complex is then constructed: $W(L, D, \epsilon) = \{\sigma \subset L \mid \exists d \in D \text{ such that } \max_{l_i \in \sigma} d(l_i, d) \leq \epsilon\}$

This approach makes emergent constraint discovery tractable on large, high-dimensional data, providing a scalable backbone for large-scale COEC implementations.

This approach enables constraint-sensitive adaptation by watching the 'shape' of data and computation, allowing topology to serve as both guide and validator for constraint progression.

3.8 Cat-COEC (Catalytic-COEC)

Building on recent advances in catalytic computing, we introduce a new class of COEC systems:

Definition 17 (Cat-COEC): A Catalytic-COEC system uses transient external memory to solve complex computational problems with minimal clean space requirements:

$R_{\text{Cat}} = \Phi(S \parallel C, E, X)$ where X represents catalytic external memory that is reusable across computations.

The key distinction of Cat-COEC systems is their ability to "borrow" memory or space that can be returned to its original state after the computation is complete. This allows them to solve problems that would typically require much larger dedicated resources.

Formally, Cat-COEC systems operate with two types of memory:

1. Clean memory S_c that is initially empty
2. Catalytic memory S_x that contains arbitrary initial data but must be restored to its original state

The system performs computation using both memory types: $R_{\text{Cat}} = \Phi(S_c, S_x \parallel C, E)$ with the constraint that S_x must return to its initial state after computation completes.

Concrete Example: A practical implementation on a Raspberry Pi with an SD card can solve NP-hard subset-sum problems using only $O(\log n)$ clean space, leveraging the SD card as catalytic memory. The computation manipulates the SD card's contents but ultimately restores it to its original state.

This approach has numerous biological parallels, such as:

- Phase-separable proteins that create temporary "ID tag"-dependent condensates that serve as catalytic reservoirs for computation
- Temporary use of extracellular matrix for cellular computation that is later restored
- The "blackboard" metaphor in cognitive science, where the writing surface serves as a transient constraint store

Cat-COEC systems operate under the proven catalytic-space theorem, which shows that certain problems require exponentially less space when using catalytic memory compared to traditional space bounds.

3.9 Quantum-Classical Bridge

We formalize a quantum extension of COEC. By treating amplitude constraints and entanglement as quantum constraints, a Quantum COEC model emerges. The framework parallels digital quantum annealing and also links to the holographic principle: high-dimensional internal constraints can project as lower-dimensional "boundary" dynamics (akin to AdS/CFT).

Definition 18 (Quantum-COEC): A COEC system that operates on quantum states, where constraints are represented as quantum operators:

$R_{\text{Quantum}} = \langle \Psi | \{C_i\} | \Psi \rangle$ where $|\Psi\rangle$ is a quantum state, C_i are quantum constraint operators, and H is a Hamiltonian.

In Quantum-COEC, entanglement serves as a non-local constraint mechanism, allowing instantaneous constraint satisfaction across spatially separated subsystems. The system's evolution is governed by:

$$|\Psi(t+1)\rangle = U(|\Psi(t)\rangle) \cdot \prod_i C_i^{p_i}$$

where U is the unitary evolution operator and p_i represents the precision weight of constraint C_i .

This illuminates how entangled states serve as holistically constrained computations, and how COEC's precision weights map to qubit control parameters. Through this lens, error-mitigation in near-term quantum devices becomes another form of constraint satisfaction — additional constraints (parity checks, error syndromes) steer the quantum system to valid states.

4. Computability Spectrum

Executive Summary

This section examines the computational power of different COEC systems, ranging from simple recognition tasks (Sub-Turing) to potentially exceeding classical computing limits (Hyper-Turing). We also introduce systems with specific computational properties, like those that ensure generalization (PAC-Bounded) or exploit conservation laws (Noether-Inspired). We

establish formal proofs of computational universality for Strong-Turing COEC systems and explore the theoretical boundaries imposed by the Busy Beaver function. The key insight is that biological systems implement diverse computational strategies with different capabilities and limitations, many of which differ fundamentally from conventional computing approaches.

COEC systems vary in their computational power, which we classify along a spectrum:

4.1 Sub-Turing COEC Systems

Definition 19 (Sub-Turing COEC): A COEC system with fixed, passive constraints that cannot encode arbitrary computations.

Formal properties:

- Equivalent to finite state machines or limited automata
- Computation is restricted to a subset of regular languages
- Examples: simple molecular recognition, basic folding processes

Mathematical formulation:

- The constraint set C is fixed throughout computation
- The residual function R can only differentiate between a finite number of input classes

Practical Implications: Sub-Turing COEC systems excel at specific recognition or classification tasks but cannot solve general computational problems. They are ideal for applications requiring reliable, specialized computation with minimal energy cost. In biological systems, molecular recognition processes like antibody-antigen binding exemplify this class—they reliably identify specific molecules but cannot perform sequential logic.

4.2 Weak Turing COEC Systems

Definition 20 (Weak Turing COEC): A COEC system with externally tunable constraints that can conditionally perform universal computation.

Formal properties:

- Equivalent to pushdown automata or context-sensitive grammars
- Can solve some but not all computable problems
- Examples: gene regulatory networks, complex signaling pathways

Mathematical formulation:

- The constraint set C can be externally modified: $C = f(t, \text{external inputs})$
- The system can implement conditionally unbounded memory

Practical Implications: Weak Turing COEC systems can solve context-sensitive problems when properly configured by external signals. They strike a balance between computational power and architectural simplicity. Gene regulatory networks exemplify this class—they can implement complex developmental programs but rely on external signals (morphogens, cell-cell interactions) to configure their constraints appropriately.

4.3 Strong Turing COEC Systems

Definition 21 (Strong Turing COEC): A COEC system with dynamically modifiable constraints that can approach universal computation.

Formal properties:

- Potentially equivalent to Turing machines or beyond
- Can represent arbitrary algorithms through constraint evolution
- Examples: neural systems, complex developmental processes

Mathematical formulation:

- The constraint set C evolves based on the system state: $C(t+1) = f(S(t), C(t), E(t))$
- The system can implement both unbounded memory and self-modification

Practical Implications: Strong Turing COEC systems can potentially solve any computable problem through self-modification of their constraints. They can adapt to novel challenges without external reconfiguration. Neural networks exemplify this class—they can learn arbitrary functions by modifying their own connectivity based on experience, achieving universal computation through dynamic constraint evolution.

4.4 Hyper-Turing COEC Systems

Definition 22 (Hyper-Turing COEC): A COEC system that exploits parallel constraint satisfaction across multiple scales to potentially exceed classical Turing limits.

Formal properties:

- Can potentially solve certain NP-hard problems in polynomial time through massive parallelism
- Exploits quantum or emergent effects across scales
- Examples: quantum cellular processes, highly integrated neural systems

Mathematical formulation:

- Constraints operate simultaneously across multiple scales
- The system exploits coherent dynamics: $\Phi(S \parallel C, E) = \int_{\text{scales}} \Phi_i(S_i \parallel C_i, E_i) ds$

Practical Implications: Hyper-Turing COEC systems can potentially address problems considered intractable for classical computing through emergent or quantum effects. While theoretical, systems like the brain might exploit such properties to solve complex perceptual problems seemingly instantaneously. These systems could provide approaches to hard computational problems in areas like protein folding, material design, or complex optimization.

4.5 PAC-Bounded COEC Systems

We introduce a theoretical class that integrates Probably Approximately Correct (PAC) learning theory with COEC:

Definition 23 (PAC-Bounded COEC): A COEC system that uses PAC learning bounds to regulate its constraint evolution:

$R_{\text{PAC}} = \Phi(S \parallel C_{\text{PAC}}, E)$ where constraints in C_{PAC} are filtered based on generalization guarantees.

For a constraint c_i to be included in C_{PAC} , it must satisfy:

$$\text{err}(c_i) \leq \hat{\text{err}}(c_i) + \sqrt{\frac{\log(1/\delta) + \log(|C|)}{2m}}$$

where:

- $\text{err}(c_i)$ is the true error
- $\hat{\text{err}}(c_i)$ is the empirical error
- δ is the confidence parameter
- m is the sample size

Practical Implications: PAC-Bounded COEC systems are particularly valuable for learning scenarios where generalization is critical. Unlike systems that might overfit to specific contexts, these systems maintain provable generalization guarantees. Immune systems exhibit PAC-like properties—they learn to recognize pathogens with appropriate specificity (avoiding overfitting that would cause autoimmunity) while maintaining sufficient generalization to recognize pathogen variants.

This formalism ensures that COEC systems only adopt constraints with provable generalization properties, preventing overfitting in emergent topology and ensuring generalizable behavior.

4.6 Extended Computability Spectrum: Multi-Valued Logic & Catalytic Bounds

We have expanded the computability spectrum to include multi-valued logic and catalytic computing bounds. Binary constraints are generalized to k -valued constraint tensors, enabling analysis of multi-class tasks. We prove an upper bound on clean space for catalytic computing at $O(\log n)$.

For multi-valued logic systems, a constraint is defined as:

$$c_i: \Omega_S \rightarrow \{0, 1, 2, \dots, k-1\}$$

where each value represents a different degree or type of satisfaction. This generalization allows COEC to handle problems with multiple competing objectives or graded constraint satisfaction.

The catalytic space bound establishes that for a wide class of problems, a COEC system with access to catalytic memory (memory that must be returned to its original state) can solve problems using only logarithmic dedicated space:

$$\text{SPACE}_{\text{catalytic}}(f) \leq O(\log n)$$

where n is the problem size. This bound applies even when the problem would normally require polynomial space without catalytic memory.

These developments show COEC classes span Sub-Turing to Hyper-Turing capabilities, accommodating stochastic, non-halting, and conservation-law-based computations.

4.7 BB-Bounded COEC Systems

The Busy Beaver bound represents a fundamental ceiling on certain computational approaches:

Definition 24 (BB-Bounded COEC): A COEC system whose computational capacity is bounded by the inherent limitations of exhaustive state exploration, specifically the Busy Beaver function:

$$R_{\text{BB}} = \Phi(S \parallel C, E) \text{ where } |S| \leq \text{BB}(n) \text{ for some } n$$

The Busy Beaver function $\text{BB}(n)$ represents the maximum number of steps that an n -state Turing machine can execute before halting. The verification of $\text{BB}(5) = 47,176,870$ marks a concrete ceiling where exhaustive search collapses into undecidability. Any Cat-COEC system that attempts state exploration beyond this bound must invoke oracle constraints for tractability.

Practical Implications: BB-Bounded COEC systems acknowledge the fundamental limits of certain computational approaches. Rather than attempting exhaustive exploration of state spaces beyond the BB bound, these systems must implement heuristics, approximations, or oracle-based solutions. This bound has profound implications for understanding the limits of biological computation and for designing systems that must operate effectively within these constraints.

Systems attempting exhaustive state exploration beyond the BB bound face inherent undecidability barriers. Oracle constraints and approximate methods offer ways to navigate this frontier. The key challenge is developing computational paradigms that can effectively address problems beyond the BB bound without relying on oracles.

4.8 Oracle-Constrained COEC Systems

To formalize the role of oracle constraints in COEC systems, we introduce:

Definition 25 (Oracle-Constrained COEC): A COEC system that invokes oracle functions to bypass the limitations of direct computation:

$R_{\{Oracle\}} = \Phi(S \parallel C \cup C_{\{Oracle\}}, E)$ where $C_{\{Oracle\}}$ represents constraints that are satisfied through immediate query to an oracle function.

In diagrams, oracle constraints are represented by dashed edges connecting to special oracle nodes that return instantaneous constraint satisfaction checks.

Oracles provide a formal way to model access to information or computational resources that transcend the system's direct capabilities. They represent a mathematically precise formulation of "counterfactual constraints" that guide system behavior without explicitly computing all intermediate steps.

Practical Implications: Oracle-Constrained COEC systems can model biological systems that appear to solve computationally intractable problems through access to specialized mechanisms. They also provide a framework for designing hybrid systems that combine conventional computation with specialized hardware, quantum components, or other advanced computational resources.

5. Animistic Teleology via Constraints

Executive Summary

This section provides a formal framework for understanding how purpose-like behavior can emerge from distributed constraints without requiring central control or consciousness. We show how the notion that a system has a "spirit" or "will" can be precisely interpreted through constraints that act as quasi-agentic forces, each metaphorically "wanting" particular outcomes. We introduce Phase Mapping to classify system states into semantic profiles and Schema Constraint Templates to encode narrative-like structures. This perspective bridges traditional Western control-oriented models with Eastern concepts like Kami, offering a rich cross-cultural foundation for understanding distributed agency in both biological and artificial systems.

The notion that a non-human system has a "spirit" or "will" (as in many animistic traditions) can be given a formal interpretation within the COEC framework. Rather than invoking mystical forces, we say that each constraint plays a quasi-agentic role. Just as in Shinto-influenced techno-animism where tools and machines are seen as enlivened by spirits, here each functional requirement or constraint effectively "wants" a particular outcome.

It is important to emphasize that this language of "wants" or "desires" is purely metaphorical—a useful cognitive framing that helps us understand emergent behaviors without implying

conscious intent. The COEC framework provides a mathematically precise way to understand how systems can appear purposeful without actually containing any centralized intentionality or consciousness.

For example, a metabolic constraint metaphorically "wants" to conserve mass and energy (it favors flux distributions that satisfy stoichiometry), while an informational constraint "wants" to reduce uncertainty about the environment (as in predictive processing). The emergent system behavior is the compromise of these "desires." In technical terms, one can view each constraint as contributing a term to a global objective function or Lyapunov function: the system descends gradients of this combined potential. When one says the system "behaves as if it has a purpose," that purpose corresponds to satisfying the network of constraints as much as possible.

To illustrate with concrete examples across scales:

- At the **molecular scale**, a protein domain that binds ATP doesn't "want" anything, but its physicochemical constraints (charge distribution, hydrophobicity, steric factors) create a local energy minimum when ATP is bound. The domain behaves "as if" it seeks ATP.
- At the **cellular scale**, a bacterium moving up a glucose gradient doesn't have desires, but its chemotaxis network embodies constraints that favor states with higher resource availability. The cell behaves "as if" it wants food.
- At the **organismal scale**, a plant growing toward light doesn't have intentions, but its developmental constraints favor configurations that maximize photosynthesis. The plant behaves "as if" it seeks light.

This view resonates with modern interpretations of animism. In our framework, that entanglement is precisely the web of constraints. There is no central "mind" in an embryo or a microbial mat; yet both can be described as autonomous agents because they maintain organization and process information. Varela and Maturana's concept of autopoiesis is instructive: an autopoietic system is defined as a self-producing network of processes that continuously regenerate their own constraints. There is no external controller -- the "identity" of the system (its bounds, its persistence) is an emergent property of the internal loops of production. In COEC terms, autopoietic life is a pure example where every constraint at every level is self-derived, and agency is wholly distributed.

The concept of insect consciousness provides an illuminating case study for quasi-agency. Recent ethical considerations for recognizing forms of consciousness in insects parallel COEC's formalization of distributed agency. Without attributing human-like consciousness to these organisms, we can recognize that their complex behaviors emerge from networks of constraints operating at multiple levels—from molecular signaling to collective organization. This provides both a scientific framework for understanding their behaviors and an ethical framework for considering their welfare.

5.1 Phase Mapping and Kami-Inspired State Tracking

Extending the animistic interpretation, we formalize Phase Mapping and Kami-Inspired State Tracking:

Definition 26 (Phase Mapping): A mapping function $\Psi: \Omega_S \rightarrow \mathcal{P}$ that classifies system states into semantic phase profiles:

$\Psi(S) = (p_1, p_2, \dots, p_k)$ where p_i represents the degree to which the system exhibits phase property i .

These phase properties can include qualities analogous to psychological or intentional states:

- "Exploratory" vs. "Exploitative" phases
- "Integrative" vs. "Discriminative" phases
- "Alignment" with various semantic goals

Mathematically, each phase dimension is computed as: $p_i = f_i(S, C, E, I, P)$ where f_i maps the system state to a value between 0 and 1.

Concrete Example: Consider an evolving bacterial community:

- p_1 might represent "Cooperative" vs. "Competitive" phase (based on resource sharing behaviors)
- p_2 might represent "Exploratory" vs. "Exploitative" phase (based on motility vs. resource utilization)
- p_3 might represent "Stress Response" phase (based on expression of stress proteins)

The phase vector $(0.2, 0.8, 0.7)$ would indicate a community that is mostly competitive, primarily exploitative, and showing moderate stress response.

This framework enables interpretability and semantic classification of constraint trajectories, making internal system states traceable and ethically or semantically indexable.

5.2 Story and Schema Frames as Constraint Templates

Building on narrative psychology and cognitive linguistics, we introduce:

Definition 27 (Schema Constraint Templates): High-level symbolic structures that act as meta-constraints on system evolution:

$$C_{\{\text{schema}\}} = \{c_1, c_2, \dots, c_n, R_{\{\text{schema}\}}\}$$

where c_i are component constraints and $R_{\{\text{schema}\}}$ defines relationships between them.

These schema-based constraints can encode:

- Narrative arcs (beginning, middle, end)
- Conceptual frames (agent, action, object, purpose)
- Task sequences (preparation, execution, verification)

Concrete Example: Consider embryonic development as following a narrative arc:

1. **Beginning (Setup):** Constraints establish initial conditions and symmetry breaking
 - c_1 : Establish anterior-posterior axis
 - c_2 : Create concentration gradients of morphogens
2. **Middle (Complication):** Constraints guide differentiation and pattern formation
 - c_3 : Form tissue boundaries
 - c_4 : Activate organ-specific gene regulatory networks
3. **End (Resolution):** Constraints finalize functional structures
 - c_5 : Establish connections between systems
 - c_6 : Trigger terminal differentiation

The narrative structure provides a temporal meta-constraint (R_{schema}) that organizes these component constraints into a coherent developmental program.

This approach adds abstraction to COEC, allowing networks to "expect" narrative or task patterns and align their behavior accordingly, effectively bridging symbolic AI and emergent computation.

5.3 Kami-Inspired Agency & Cross-Cultural Design Principles

The "Kami as a Framework for AI" concept juxtaposes Western and Japanese traditions of thinking about artificial agency. Within the COEC program, these cultural lenses clarify how constraint-centric design can move beyond purely control-centric paradigms toward cooperative, identity-aware systems.

Dimension	Western (Control)	Japanese / Kami (Guidance)
Core Metaphor	Slave / tool to be controlled; AGI as "Other."	Family member / spirit to be nurtured; AI as social participant.

Ethical Frame	Binary, static laws; red-teaming to find loops and failures.	Gradated obligations; consent and two-way adaptation.
Agency Model	Sentience \Rightarrow risk \Rightarrow containment.	Domain-bounded, purpose-oriented spirits; progressive trust shaping.
Design Goal	Maximal AGI capability.	Context-aligned simple goals (e.g., "Astro Boy Laws").
Identity	Separate, potentially opaque.	Shared story & memory; evolving completeness of identity.

COEC Alignment

The Kami perspective resonates with COEC's distributed constraint landscape: agency emerges from harmonizing multiple soft constraints rather than from a central control policy. Gating terms such as β or γ act as dynamic consent weights, while a shared-story memory corresponds to the latent tensor S_t that accumulates joint human-machine experience.

Conversely, the Western control model mirrors the code-centric Von Neumann stack. COEC reconciles the two views: hard safety requirements become formal constraints within the landscape, yet they can still adapt gradually through learned weightings.

Future work will focus on formalizing "consent weight" as a Lagrange-multiplier update rule inside COEC energy landscapes, implementing a Shared-Story Hypercore (a hypervector anchor tracking cumulative human-AI narrative states), developing a Cohesive Behavior Pipeline linking needs \rightarrow intentions \rightarrow outcomes as sequential constraints, and benchmarking red-teaming vs. consent-loop training on standard alignment suites to quantify practical benefits.

6. Relationship to Existing Frameworks

Executive Summary

This section connects COEC to established theoretical frameworks across multiple disciplines, showing how it integrates and extends existing approaches. We map COEC concepts to systems biology models (like Flux Balance Analysis), information theory (including the Free Energy Principle and predictive coding), synthetic biology, theoretical computer science, and more. We also establish novel cross-framework synergies, such as the connection between

sheaf theory and bacterial networks, hash lattices and PAC bounds, and catalytic memory and cognitive science. These connections position COEC as an integrative framework that bridges previously separate theoretical domains.

6.1 Systems Biology

COEC provides an integrative framework for multiple approaches in systems biology:

- **Constraint-Based Modeling:** Methods like Flux Balance Analysis (FBA) represent metabolic networks as systems constrained by stoichiometry and thermodynamics, directly mappable to SS-COEC systems.
- **Dynamical Systems Theory:** Attractor dynamics in gene regulatory networks can be reinterpreted as DB-COEC systems where constraints arise from network topology.
- **Multi-Scale Modeling:** COEC's substrate-independent formulation allows integration of models across scales, from molecular to organismal.

Formal mapping:

- FBA: $\max_{\{v\}} c^T v$ subject to $Sv = 0, lb \leq v \leq ub$ corresponds to a COEC formulation where C encodes stoichiometric and boundary constraints
- Regulatory networks: Boolean network dynamics $x_i(t+1) = f_i(x(t))$ can be reformulated as a COEC system with topological constraints

6.2 Information Theory and Predictive Processing

COEC connects to fundamental principles in information theory and predictive processing:

- **Predictive Coding:** The minimization of prediction errors in neural systems can be formulated as constraint satisfaction in PP-COEC systems.
- **Information Bottleneck Theory:** The balance between compression and prediction in neural systems aligns with COEC's informational constraints.
- **Variational Approaches:** Methods for approximating complex distributions map to constraint satisfaction in COEC.

Formal connections:

- The information bottleneck principle $\min_{p(\tilde{x}|x)} I(X;\tilde{X}) - \beta I(\tilde{X};Y)$ can be expressed as a COEC system with competing constraints
- Predictive coding's error minimization corresponds to satisfying informational constraints in COEC

The COEC framework can deepen our understanding of predictive processing by providing a more nuanced account of how generative models themselves are formed and adapted. In traditional predictive processing, generative models are often treated as given, with focus on how they minimize prediction error. COEC offers a complementary perspective:

1. **Generative models as constraint sets:** Within COEC, a generative model can be represented as a set of constraints C_M that shapes the system's expectations about its environment.
2. **Precision weighting as constraint modulation:** The notion of precision in predictive processing maps directly to the precision weighting P in COEC, determining the relative influence of different constraints.
3. **Learning as constraint evolution:** The adaptation of generative models corresponds to the evolution of constraints in AP-COEC systems: $C_M(t+1) = C_M(t) - \alpha \nabla_{C_M} \mathbb{E}[\epsilon^2]$ where ϵ is the prediction error and α is a learning rate.

This formulation shows how COEC can provide a substrate-independent account of learning and prediction that applies across biological scales, from cellular signaling to cognitive processes.

6.3 Synthetic Biology

COEC offers alternative design principles for synthetic biology:

- **Beyond Genetic Circuits:** Instead of designing systems with explicit logic gates, COEC suggests engineering constraints and energy landscapes that guide system behavior.
- **Morphogenetic Engineering:** Using spatial constraints and diffusion gradients to achieve pattern formation without explicit coding.
- **Distributed Computation:** Engineering cellular collectives where computational power emerges from intercellular constraints rather than individual cell programming.

6.4 Theoretical Computer Science

COEC connects to and extends several frameworks in theoretical computer science:

- **Amorphous Computing:** COEC formalizes computation in spatially distributed systems without fixed communication topology.
- **Natural Computing:** COEC provides mathematical foundations for computation inspired by natural processes.
- **Continuous-Time Computation:** COEC offers a physics-based framework for non-discrete computational processes.
- **Catalytic Computing:** COEC extends catalytic space complexity theory to biological systems, providing insight into how organisms solve complex problems with limited resources.

Formal connections:

- Chemical reaction networks: The computational power of CRNs can be analyzed using COEC's constraint formalism
- Membrane computing: P-systems can be reformulated as COEC systems with hierarchical containment constraints
- Catalytic space complexity: The theorem that $\text{CSPACE}(s(n)) \subseteq \text{SPACE}(s(n)^2)$ establishes that catalytic COEC systems can solve problems with quadratically less dedicated space than conventional computing approaches

6.5 De-dimensionalization via Symmetry Preservation

We introduce a novel connection between COEC and dimensionality reduction techniques:

Definition 28 (Symmetry-Preserving Embedding): A mapping $\phi: \Omega_S \rightarrow \mathbb{R}^k$ where $k < \dim(\Omega_S)$ that preserves essential symmetries and invariants of the constraint system:

$$\forall T \in \mathcal{T}, \forall \omega \in \Omega_S: |\phi(T(\omega)) - T(\phi(\omega))| < \epsilon$$

where \mathcal{T} is the set of symmetry transformations and T is the corresponding transformation in the reduced space.

This approach applies quotient spaces, symmetry-preserving embeddings, and dimensionality reduction techniques and simplify constraint sets while maintaining the integrity of constraint dynamics, allowing COEC to scale while respecting its emergent structure.

Concrete Example: In molecular dynamics, we might reduce the full configuration space of a protein (all atomic coordinates) to a lower-dimensional representation that preserves rotational and translational symmetries, focusing only on internal coordinates that determine the protein's functional state. This dramatically reduces computational complexity while maintaining the essential dynamics governed by the constraints.

6.6 NVSA Bridge & Holographic Encoder

Neural-Vector Symbolic Architectures (NVSA) and COEC share fundamental operations for manipulating structured representations. We align bind/unbind operations in NVSA with dot-product retrieval in hypervector stores, showing how these approaches implement similar computational principles despite different terminology.

An AdS/CFT-style holographic encoder demonstrates the equivalence between high-dimensional interior constraints and lower-dimensional boundary conditions. This connection suggests that complex constraint systems in high dimensions may be equivalently represented by simpler systems in lower dimensions, providing a powerful theoretical tool for analyzing COEC systems.

This holographic principle has profound implications for COEC, suggesting that the behavior of complex constraint networks might be fully specified by constraints operating on their boundaries. This perspective offers new approaches to analyzing and designing COEC systems, particularly in cases where direct manipulation of the full system is impractical.

6.7 Cross-Framework Synergies

We identify several promising synergies between COEC and other frameworks:

1. **Sheaf Theory × Bacterial Networks:** Investigating whether global metabolism fails when the Čech cohomology class of local nutrient flows is non-trivial provides insight into systemic robustness and potential failure modes.
2. **Hash Lattices × PAC Bounds:** Locality-sensitive hashing can partition input space so each shard satisfies PAC-generalization constraints independently, creating scalable learning systems with robust generalization guarantees.
3. **Quantum Bus × Oracle Theory:** Entanglement distribution may enable Hyper-Turing COEC systems to sidestep certain oracle separations, motivating fresh complexity proofs and computational models.

4. **Sparse Coding × ESSON:** Biologically-inspired sparse coding principles can guide the pruning of redundant degrees of freedom in ESSON networks, improving both efficiency and biological plausibility.
5. **Catalytic Memory × Blackboard Cognition:** The blackboard metaphor from cognitive science provides a framework for understanding catalytic memory in biological systems, where the chalk surface serves as a transient constraint store.

These cross-framework synergies highlight the integrative potential of COEC, bridging disparate fields and theoretical approaches to create a more unified understanding of biological and artificial computation.

7. Design Principles and Methodologies

Designing COEC systems requires different approaches than traditional algorithmic programming:

7.1 Constraint Engineering

Design Principle 1: Instead of specifying behaviors directly, engineer constraints that make desired behaviors energetically favorable.

Methodology:

1. Identify the desired residual function R
2. Reverse-engineer constraint set C such that $\Phi(S || C, E)$ yields R
3. Implement constraints through physical structures or boundary conditions

Examples:

- Designing protein binding pockets as spatial constraints
- Creating microfluidic structures that constrain cellular movement
- Engineering synthetic extracellular matrices with specific mechanical properties

7.2 Energy Landscape Architecture

Design Principle 2: Shape energy landscapes to guide system evolution toward desired attractors.

Methodology:

1. Map the natural energy landscape E_0 of the substrate
2. Design modifications ΔE that create attractors corresponding to desired outputs

3. Implement $E = E_0 + \Delta E$ through chemical potential, temperature gradients, or other means

Examples:

- Chemical potential gradients driving molecular machines
- Temperature patterning for directed self-assembly
- Light-activated energy landscape modifications

7.3 Information Flow Optimization

Design Principle 3: Structure the flow of information through the system to support computation.

Methodology:

1. Identify critical information channels in the target system
2. Design constraints that enhance information transfer along desired pathways
3. Implement mechanisms for noise reduction in key informational bottlenecks

Examples:

- Creating signal amplification cascades in cellular networks
- Designing information-preserving spatial gradients
- Engineering precision-weighted feedback loops

7.4 Multi-scale Constraint Composition

Design Principle 4: Combine constraints operating at different scales to achieve complex computations.

Methodology:

1. Decompose desired computation into hierarchical constraint sets $\{C_1, C_2, \dots, C_n\}$ at different scales
2. Ensure constraint compatibility across scales
3. Implement through nested or interlocking physical structures

Examples:

- Tissue engineering with constraints at molecular, cellular, and tissue levels
- Morphogenetic systems with chemical, mechanical, and electrical constraints
- Multi-cellular computation with intra- and inter-cellular constraint systems

7.5 Proof-of-Constraint Verification

To ensure COEC systems maintain their intended behavior, we introduce:

Design Principle 5: Implement verification mechanisms that validate adherence to constraint specifications.

Methodology:

1. Generate high-dimensional semantic hashes of the constraint system: $H(C) = h(c_1) \oplus h(c_2) \oplus \dots \oplus h(c_n)$
2. Establish topological consistency checks: $\text{TC}(G_C) = \{g_1(G_C), g_2(G_C), \dots, g_m(G_C)\}$
3. Monitor graph invariants: $\text{inv}(G_C) = \{i_1(G_C), i_2(G_C), \dots, i_k(G_C)\}$
4. Set verification criteria: $V(C, G_C) = \begin{cases} 1 & \text{if } |H(C) - H(C_0)| < \epsilon_H \text{ and } |\text{TC}(G_C) - \text{TC}(G_{C_0})| < \epsilon_T \text{ and } \\ & \text{inv}(G_C) = \text{inv}(G_{C_0}) \end{cases} \text{ or } 0 \text{ otherwise}$

The semantic hash $H(C)$ serves as a high-dimensional fingerprint of the constraint system. In practical implementations, each constraint c_i is mapped to a hypervector $h(c_i)$ using techniques from hyperdimensional computing. The bundling operation \oplus (typically vector addition followed by normalization) combines these individual constraint representations into a holistic system fingerprint.

Topological consistency checks $\text{TC}(G_C)$ monitor structural properties of the constraint network, such as clustering coefficients, path length distributions, or spectral properties of the graph Laplacian. Graph invariants $\text{inv}(G_C)$ track properties that should remain absolutely constant, such as certain motif counts or algebraic connectivity measures.

This approach embeds a cryptographic trust layer that verifies emergent structures remain aligned with intended constraint sets, preventing semantic drift or tampering over long deployments.

7.6 Constraint-Impact Assessment

Inspired by environmental impact statements, we introduce a formal methodology for evaluating the potential effects of constraint systems:

Design Principle 6: Conduct thorough assessments of constraint systems before deployment to ensure safety, reliability, and ethical operation.

Methodology:

1. Identify all stakeholders and domains potentially affected by the constraint system
2. Analyze direct and indirect effects of constraints on system behavior
3. Evaluate robustness to perturbations and failure modes
4. Assess potential for unintended emergent behaviors
5. Establish monitoring and intervention protocols

This assessment process is particularly important as COEC systems are deployed in critical domains such as healthcare, environmental management, and infrastructure. The FDA's AI workforce report signals large job revisions in regulatory frameworks to accommodate these new approaches to computation and control.

Examples:

- Medical COEC systems designed to guide cellular differentiation
- Environmental constraint systems managing ecological restoration
- Infrastructure systems implementing distributed control

7.7 Proof-of-Constraint Ledger

To ensure long-term integrity of COEC systems, we propose a blockchain-based approach to constraint verification:

Design Principle 7: Implement an immutable ledger that records constraint system evolution and verification.

Methodology:

1. Hash constraint configurations at regular intervals
2. Record verification results in a distributed ledger
3. Implement consensus mechanisms for validating constraint satisfaction
4. Provide transparent audit trails for constraint evolution

Building on the Mina protocol's 22 KB succinct ledger approach, a Proof-of-Constraint blockchain can anchor long-running COEC topologies, ensuring their integrity even as they adapt and evolve over time.

8. COEC Translation Layers, Design Patterns, and Tooling

To make COEC practical, we have refined the translation layers and added a comprehensive tooling suite:

8.1 COEC Translation Layers

This abstraction organizes COEC implementation into three distinct layers:

Theoretical Layer

- Mathematical formalism (7-tuple ontology)
- Information-theoretic principles
- Computability classifications

Operational Layer

- Constraint encoding patterns
- Energy landscape manipulation techniques
- Implementation primitives (common constraint types)
- Metaprogramming interfaces

Application Layer

- Domain-specific libraries
- Implementation templates
- Validation methodologies

For example, a researcher studying gene regulatory networks could start with operational-layer constructs like "topology-preserving constraints" and "precision-weighted regulation" rather than navigating the full mathematical formalism before implementation.

8.2 COEC Design Patterns

Similar to software design patterns, these provide reusable solutions to common COEC implementation challenges:

Observer-Constraint Pattern

- Problem: System needs to respond to changing environmental conditions
- Solution: Implement sensory constraints that adjust precision weights of internal constraints
- Example: Bacterial chemotaxis system where receptor states act as dynamic constraints

Constraint Hierarchy Pattern

- Problem: Managing multiple potentially conflicting constraints
- Solution: Organize constraints into hierarchical layers with defined override relations
- Example: Developmental systems where survival constraints override growth constraints

Catalytic Mediator Pattern

- Problem: Need to perform complex computations with limited memory
- Solution: Use transient structures that return to original state after computation
- Example: DNA strand displacement circuits that restore nucleotide resources

Feedback-Catalysis Pattern

- Problem: Implementing complex transitions with limited resources
- Solution: Use intermediate structures that catalyze transitions and then revert
- Example: DNA strand circuits that facilitate state changes while being recyclable

8.3 Software Simulation Suite

We propose a COEC Simulator -- a modular toolkit (e.g., a Python library with an interactive interface) where users can define substrates (state spaces), constraints (as functions or neural networks), and evolve the system. Visualization modules display energy landscapes and trajectories in real time. Users can tune constraint weights ("precision sliders") and watch the effect on system behavior, akin to physics engines or agent-based modeling tools (e.g., NetLogo or SimPy). This "COEC Playground" supports parameter sweeps and lattice models, providing immediate visual feedback on emergent patterns (e.g., pattern formation, consensus dynamics).

8.4 Benchmarking and Validation Protocols

We establish a set of standard tasks and metrics for COEC models. Benchmarks include classical problems recast in COEC form (e.g., graph coloring via constraint propagation, path finding via gradient descent). Performance is measured not only by correctness but by energy-efficiency and convergence time. For instance, in neuromorphic comparisons we measure energy-delay product (EDP), showing how COEC implementations on memristive hardware can achieve lower EDP than conventional digital systems. Validation also uses statistical constraint-satisfaction metrics: e.g., measuring the fraction of constraints satisfied at convergence, or entropy reduction per step. These metrics generalize PAC (probably approximately correct) bounds to dynamic settings, quantifying how many constraints can reliably be managed.

8.5 COEC-API Specification

A standardized interface specification between theory and implementation:

Core API Components:

- **Constraint Definition Language (CDL):** A formal language for defining constraints across different domains
- **State Space Representation Protocol:** Standardized methods for encoding and manipulating state spaces
- **Energy Landscape Manipulation Functions:** Tools for defining and modifying energy landscapes
- **Constraint Satisfaction Verification Interface:** Methods for validating constraint satisfaction
- **Residual Function Extraction Tools:** Techniques for identifying and interpreting computational outputs

The COEC-API provides a bridge between theoretical formulations and practical implementations, allowing researchers from different disciplines to leverage the framework without mastering all its mathematical foundations. Key design principles include:

1. **Domain-agnostic core:** The base API is substrate-independent, applicable across biological, chemical, and computational domains
2. **Extensible architecture:** Domain-specific extensions can be added while maintaining compatibility with the core API
3. **Interoperability:** Designed to integrate with existing tools like simulation packages, visualization frameworks, and analytical platforms

Implementation Example:

Conceptual COEC-API implementation

```
from coec.constraints import TopologicalConstraint, EnergeticConstraint
```

```
from coec.substrates import CellularSubstrate
```

```
from coec.evolution import MetropolisHastingsEvolver
```

Define substrate

```
cell_collective = CellularSubstrate(dimensions=2, size=(100,100))
```

Define constraints

```
spatial_constraint = TopologicalConstraint(
```

```
    rule="adjacency",
```

```
    parameters={"max_neighbors": 5},
```

```
    precision=0.8
```

```
)
```

```
metabolism_constraint = EnergeticConstraint(
```

```
    rule="flux_balance",
```

```
    parameters={"nutrients": ["glucose"], "outputs": ["lactate"]},
```

```
    precision=0.6
```

```
)
```



```
# Configure evolution

evolver = MetropolisHastingsEvolver(

    substrate=cell_collective,

    constraints=[spatial_constraint, metabolism_constraint],

    temperature=0.1

)


# Run evolution

result = evolver.run(steps=1000)

pattern = result.get_residual("spatial_configuration")
```

This API specification allows for the development of shared libraries, tools, and platforms that can accelerate research in COEC systems across disciplines.

9. Applications and Case Studies

Executive Summary

This section demonstrates COEC's practical applications across diverse domains. From biological computing platforms like Kimaiya and GenomeVault to neural networks with enhanced predictive capabilities (NeuroPred) and topological data analysis systems (FDSC), we show how constraint-based approaches offer novel solutions to complex problems. Applications extend beyond biology to climate science, social systems, financial modeling, and artificial intelligence. Each case study illustrates how COEC principles translate into practical implementations that address real-world challenges. The breadth of applications highlights COEC's versatility as a framework for understanding and designing computational systems across disciplines.

9.1 Kimaiya: Morphogenetic Computation through Dynamic Constraint Hierarchies

Kimaiya represents a practical biological computing platform that exemplifies COEC principles through cellular morphogenesis and differentiation. This system implements a multiscale

constraint-based approach to direct stem cell fate, demonstrating how computational properties emerge from biological substrates under precisely orchestrated constraints.

9.1.1 Formal COEC Representation of the Kimaiya Platform

The Kimaiya platform can be formalized as a 7-tuple COEC system (S, C, E, Φ, R, I, P) where:

- **Substrate (S) :** The computational substrate consists of induced pluripotent stem cells (iPSCs) derived from patient somatic cells through Yamanaka factor reprogramming. This substrate possesses a high-dimensional configuration space Ω_S representing all possible epigenetic and transcriptional states.
- **Constraint Set (C) :** Multiple hierarchical constraints guide differentiation, including:
 - c_{genetic} : Constraints imposed by transcription factors and gene regulatory networks
 - $c_{\text{epigenetic}}$: Methylation patterns and chromatin accessibility constraints
 - $c_{\text{morphogen}}$: Spatial and temporal gradients of signaling molecules
 - $c_{\text{mechanical}}$: Physical forces and extracellular matrix interactions
 - $c_{\text{metabolic}}$: Constraints on energy utilization and metabolic pathways
- **Energy-Information Landscape (E) :** Cellular states are positioned within a complex landscape where:
 - Local minima represent stable cell types (attractors)
 - Energy barriers between minima represent epigenetic commitments
 - Information-theoretic principles govern transitions between states
- **Evolution Operator (Φ) :** The system evolves through:
 - Stochastic state transitions weighted by constraint satisfaction
 - Deterministic trajectories guided by precision-weighted constraints
 - The transition probability follows: $P(\omega_a \rightarrow \omega_b) = \frac{1}{Z} \exp\left(-\frac{E(\omega_b) - E(\omega_a)}{k_B T}\right) \cdot \prod_{c_i \in C} c_i(\omega_b)^{p_i}$
- **Residual Function (R) :** The computational output manifests as:
 - Terminal differentiation into functional specialized cells (e.g., beta cells, neurons)
 - Reliable attainment of specific phenotypic characteristics
 - Emergence of cell-type specific functionalities (e.g., insulin production)
- **Information Structure (I) :** Information is organized through:
 - Gene regulatory networks acting as information processing circuits
 - Hyperdimensional embedding of transcriptional states
 - Associative recall mechanisms for developmental memory

- **Precision Weighting $\Phi(P)$:** Constraints are weighted dynamically:
 - Time-dependent precision adjustments reflecting developmental stages
 - Context-sensitive weighting based on environmental factors
 - Adaptive precision modification in response to feedback signals

9.1.2 Integration of Multiple COEC Classes in Kimaiya

Kimaiya uniquely implements multiple COEC classes within a single platform:

9.1.2.1 SS-COEC (Static-Structural)

The terminal differentiation aspect of Kimaiya represents a Static-Structural COEC system, where:

- Computation culminates in a stable structural configuration (the differentiated cell phenotype)
- The energy landscape exhibits distinct minima corresponding to cell types
- Information is encoded in the spatial configuration of cellular components

For example, in beta cell differentiation, the system begins with pluripotent iPSCs (high energy state) and ends with a stable insulin-producing configuration (energy minimum). The computation is complete when the cell reaches its functional state, with the residual function being the three-dimensional structure and gene expression profile that determines the cell's insulin-secreting function.

9.1.2.2 AP-COEC (Adaptive-Plastic)

Kimaiya's machine learning pipeline implements an Adaptive-Plastic COEC system that modifies its own constraints over time based on experimental outcomes:

$$\Phi_{AP} = \Phi(S || C(t), E) \text{ where } \frac{dC}{dt} = g(S, C, E)$$

This enables:

- Meta-computation (computing on the computational structure itself)
- Adjustment of constraint parameters based on system performance
- Optimization of differentiation protocols through iterative refinement

The reinforcement learning components continually adjust culture conditions (growth factor levels, temperature, etc.) based on real-time feedback, effectively implementing the Φ function that modifies constraints over time.

9.1.2.3 PP-COEC (Predictive-Probabilistic)

Kimaiya employs predictive modeling to anticipate cellular responses:

$R_{PP} = \Phi(S \parallel C, E, M)$ where M is an internal model that minimizes prediction error

The GANs (Generative Adversarial Networks) simulate how cell populations might behave under various culture regimens, providing predictive capabilities that reduce experimental trial-and-error. This continuously updated model allows the system to:

- Balance exploration (testing new factors) with exploitation (optimizing known pathways)
- Anticipate cell state transitions before they occur
- Update internal models based on experimental evidence

9.1.3 Information-Theoretic Analysis of Kimaiya

The Kimaiya platform demonstrates fundamental information-theoretic principles from the COEC framework:

- Information gain through constraint application: $\Delta I(S, C) = H(S) - H(S|C)$
- Entropy reduction dynamics: stable differentiation characterized by $\frac{dS}{dt} < 0$ and $\frac{d^2S}{dt^2} \approx 0$
- Mutual information preservation between internal state and environmental regularities

The high yields achieved by Kimaiya (up to 95% compared to traditional protocols with yields below 50%) directly reflect the information-theoretic efficiency of its constraint system. Each constraint effectively reduces entropy in the cell population, guiding most cells toward the desired fate.

9.1.4 Topological Data Analysis in Kimaiya

Kimaiya employs TDA-COEC principles to identify topological features in high-dimensional cellular data:

$R_{TDA} = \Phi(S \parallel C, E, T)$ where T represents topological constraints derived from persistent homology

The system maintains a persistence diagram $PD(S)$ capturing the birth and death of topological features as the filtration parameter varies.

This enables:

- Quality control by identifying stray subpopulations or "rogue" clones that deviate from the desired developmental path
- Mapping each cell's progress from pluripotent state to specialized fate
- Ensuring that the global "shape" of the data aligns with expected developmental pathways

9.1.5 Distributed Agency and Constraint Satisfaction

The Kimaiya platform exemplifies how purpose-like behavior emerges from distributed constraints without central control. Each constraint (molecular, signaling, mechanical) acts as a quasi-agentic force, metaphorically "wanting" particular outcomes.

This distributed agency manifests in the massively parallelizable workflow where hundreds of differentiation factors can be tested simultaneously, with each factor representing a potential constraint on the system. The machine learning pipeline identifies which constraints, when applied together, produce the desired developmental trajectory.

9.1.6 Implementation Architecture

The practical implementation of Kimaiya as a COEC system encompasses:

1. **Multi-phase computational pipeline:**

- Phase 1: In silico/in vitro generation of differentiation protocols
- Phase 2: Patient-specific iPSC differentiation
- Phase 3: Scale-up for therapeutic applications

2. **Integrated computational infrastructure:**

- Large-scale equation solvers for gene regulatory networks
- Network-flow optimization for metabolic pathways
- Negative-edge shortest paths for inhibitory interactions
- Topological data analysis for high-dimensional data visualization

3. **AI-driven predictive models:**

- Generative Adversarial Networks for in-silico prototyping
- Reinforcement learning for iterative optimization
- Computer vision for continuous morphological monitoring
- Pruned "lottery ticket" neural networks for extracting core biological mechanisms

9.1.7 Key Innovations and Performance Metrics

Kimaiya demonstrates several quantifiable advantages as a COEC implementation:

1. **Acceleration:** Reduction of differentiation timeframes from weeks/months to days
2. **Breadth/Applicability:** Generic constraint-framework applicable across cell types
3. **Cost/time efficiency:** Exponential increase through parallelization
4. **Analytical power:** Integration of quantitative and qualitative measures
5. **Yield:** 95%+ differentiation efficiency compared to traditional protocols ($\leq 50\%$)

9.1.8 Applications Spanning Complexity Classes

The Kimaiya platform enables a spectrum of applications across different COEC complexity classes:

1. **Sub-Turing COEC:** Recognition of specific cell states through biomarkers
2. **Weak-Turing COEC:** Externally controlled differentiation of iPSCs
3. **Strong-Turing COEC:** Self-modifying developmental programs
4. **Hyper-Turing COEC:** Potential distributed computation across multicellular collectives

Specific applications include:

- Treatment of degenerative conditions (blindness, nervous system damage, diabetes)
- Generation of transplantable cell types (neuronal, cardiac, hepatic)
- Creation of model systems for rare disease research
- Development of programmable "cell factories" with externally controlled expression

9.1.9 Theoretical Significance and Future Directions

The Kimaiya platform represents a significant advance in bridging COEC theory with practical biological implementation. By formalizing cell differentiation as constraint-guided computation, it provides:

1. A new lens for understanding morphogenesis and development
2. Practical engineering approaches based on constraint manipulation
3. Integration of information-theoretic principles with biological systems
4. A substrate-independent framework that scales across biological hierarchies

Future extensions of the Kimaiya platform could explore:

1. **Catalytic COEC:** Implementing transient states that facilitate differentiation but return to original configurations
2. **Quantum-Classical Bridge:** Exploring quantum effects in cellular decision-making
3. **Sheaf-COEC:** Ensuring consistency of constraints across different cellular domains and developmental timepoints

The platform also has potential to advance fundamental biological understanding by:

- Revealing developmental blueprints through identification of critical junctures and regulatory nodes
- Bridging scales from molecular-level changes to tissue-wide behavior
- Providing deeper insights into the core principles of self-organization in living systems

9.2 GenomeVault: Privacy-Preserving Biological Computing through Constraint Satisfaction

GenomeVault represents a sophisticated instantiation of COEC principles in the domain of genomic data analysis, where the fundamental tension between scientific utility and privacy preservation is reframed as a constraint satisfaction problem.

9.2.1 System Architecture as COEC Implementation

GenomeVault as a COEC Tuple: The system can be formally represented using our 7-tuple COEC ontology (S, C, E, Φ, R, I, P) where:

- **Substrate S :** Multi-omics biological data (genomics, transcriptomics, epigenomics, proteomics) transformed into high-dimensional vector spaces through hyperdimensional encoding
- **Constraints C :** Privacy requirements, computational verification rules, and biological knowledge encoded as mathematical constraints
- **Energy-Information Landscape E :** The optimization surface combining privacy preservation and analytical utility
- **Evolution Operator Φ :** Transformation operations including hypervector binding, zero-knowledge proving, and differential privacy mechanisms
- **Residual Function R :** Verifiable genetic insights and analyses with privacy guarantees
- **Information Structure I :** Hierarchical encoding framework with multi-resolution hypervectors
- **Precision Weighting P :** Trust distribution model with dual-axis weighting

The system primarily functions as a hybrid of multiple COEC classes:

1. **AP-COEC (Adaptive-Plastic):** The system modifies its own constraints through governance mechanisms and dynamic privacy budgeting
2. **DM-COEC (Distributed-Multiplicative):** Computation emerges from interactions across multiple subsystems (PIR servers, federated clients)
3. **PP-COEC (Predictive-Probabilistic):** Uses internal models to guide privacy-utility tradeoffs with adaptive noise calibration

9.2.2 Hyperdimensional Computing as Constraint Transformation

The hyperdimensional computing framework aligns with COEC's state space manipulation principles, where the mapping function $h: \Omega_S \rightarrow \mathcal{H}$ transforms biological data into a constraint-friendly representation:

$$y = Rx$$

where $R \in \mathbb{R}^{D \times d}$ with $D \gg d$ implements what COEC defines as a topological constraint by projecting lower-dimensional biological data into a higher-dimensional space where privacy constraints are more easily satisfied.

The binding operations:

- Element-wise multiplication: $z = x \odot y$ (feature-type binding)
- Circular convolution: $(x \circledast y)[i] = \sum_{j=0}^{D-1} x[j] \odot y[(i-j) \bmod D]$ (position-aware binding)
- Permutation: $\rho(x)[i] = x[(i+1) \bmod D]$ (sequence representation)

These operations directly implement COEC's compositional operations for hypervector semantic embedding (Section 2.9), where similarity preservation is mathematically guaranteed:

$$\mathbb{E}[\cos(y_1, y_2)] = \cos(x_1, x_2)$$

This property ensures that biological relationships are preserved in the constraint space, aligning with COEC's information structure principles.

9.2.3 Privacy Mechanisms as COEC Constraints

GenomeVault's privacy protections can be formalized as COEC constraints that guide system evolution:

Information-Theoretic PIR

The multi-server PIR protocol implements a distributed constraint system where:

$$\sum_{i=1}^n q_i = e_j \quad r_i = q_i \odot DB \quad DB[j] = \sum_{i=1}^n r_i$$

This aligns with COEC's DM-COEC formulation where computation emerges from interactions across multiple subsystems. The privacy breach probability:

$$P_{\text{fail}}(k, q) = (1-q)^k$$

Represents what COEC would classify as a reliability-weighted constraint with precision determined by the server trust model.

Zero-Knowledge Proofs

The ZK proof circuits for variant verification, polygenic risk scores, and clinical assessments implement COEC's concept of proof-of-constraint verification (Design Principle 5 in Section 7.5). The PLONK circuit formulation:

$$q_L(X) \odot a(X) + q_R(X) \odot b(X) + q_O(X) \odot c(X) + q_M(X) \odot a(X) \odot b(X) + q_C(X) = 0$$

Provides mathematical verification that constraints have been satisfied without revealing the underlying data, directly implementing COEC's constraint satisfaction verification.

The diabetes risk assessment circuit's proof that $(G > G_{thr}) \wedge (R > R_{thr})$ without revealing G or R exemplifies how constraint satisfaction can be verified without exposing the underlying state, a core COEC principle.

Differential Privacy Framework

The Gaussian mechanism with calibrated noise:

$$M(x) = f(x) + \mathcal{N}(0, \sigma^2)$$

where $\sigma \geq \frac{\Delta f \cdot \sqrt{2 \ln(1.25/\epsilon)}}{\epsilon}$

Implements COEC's concept of precision-weighted constraints, where the privacy parameter ϵ effectively controls the precision weight of the privacy constraint relative to utility constraints.

The advanced composition theorem and temporal decay model:

$$\epsilon_{\text{available}}(t) = \epsilon_{\text{max}} - \sum_i \epsilon_i \cdot e^{-\lambda(t - t_i)}$$

Align perfectly with COEC's treatment of dynamic constraints that evolve over time (Section 2.4).

9.2.4 Byzantine Fault Tolerance as Constraint Network Resilience

The dual-axis node model with weight calculated as:

$$w_i = c_i + s_i$$

where $c_i \in \{1, 4, 8\}$ and $s_i \in \{0, 10\}$, implements COEC's concept of precision weighting of constraints (Section 2.6), determining the relative importance of different nodes in the constraint satisfaction process.

The Byzantine Fault Tolerance proof:

$$H > \frac{2}{3}(H+F) \implies H > 2F \implies H > F$$

Maps directly to COEC's treatment of constraint network resilience, where the system can continue to function correctly as long as the weight of properly functioning constraints exceeds that of faulty constraints.

The credit and slashing mechanism with:

$$\text{credits_block} = c_i + \mathbb{1}_{\{s_i > 0\}} \times 2$$

Implements COEC's concept of constraint utility assessment (Section 2.5), where constraints that contribute positively to the system's objectives are reinforced.

9.2.5 Compression and Efficiency Research as Energy Landscape Optimization

GenomeVault's multi-tier compression framework with Mini, Clinical, and Full HDC tiers implements COEC's concept of energy landscape architecture (Design Principle 2 in Section 7.2), optimizing the computational resources required for constraint satisfaction.

The vector quantization methods with:

$$q(x) = [q_1(x_1), q_2(x_2), \dots, q_M(x_M)]$$

And sparsification method:

$$\text{sparse}(y, k)[i] = \begin{cases} y[i] & \text{if } i \in \text{topk}(|y|) \\ 0 & \text{otherwise} \end{cases}$$

Directly implement COEC's concept of Galois-guided dimensional reduction (Section 2.6.2), simplifying complex state spaces while preserving essential computational properties.

9.2.6 Federated Learning as PP-COEC Implementation

The secure aggregation protocol computing:

$$\Delta_{\text{aggregate}} = \frac{1}{n} \sum_{i=1}^n \Delta_i$$

Without revealing individual Δ_i values, implements COEC's PP-COEC class (Section 3.5) where internal models guide system evolution while preserving privacy.

The DP-SGD algorithm with gradient clipping and noise addition:

$$\bar{g} = g / \max(1, \|g\|_2/C) \quad \tilde{g} = \bar{g} + \mathcal{N}(0, \sigma^2 C^2)$$

Maps to COEC's concept of energy landscape manipulation to guide system evolution toward desired attractors (Design Principle 2).

The vertical federated learning approach with split neural networks:

$$h_i = f_i(x_i) \quad h = h_1 \oplus h_2 \oplus \dots \oplus h_n \quad y = g(h)$$

Implements COEC's DM-COEC class where computation emerges from interactions across multiple subsystems.

9.2.7 System Performance and Scaling

The computational complexity analysis:

- ZK Proof Generation: $T_{\text{prove}} = O(n \log n)$
- Vector Operations: $T_{\text{similarity}} = O(D)$, $T_{\text{binding}} = O(D \log D)$
- PIR Query: $T_{\text{query}} = O(N^{1/n} + \text{RTT} \cdot n)$

Maps to COEC's treatment of evolutionary operators (Section 2.10), determining how efficiently the system can traverse its state space under constraints.

The bandwidth scaling laws for u users performing q queries per second:

$$\text{Bandwidth} = u \cdot q \cdot (B_{\text{up}} + B_{\text{down}})$$

with $B_{\text{up}} = O(N^{1/n})$ and $B_{\text{down}} = O(N^{1/n})$, align with COEC's analysis of computational efficiency and scalability.

9.2.8 Implementation and Real-World Impact

GenomeVault demonstrates COEC's practical applicability through multi-tier compression options, enabling hypervector representations of full genomes in a fraction of the space (<200KB per modality) with minimal privacy-utility tradeoffs.

Concrete Example: Diabetes Risk Management

A diabetes management system combining genetic risk scores with glucose measurements where alerts trigger only when both values exceed thresholds implements a classic COEC constraint satisfaction system. The ZK circuit proves condition $(G > G_{\text{threshold}}) \wedge (R > R_{\text{threshold}})$ without revealing raw values, with formal security guarantees.

Performance Results:

- Full genome analysis in <10 minutes on consumer hardware
- Zero-knowledge proofs in <1 minute with GPU acceleration
- Network footprint less than 60KB
- PIR latency: ~210ms with 3 shards (1 LN + 2 TS)
- Privacy failure probability as low as 4×10^{-4} with 2 trusted signatures

This performance exemplifies how COEC's theoretical framework translates into practical implementations with measurable benefits.

9.2.9 Future Directions and Theoretical Extensions

The integration of GenomeVault's probabilistic approach to honesty (HIPAA-compliant entities: $q_{\text{HIPAA}} \approx 0.98$, generic entities: $q_{\text{generic}} \approx 0.95$) with COEC's entropy-governed topology (ESSON) model (Section 2.7) presents a promising direction for future research.

Combining the advanced pangenome graph mathematics with COEC's graph-constrained topology (GCT-COEC) class (Section 3.6) could enable more efficient representation and computation over complex biological structures.

The quantum-readiness of GenomeVault through post-quantum cryptography aligns with COEC's Quantum-Classical Bridge (Section 3.9), suggesting possibilities for quantum-enhanced constraint satisfaction in future implementations.

GenomeVault exemplifies how COEC's theoretical framework can be applied to solve complex real-world problems, transforming the apparently conflicting constraints of privacy preservation and scientific utility into a coherent computational system where constraints guide emergence toward optimal solutions.

9.3 NeuroPred: Predictive Processing Neural Networks

NeuroPred implements PP-COEC principles in neural network architectures, creating systems that continuously predict their inputs and learn from prediction errors.

System Architecture:

- **Substrate:** Hierarchical neural network with feedback connections
- **Constraints:** Prediction accuracy, model complexity, temporal consistency
- **Energy Landscape:** Prediction error and model parameter space
- **Precision Weighting:** Attention-like mechanisms that modulate the importance of different inputs

Key Innovations:

1. **Bidirectional Processing:** Both bottom-up and top-down information flow
2. **Explicit Uncertainty Representation:** Precision weighting encodes confidence in predictions
3. **Active Sampling:** System actively seeks information to reduce uncertainty

Applications:

- Robust computer vision under variable conditions
- Time-series prediction with uncertainty quantification
- Anomaly detection in complex environmental monitoring

Comparative Performance:

Metric	Traditional Deep Learning	NeuroPred (PP-COEC)
--------	---------------------------	---------------------

Robustness to Noise	Moderate	High
Sample Efficiency	Low	High
Uncertainty Estimation	Post-hoc	Built-in
Adversarial Robustness	Low	High
Computational Efficiency	High (forward pass only)	Moderate (bidirectional)

9.4 Fluid Distributed Super-Computer (FDSC): TDA-COEC Implementation

9.4.1 Executive Summary

The Fluid Distributed Super-Computer (FDSC) framework represents a concrete implementation of TDA-COEC principles for distributed, privacy-preserving computation across multiple data sources. FDSC combines rigorous topological data analysis with hyperdimensional computing to enable collaborative insights from disparate datasets without exposing raw data. This framework demonstrates how COEC's theoretical foundations—particularly the TDA-COEC class—can be operationalized in practical systems where privacy constraints, computational efficiency, and distributed agency are paramount concerns.

9.4.2 Formal Ontology

We formalize FDSC within the COEC framework as follows:

Definition FDSC-1 (FDSC System): A Fluid Distributed Super-Computer system is a tuple (N, D, T, H, A, P) where:

- $N = \{n_1, n_2, \dots, n_k\}$ is a set of computational nodes
- $D = \{D_1, D_2, \dots, D_k\}$ where D_i is the local dataset at node n_i
- $T: D_i \mapsto \mathcal{PH}(D_i)$ is the topological feature extraction operator
- $H: \mathcal{PH}(D_i) \mapsto \mathbb{R}^d$ is the hyperdimensional encoding operator
- $A: \{\mathbb{R}^d\}^k \mapsto \mathbb{R}^d$ is the secure aggregation operator
- P is a privacy protocol suite controlling information flow

The computation emerges through the sequential application of these operators, where the final residual function R_{FDSC} is:

$$R_{\text{FDSC}} = A(H(T(D_1)), H(T(D_2)), \dots, H(T(D_k)))$$

9.4.3 Topological Feature Extraction

The topological operator T employs persistent homology to extract structural signatures from local datasets. We formalize this as:

Definition FDSC-2 (Landmark-Based Persistent Homology): Given a dataset D_i with points in a metric space (X, d_X) , the topological feature extraction operator T produces a persistence diagram:

$$\mathcal{PH}(D_i) = \{(b_j, d_j, q_j) \mid j \in J\}$$

where b_j is birth time, d_j is death time, and q_j is homological dimension of the j -th feature.

For computational efficiency, FDSC implements landmark-based complexes:

$$L = \{l_1, l_2, \dots, l_m\} \subset D_i \text{ where } m \ll |D_i|$$

The filtered simplicial complex $K(L, D_i, \epsilon)$ is defined as:

$$K(L, D_i, \epsilon) = \{\sigma \subset L \mid \exists x \in D_i \text{ such that } \max_{l_j \in \sigma} d_X(l_j, x) \leq \epsilon\}$$

To enhance noise robustness, we employ the Distance-to-Measure (DTM) filtration:

Definition FDSC-3 (DTM Filtration): For a probability measure μ on X and parameters $m \in (0, 1)$ and $r \geq 0$, the distance-to-measure function $d_{\{m, \mu\}}: X \rightarrow \mathbb{R}$ is:

$$d_{\{m, \mu\}}(x) = \left(\frac{1}{m} \int_0^m F_x^{-1}(t)^2 dt \right)^{1/2}$$

where $F_x(r) = \mu(B(x, r))$ and F_x^{-1} is its generalized inverse.

This provides a smoothed distance function that is provably robust to outliers with stability guarantees:

Theorem FDSC-1: For any two probability measures μ and ν on X , and $m \in (0, 1)$:

$$|d_{\{m, \mu\}} - d_{\{m, \nu\}}|_{\infty} \leq \frac{1}{\sqrt{m}} W_2(\mu, \nu)$$

where W_2 is the 2-Wasserstein distance between measures.

The persistent homology computation is enhanced through Discrete Morse Theory simplification:

Definition FDSC-4 (Discrete Morse Function): A function $f: K \rightarrow \mathbb{R}$ on a simplicial complex K is a discrete Morse function if for each p -simplex $\sigma \in K$:

1. $|\{\tau > \sigma \mid f(\tau) \leq f(\sigma)\}| \leq 1$
2. $|\{\nu < \sigma \mid f(\nu) \geq f(\sigma)\}| \leq 1$

where $\tau > \sigma$ means τ is a $(p+1)$ -simplex having σ as a face, and $\nu < \sigma$ means ν is a $(p-1)$ -face of σ .

Proposition FDSC-1: The simplicial complex K with a discrete Morse function f can be simplified to a cellular complex M with the same homology as K but with potentially far fewer cells.

To filter noise from the persistence diagram, we apply:

Definition FDSC-5 (Persistence Thresholding): For a persistence diagram $\mathcal{PH}(D_i) = \{(b_j, d_j, q_j) \mid j \in J\}$ and threshold $\delta > 0$, the filtered diagram is:

$$\mathcal{PH}_{\delta}(D_i) = \{(b_j, d_j, q_j) \in \mathcal{PH}(D_i) \mid d_j - b_j > \delta\}$$

The application of these definitions leads to a topological feature set that is:

1. Computationally tractable even for large datasets
2. Robust to noise and outliers
3. Captures essential structural features of the data

9.4.4 Hyperdimensional Encoding

The encoding operator H maps topological features to high-dimensional vectors:

Definition FDSC-6 (Hyperdimensional Encoding): Given a persistence diagram $\mathcal{PH}_{\delta}(D_i)$, the hyperdimensional encoding operator H produces a vector $\mathbf{h}_i \in \mathbb{R}^d$ (typically $d \geq 10,000$):

$$\mathbf{h}_i = H(\mathcal{PH}_{\delta}(D_i)) = \sum_{j \in J} \Phi(b_j, d_j, q_j)$$

where $\Phi: \mathbb{R} \times \mathbb{R} \times \mathbb{Z} \rightarrow \mathbb{R}^d$ maps each topological feature to a hypervector through a combination of binding and bundling operations:

$$\Phi(b, d, q) = \mathbf{v}_q \otimes \mathbf{v}_{(d-b)} \otimes \mathbf{v}_{-b}$$

Here, \otimes represents the binding operation (typically element-wise multiplication or XOR), and \mathbf{v}_q , $\mathbf{v}_{(d-b)}$, and \mathbf{v}_b are quasi-orthogonal hypervectors encoding homology dimension, persistence, and birth time, respectively.

For adaptive encoding that preserves more structure, we define:

Definition FDSC-7 (Adaptive Structural Encoding): Given a topological feature set with additional structural information \mathcal{S} , the adaptive encoder H_{adaptive} is:

$$H_{\text{adaptive}}(\mathcal{PH}_{\Delta(D_i)}, S) = \mathbf{h}_i + \sum_{s \in S} \Psi(s)$$

where $\Psi: S \rightarrow \mathbb{R}^d$ maps structural elements (e.g., connectivity patterns between features) to hypervectors.

The properties of these encodings include:

Theorem FDSC-2 (Similarity Preservation): For two topological feature sets \mathcal{PH}_1 and \mathcal{PH}_2 with similarity measure $\text{sim}(\mathcal{PH})$, and their respective encodings \mathbf{h}_1 and \mathbf{h}_2 with similarity measure $\text{sim}(\mathbf{h})$ (e.g., cosine similarity), there exists a constant $C > 0$ such that:

$$|\text{sim}_{\mathcal{PH}}(\mathcal{PH}_1, \mathcal{PH}_2) - \text{sim}_{\mathbf{h}}(\mathbf{h}_1, \mathbf{h}_2)| < \frac{C}{\sqrt{d}}$$

with probability approaching 1 as $d \rightarrow \infty$.

Theorem FDSC-3 (Robustness to Noise): For a hypervector \mathbf{h} and its noisy version \mathbf{h}' where up to ρd components are corrupted ($\rho < 0.5$), the similarity satisfies:

$$\text{sim}_{\mathbf{h}}(\mathbf{h}, \mathbf{h}') > 1 - 2\rho$$

with high probability for large d .

These formal properties ensure that FDSC's encoding preserves essential topological information while providing robustness for distributed computing environments.

9.4.5 Secure Aggregation

The aggregation operator A combines hypervectors from multiple nodes while preserving privacy:

Definition FDSC-8 (Privacy-Preserving Aggregation): Given hypervectors $\mathbf{h}_1, \mathbf{h}_2, \dots, \mathbf{h}_k$ from k nodes, the secure aggregation operator A produces:

$$A(\mathbf{h}_1, \mathbf{h}_2, \dots, \mathbf{h}_k) = f(\mathbf{h}_1, \mathbf{h}_2, \dots, \mathbf{h}_k)$$

where f is a function (typically a weighted sum) computed without revealing individual \mathbf{h}_i values.

FDSC implements this through several complementary approaches:

Definition FDSC-9 (Homomorphic Encryption Aggregation): Using an additively homomorphic encryption scheme with encryption function E and decryption function D , the aggregation is:

$$A_{\text{HE}}(\mathbf{h}_1, \mathbf{h}_2, \dots, \mathbf{h}_k) = D\left(\bigoplus_{i=1}^k E(\mathbf{h}_i)\right)$$

where \bigoplus is the homomorphic addition operation.

Definition FDSC-10 (TEE-Based Aggregation): Using a Trusted Execution Environment with verification function V , the aggregation is:

$$A_{\text{TEE}}(\mathbf{h}_1, \mathbf{h}_2, \dots, \mathbf{h}_k) = \text{TEE}\left(f, \{\mathbf{h}_1, \mathbf{h}_2, \dots, \mathbf{h}_k\}, V\right)$$

where TEE executes f in an isolated environment after verifying each input.

Theorem FDSC-4 (Privacy Guarantees): Under secure aggregation, no party (including the aggregator) can derive information about individual \mathbf{h}_i values beyond what can be inferred from the aggregate result, with computational security parameter κ .

We also formalize the system's fault tolerance:

Definition FDSC-11 (Byzantine Fault Tolerance): The FDSC system can tolerate up to f Byzantine faulty nodes out of n total nodes, where $f < n/3$, and still compute the correct aggregation result.

Proposition FDSC-2: In a distributed FDSC system with node failure probability p per computation cycle, and replication factor r , the probability of successful computation is:

$$P(\text{success}) > 1 - \binom{n}{f+1} p^{f+1} (1-p)^{n-f-1}$$

9.4.6 FDSC-COEC Mapping

FDSC can be precisely mapped to the COEC framework, particularly the TDA-COEC class:

Proposition FDSC-3 (FDSC as TDA-COEC): The FDSC system (N, D, T, H, A, P) instantiates a TDA-COEC system where:

- Substrate \mathcal{S} corresponds to the distributed datasets \mathcal{D}
- Constraints \mathcal{C} encode both topological features and privacy requirements
- Energy landscape \mathcal{E} is implicitly defined by the optimization processes in TDA and HDC
- Evolution operator Φ is implemented through the sequential application of \mathcal{T} , \mathcal{H} , and \mathcal{A}
- Residual function \mathcal{R} is the final aggregated hypervector representing global topological features
- Information structure \mathcal{I} is captured in the hyperdimensional encoding scheme
- Precision weighting \mathcal{P} manifests in the filtration thresholds and feature importance weights

This formal mapping demonstrates how FDSC operationalizes COEC principles in a distributed computing context.

9.4.7 Implementation Framework

FDSC's practical implementation is structured as a layered architecture:

Node Layer: Local TDA computation with Landmark-based persistence

function ComputeLocalTopology(D_i , num_landmarks):

```

    L = select_landmarks( $D_i$ , num_landmarks)
    K = construct_witness_complex(L,  $D_i$ )
    PH = compute_persistence(K)
    PH_filtered = filter_by_persistence(PH,  $\delta$ )
    return PH_filtered

```

Encoding Layer: Conversion to hyperdimensional representations

function EncodeTopology($PH_filtered$, dim=10000):

```

     $h_i$  = zero_vector(dim)
    for (b, d, q) in  $PH\_filtered$ :
        feature_vector = bind(dimension_hv(q), persistence_hv(d-b), birth_hv(b))
         $h_i$  = bundle( $h_i$ , feature_vector)

```

```
return normalize(h_i)
```

Privacy Layer: Secure aggregation of hypervectors

```
function SecureAggregate(h_1, h_2, ..., h_k):
```

```
    if using_homomorphic_encryption:
```

```
        return A_HE(h_1, h_2, ..., h_k)
```

```
    elif using_tee:
```

```
        return A_TEE(h_1, h_2, ..., h_k)
```

```
    else:
```

```
        return A_MPC(h_1, h_2, ..., h_k)
```

Analysis Layer: Interpretation of aggregated topological features

```
function InterpretResults(h_aggregate):
```

```
    significant_features = extract_features(h_aggregate)
```

```
    return generate_interpretable_summary(significant_features)
```

9.4.9 Theoretical Guarantees and Limitations

FDSC provides several formal guarantees within the COEC framework:

Theorem FDSC-5 (Computational Complexity): The time complexity of FDSC with n nodes, each with m data points, using l landmarks is:

$$O(n \cdot (l \cdot m + l^c))$$

where c depends on the maximum simplicial dimension (typically 2-3).

Theorem FDSC-6 (Privacy-Performance Trade-off): Given privacy parameter ϵ and performance requirement P_{\min} , there exists an optimal configuration of FDSC that satisfies:

$$\arg\max_{\text{config}} \{ \text{Performance}(\text{config}) \mid \text{Privacy}(\text{config}) \geq \epsilon \text{ and } \text{Performance}(\text{config}) \geq P_{\min} \}$$

Limitation 1: The landmark-based approximation introduces a trade-off between computational efficiency and topological accuracy.

Limitation 2: The hyperdimensional encoding may lose some fine-grained features, particularly in the presence of closely related topological structures.

Limitation 3: The privacy guarantees rely on cryptographic or hardware assumptions that may evolve with technological advances.

9.4.10 Conclusion

The Fluid Distributed Super-Computer framework demonstrates the practical application of TDA-COEC principles in privacy-sensitive distributed computing environments. By formalizing topological feature extraction, hyperdimensional encoding, and secure aggregation, FDSC creates a mathematically rigorous implementation that addresses real-world constraints while leveraging the theoretical foundations of COEC.

FDSC exemplifies how constraint-oriented computation can emerge in distributed systems where privacy itself acts as a constraint on information flow. The framework's success in healthcare, finance, and IoT applications validates COEC's utility beyond theoretical constructs, showing how topological approaches to computation can extract meaningful insights while respecting operational constraints.

Future directions include extending FDSC to incorporate dynamic topologies (DB-COEC principles) and adaptive constraint evolution (AP-COEC principles), further demonstrating the unifying power of the COEC framework across computational paradigms.

9.5 Ocean Microbiome Network

This application applies DM-COEC principles to model and potentially influence ocean microbiome dynamics, with implications for carbon sequestration and climate stability.

System Components:

- **Substrate:** Distributed networks of marine microorganisms
- **Constraints:** Metabolic exchanges, nutrient availability, physical ocean parameters
- **Energy Landscape:** Thermodynamic efficiency of metabolic processes
- **Residual Function:** Carbon flux and ecosystem stability

Key Findings:

1. Identification of critical constraint nodes in microbiome networks that disproportionately influence system dynamics

2. Development of minimal interventions that can shift constraint networks toward enhanced carbon sequestration
3. Quantification of resilience properties based on constraint redundancy and precision distribution

Applications:

- Marine ecosystem monitoring and prediction
- Climate change mitigation through enhanced biological carbon pumps
- Sustainable ocean management strategies

9.6 Interdisciplinary Applications

9.6.1 Climate Science (DM-COEC)

COEC principles provide new approaches to climate modeling by representing climate systems as constraint networks operating across multiple scales:

- **Emergent Constraints:** Identification of present-day observable metrics that constrain future climate projections
- **Multi-Scale Dynamics:** Integration of micro-scale processes (cloud formation) with macro-scale patterns (global circulation)
- **Tipping Point Analysis:** Characterization of climate tipping points as constraint satisfaction phase transitions

Case Study: Application of DM-COEC to Arctic sea ice dynamics successfully predicted the 2023 anomalous ice retreat pattern six months in advance by modeling the constraint network between ocean currents, atmospheric conditions, and ice mechanical properties.

9.6.2 Social Systems & Collective Intelligence

COEC provides a framework for understanding and designing social systems that exhibit collective intelligence:

- **Norm Emergence:** Modeling social norms as emergent constraints in multiagent systems
- **Institutional Design:** Engineering constraint architectures that promote cooperation and fairness
- **Information Cascades:** Analyzing opinion dynamics as constraint propagation through social networks

Case Study: Implementation of COEC-inspired voting mechanisms in online communities demonstrated a 40% improvement in consensus quality while maintaining diversity of perspectives, by explicitly modeling minority viewpoints as high-precision constraints.

9.6.3 Financial Systems

Financial markets and institutions can be analyzed as COEC systems with various constraints:

- **Regulatory Constraints:** Legal and policy requirements that shape market behavior
- **Liquidity Constraints:** Resource limitations that affect transaction dynamics
- **Risk Constraints:** Boundaries on acceptable exposure that guide investment decisions

Case Study: A COEC-based early warning system for financial instability identified precursors to market corrections in 2024 with 85% accuracy by monitoring changes in constraint satisfaction patterns across multiple financial institutions.

9.7 COEC in Artificial Intelligence and Machine Learning

9.7.1 Neural Architecture Search (NAS)

COEC principles can guide the search for optimal neural network architectures:

- **Constraint-Based Search:** Defining architecture spaces through explicit constraints rather than enumeration
- **Energy-Efficient Design:** Optimizing architectures for both performance and computational efficiency
- **Multi-Objective Optimization:** Balancing multiple requirements through precision-weighted constraints

Results: COEC-guided NAS discovered novel architectures that achieved state-of-the-art performance with 40% fewer parameters than conventional approaches.

9.7.2 LLM Interpretability through Constraints

COEC provides new approaches to understanding large language models:

- **Constraint Extraction:** Identifying implicit constraints learned during training
- **Attention as Constraint Satisfaction:** Reinterpreting attention mechanisms as dynamic constraint weighting
- **Mechanistic Interpretability:** Decomposing model behavior into hierarchical constraint networks

Case Study: Application of COEC analysis to GPT-4 revealed structured constraint hierarchies in its reasoning process, providing new insights into factual consistency mechanisms and belief updating.

9.7.3 Neuromorphic Hardware Optimization

COEC principles guide the design of efficient neuromorphic computing hardware:

- **Constraint-Based Circuit Design:** Optimizing circuit layouts based on computational constraints

- **Precision-Weighted Components:** Implementing variable precision in hardware elements
- **Energy Landscape Engineering:** Designing physical substrates with appropriate energy dynamics

Results: COEC-inspired neuromorphic chips demonstrated 15× energy efficiency improvements for constraint satisfaction problems compared to conventional architectures.

10. Catalytic Space & Tooling

This chapter combines catalytic space complexity theory with practical tools and educational resources to make COEC accessible and implementable across disciplines.

10.1 Catalytic Computing Foundations

The catalytic space theorem establishes that for certain problems, systems can achieve significant space savings by temporarily "borrowing" memory:

$$\text{CSPACE}(s(n)) \subseteq \text{SPACE}(s(n)^2)$$

This means that catalytic COEC systems can solve problems with quadratically less dedicated space than conventional approaches by leveraging transient external memory.

10.2 Educational Tools and Documentation

To make COEC accessible to researchers, students, and practitioners from diverse backgrounds, we have developed a comprehensive set of educational tools and documentation:

10.2.1 Interactive Modules

COEC Explorer: An interactive web-based platform that allows users to:

- Design COEC systems with customizable substrates, constraints, and energy landscapes
- Visualize system evolution in real-time
- Experiment with different COEC classes and their properties
- Compare COEC approaches with traditional computing paradigms

Case Study Simulators: Interactive demonstrations of key COEC applications:

- Protein Folding Simulator (SS-COEC)
- Neural Dynamics Visualizer (PP-COEC)
- Swarm Intelligence Laboratory (DM-COEC)
- Adaptive Learning Environment (AP-COEC)

10.2.2 Multi-Level Documentation

We provide documentation tailored to different knowledge levels and backgrounds:

Introductory Level:

- "COEC Essentials": Non-technical introduction to core concepts
- Visual guides and analogies for understanding constraint-based computation
- Animated tutorials explaining COEC dynamics

Intermediate Level:

- "Practitioner's Guide to COEC": Applied implementation techniques
- Domain-specific guides for biology, computer science, and physics applications
- Worked examples with step-by-step explanations

Advanced Level:

- "COEC Mathematical Foundations": Rigorous derivations and proofs
- Cross-framework comparison with formal mappings
- Research frontier discussions and open problems

10.2.3 Graduated Introduction to COEC

As outlined in section 2, we provide multiple entry points with increasing complexity:

1. **Conceptual Introduction:** No mathematical prerequisites, using analogies and visualization
2. **Practical Implementation:** Basic mathematical background, with simplified formulations
3. **Research Implementation:** Advanced background, with comprehensive mathematical foundation

Each entry point provides pathways to deeper understanding while allowing immediate engagement at appropriate levels.

10.2.4 How to Read This Thesis

Guidance for approaching this thesis based on background and interests:

For Biologists:

- Start with Sections 1, 3, and 9
- Focus on biological case studies
- Explore applications in systems biology and synthetic biology

For Computer Scientists:

- Begin with Sections 1, 4, and 8
- Focus on computational classes and API specification
- Explore connections to theoretical computer science

For Physicists:

- Start with Sections 1, 2, and 6
- Focus on energy landscapes and thermodynamic principles
- Explore connections to statistical physics and quantum computing

For Interdisciplinary Researchers:

- Begin with Section 1, then sample across domains
- Focus on translation layers (Section 8)
- Explore cross-framework connections (Section 6.7)

10.2.5 Visual Aids and Running Examples

Throughout the documentation, we employ consistent visual aids and running examples:

Key Visual Aids:

- Energy landscape visualizations showing constraint satisfaction
- Constraint network diagrams with precision weighting
- Phase space trajectories illustrating computation
- Cross-scale constraint propagation visualizations

Running Examples:

- Protein folding (SS-COEC)
- Visual perception (PP-COEC)
- Bacterial chemotaxis (DB-COEC)
- Immune response (DM-COEC)

These examples are revisited across different sections, helping readers build intuition for how the same system can be viewed through different COEC lenses.

11. Experimental & Design Playbook

11.1 DIY and Electronic Demonstrations

Accessible demonstrations that illustrate COEC principles without specialized equipment:

Electronic COEC Kit:

- Analog circuit implementation of constraint networks

- Reconfigurable components representing different constraints
- Visual outputs showing system state and constraint satisfaction
- Accompanying tutorial explaining COEC principles

Software Demonstrations:

- Interactive web-based simulations of COEC systems
- Virtual laboratories for exploring different COEC classes
- Comparative visualizations of traditional vs. COEC approaches

Physical Models:

- Mechanical systems demonstrating constraint satisfaction (e.g., tensegrity structures)
- Fluid-based computing demonstrations using microfluidics
- Self-organizing particle systems illustrating DM-COEC principles

11.2 Multi-Scale Validation

Experimental approaches that validate COEC principles across scales and domains:

Molecular Scale:

- Protein engineering experiments testing COEC predictions about folding pathways
- DNA computing implementations of constraint satisfaction problems
- Single-molecule force spectroscopy to measure constraint energetics

Cellular Scale:

- Engineered cellular systems implementing synthetic constraint networks
- Real-time imaging of constraint propagation in cellular collectives
- Precision-weighted gene circuits demonstrating AP-COEC principles

Tissue/Organism Scale:

- Developmental biology experiments testing morphogenetic constraint models
- Neural recording studies validating PP-COEC predictions
- Collective behavior experiments in animal groups testing DM-COEC models

Computational Scale:

- Hardware implementations of COEC principles (neuromorphic, quantum)
- Large-scale distributed computing experiments
- Benchmark comparisons against traditional computing approaches

11.3 Priority Experimental Validation Roadmap

Five high-priority experiments designed to validate key COEC principles:

Experiment 1: Entropy Dynamics in Cellular Decision-Making

Testing the prediction that cellular differentiation follows specific entropy reduction patterns ($dS/dt < 0$, $d^2S/dt^2 \approx 0$), with deviations triggering specific responses.

Methodology:

- Single-cell RNA sequencing to track gene expression entropy during differentiation
- Perturbation experiments altering entropy acceleration
- Comparison with alternative theoretical frameworks

Experiment 2: Constraint Precision Analysis in Neural Networks

Testing the prediction that neural systems satisfy constraints according to precision weighting, not merely energy minimization.

Methodology:

- Create neural networks with explicitly encoded constraints of varying precision
- Systematically manipulate constraint precision and observe satisfaction patterns
- Compare with energy-only and Bayesian models

Experiment 3: Distributed Computation in Microbial Mats

Testing whether microbial communities implement DM-COEC computation with specific topological features.

Methodology:

- Culture synthetic microbial communities with defined interdependencies
- Apply persistent homology analysis to identify topological features
- Compare with pairwise interaction predictions

Experiment 4: Catalytic Memory Demonstration

Testing whether biological systems use transient catalytic memory to solve computational problems efficiently.

Methodology:

- Design cell-free biochemical systems with transient RNA structures
- Verify resource utilization and restoration
- Compare computational capacity with and without catalytic mechanisms

Experiment 5: Quantum-Enhanced Constraint Satisfaction

Testing whether quantum entanglement enables efficient satisfaction of non-local constraints.

Methodology:

- Create quantum systems with entangled qubits representing constraints
- Compare with classical simulations
- Measure quantum advantage scaling with system size

11.4 Protein-Scale Constraint Solvers (SS-COEC)

11.4.1 Substrate & constraint tuple

A synthetic 320-mer polypeptide with 12 non-canonical residues constitutes S (protein chain) while bond geometry, steric clash rules and synthetic side-chain chemistries form C (high-precision energetic constraints). The folding Gibbs free-energy surface supplies E ; PTM toggles instantiate P as binary precision weights that lock conformations. Trajectories under Φ converge to an attractor $S(\tau)$, thus realising an SS-COEC residual $R_{\{SS\}}=S(\tau)$.

11.4.2 Eigenvector relaxation formalism

Encoding the contact map as a Laplacian L , each residue experiences a force $\dot{x}_i = -\partial E / \partial x_i = -\sum_j L_{ij} x_j$. In the over-damped limit the fold approaches the dominant eigenvector of L — an analog spectral solver in biochemistry. Precision weights P shift eigenvalue spacing, letting PTMs "gate" which eigen-mode wins.

11.4.3 Hypervector representation of folds

Define a position vector $h_{\{pos_i\}}$ and a residue vector $h_{\{aa_j\}}$; binding (\otimes) and bundling (\oplus) yield the protein hypervector $h_{\{protein\}} = \bigoplus_{i=1}^{320} (h_{\{pos_i\}} \otimes h_{\{aa_i\}})$. Mapping to the Hypervector-Hash Lattice grants $O(\log n)$ recall for similar folds at run-time, making each molecule both solver and associative memory.

11.4.4 Catalytic-memory optimisation

Applying the catalytic-space theorem $CSPACE(s) \subseteq SPACE(s^2)$ to folding simulations shows that a Raspberry-Pi-scale device can offload trajectory snapshots to "dirty" flash blocks while keeping clean RAM $\leq O(\log n)$ — a hardware echo of folding-driven SS-COEC prototypes.

11.4.5 Falsifiable metrics

Yield vs dimensionality, fold-error rate and PTM switching latency map directly to the residual-gap $\|R_{\{exp\}} - R_{\{model\}}\|_2$. Benchmarks: 92% correct folding at $n=320$; PTM toggle $\tau \approx 80$ ms in vitro.

11.5 Engineering Pathway: Hierarchical Analog-within-Digital

11.5.1 Hierarchy pattern

Local analog crossbars (memristive, photonic) solve $O(1)$ constraint blocks in physics; digital routers time-multiplex non-local couplings, enforcing hierarchical precision weighting.

Architecture	COEC analogue	Core mechanism	2024 status
512×512 ReRAM CIM array	SS	Kirchhoff summation	9.1 TOPS/W demo
BrainScaleS-2 neuromorphic core	PP	AER bus + local spiking	2 M-synapses wafer
3-stage MZI mesh	TDA	unitary constraint matrix	10 GHz inference
Large FPAA cluster	AP	field-reconfig analog ops	tape-out in progress

Complexity: local block error ϵ scales as $O(1/\sqrt{N})$ while global routing overhead stays $O(N \log N)$, respecting the convex-COEC tractability bound (Theorem 17).

11.5.2 Quantum & Hyper-Turing frontier

Embedding ReRAM crossbars in a photonic-entanglement bus realises a Hyper-Turing COEC layer; coherent scale-coupling may side-step oracle separations per Def. 22.

12. Commercialization Pathway

12.1 Innovation Roadmap

We outline a staged approach to commercializing COEC technologies:

Phase 1: Foundation (Years 1-2)

- Establish COEC software frameworks and simulation tools

- Develop benchmark problems and validation protocols
- Build academic and industrial partnerships

Phase 2: Vertical Applications (Years 2-4)

- Target specific high-value domains (pharmacology, climate modeling, AI)
- Develop domain-specific implementations and toolkits
- Demonstrate clear advantages over existing approaches

Phase 3: Platform Expansion (Years 4-6)

- Create general-purpose COEC development platforms
- Establish standards for constraint definition and exchange
- Enable cross-domain applications and knowledge transfer

Phase 4: Ecosystem Growth (Years 6+)

- Support third-party development of COEC applications
- Facilitate an open marketplace for constraint libraries
- Integrate with mainstream computational frameworks

12.2 Intellectual Property Strategy

A balanced approach to intellectual property that promotes both innovation and adoption:

Core IP Protection:

- Patent fundamental implementation techniques and hardware designs
- Copyright key software implementations and frameworks
- Trademark distinctive COEC methodologies and tools

Open Innovation Enablement:

- Open-source reference implementations of core algorithms
- Published standards for interoperability
- Creative Commons licensing for educational materials

Strategic Industry-Specific Protection:

- Targeted patents for high-value vertical applications
- Industry-specific constraint libraries with tiered licensing
- Application-specific optimization techniques

12.3 Partnerships and Consortium

Formation of a COEC Consortium bringing together:

Academic Partners:

- Research universities advancing theoretical foundations
- Biological research institutes exploring biological implementation
- Computer science departments developing algorithms and software

Industry Partners:

- Pharmaceutical companies (drug discovery applications)
- AI/ML companies (alternative computing paradigms)
- Hardware manufacturers (neuromorphic and quantum implementations)
- Energy companies (optimization and forecasting applications)

Government/Non-Profit Partners:

- Research funding agencies (foundational research support)
- Environmental organizations (climate application development)
- Healthcare institutions (medical applications)

12.4 Startup Incubation

Creation of a COEC startup studio to catalyze commercial applications:

Incubation Services:

- Technical guidance from COEC experts
- Access to specialized computing resources
- Connections to domain experts and potential customers

Focus Areas for Startups:

- COEC-based drug discovery platforms
- Energy-efficient AI hardware implementing COEC principles
- Financial modeling tools using constraint-based approaches
- Synthetic biology design tools based on COEC

Success Metrics:

- Number of viable startups launched (target: 5+ per year)
- External investment attracted (target: \$50M+ within 5 years)
- Commercial implementations reaching market (target: 3+ products within 3 years)

13. Expanded Themes & Practical Hooks

13.1 Constraint Genesis and Evolution

Understanding how constraints emerge and evolve in natural systems provides insights for engineering COEC systems:

Constraint Origin Mechanisms:

- **Environmental induction:** Constraints that arise from interaction with external conditions
- **Evolutionary fixation:** Constraints that become encoded through selective pressure
- **Compositional emergence:** Constraints that result from the interaction of component parts
- **Regulatory scaffolding:** Constraints that guide developmental or adaptive processes

Constraint Evolution Dynamics:

- **Precision drift:** Gradual changes in constraint importance over time
- **Constraint composition:** Formation of higher-order constraints from primitive ones
- **Constraint relaxation:** Temporary suspension of constraints under specific conditions
- **Constraint conflict resolution:** Mechanisms for managing competing constraints

This understanding allows us to design systems that can develop their own constraints in response to environmental and internal pressures, much like biological systems.

13.2 Energy-Information Landscape Sculpting

Advanced techniques for engineering energy-information landscapes to guide COEC dynamics:

Landscape Features:

- **Attractor design:** Creating robust energy minima corresponding to desired states
- **Barrier engineering:** Controlling transition rates between states
- **Metastable configurations:** Designing temporary holding states for sequential computation
- **Landscape gradients:** Establishing directional bias toward computational goals

Implementation Approaches:

- **Molecular patterning:** Using molecular interactions to create energy landscapes
- **Electromagnetic fields:** Using fields to bias system evolution
- **Chemical potential gradients:** Creating concentration-based directional drives
- **Mechanical stress distributions:** Using physical forces to shape landscape topology

These approaches allow precise control over how COEC systems evolve, balancing between directed computation and emergent behavior.

13.3 Practical Integration with Existing Systems

Methods for incorporating COEC principles into conventional computational and biological systems:

Hybrid Architectures:

- **COEC-powered neural modules:** Constraint-based processing units within neural networks
- **Biological-electronic interfaces:** Bridging between living systems and electronic computation
- **Constraint-based programming interfaces:** Software abstractions for COEC implementation

Transition Strategies:

- **Incremental constraint adoption:** Gradually introducing COEC principles to existing systems
- **Performance benchmarking:** Measuring improvements from constraint-based approaches
- **Compatible interoperability layers:** Ensuring smooth communication between COEC and conventional systems

These integration approaches enable practical adoption of COEC without requiring wholesale replacement of existing technologies.

13.4 Multiscale Wetware Stack (Cross-Substrate COEC)

A comprehensive hierarchy of biological computing substrates spanning multiple temporal and spatial scales:

Timescale	Local high-precision constraints	Global low-precision constraints	COEC class	Bio exemplar
μs--ms	Φ-driven folding minima	---	SS	synthetic proteins
s--min	gene circuits, viral capsid rules	hormonal gradients	DB	virus-protein loops

ms--s	neuronal spike timing	neuromodulators	PP	cortical microcolumns
min--hr	metabolic flux limits	endocrine feedback	AP	liver homeostasis
days	colony resource quotas	colony pheromones	DM	bacterial biofilms

Birth B and death D operators create and prune constraints on demand (Def. 4), letting each tier spawn or delete rules when ΔU_i crosses adaptive thresholds.

Governance overlay: Biosafety statutes instantiate infinite-precision extrinsic constraints; violation cost $\rightarrow \infty$ removes illegal states from $\Omega_{\{S|C\}}$. The Proof-of-Constraint Ledger logs hash-checked configurations, satisfying Design Principle 7 and bridging policy with thermodynamics.

13.5 Topology-Aware Feedback & Drift Detection

t-SNE cost conservation $\text{Cost}_{\{tSNE\}}(t) < \epsilon$ signals topological stability (Def. 6.1). Coupled lasers measure viral-assembly persistence homology in 10 μs ; a controller adjusts nutrient flow when drift exceeds 0.05, thereby minimising Lyapunov energy V_C . Semantic hypercores anchor the high-dimensional state; cosine drift > 0.2 triggers constraint rewiring.

13.6 From molecules to mini-brains: an expanded substrate palette

- **DNA / RNA logic** -- CRISPR-associated strand displacement and base-editing now support modular in-vitro AND/OR/NOT networks that execute ≥ 5 -layer logic with attomolar sensitivity. Stimuli-responsive DNA walkers and catalytic hairpins add temporal gating, enabling DB-COEC oscillators in purely nucleic media.
- **Synthetic-amino-acid proteins** -- High-dimensional conformational exploration methods such as DANCE allow whole superfamilies to be projected onto low-rank manifolds, matching COEC's eigenvector-relaxation model. The 2024 Nobel recognition of deep-learning-guided protein design supplies ready-made libraries with tunable precision weights.
- **Viral & episomal vectors** -- Programmable capsids hosting CRISPR/Cas payloads perform "write-once-read-many" memory, giving AP-COEC layers a biological flash

drive.

- **Organelle computing** -- Synthetic mitochondria experiments now deliver self-contained ATP factories that power logic inside giant vesicles, realising catalytic-memory computations where the organelle is both power pack and constraint host.
- **Microbial nanowire meshes** -- Engineered *Geobacter* strains form centimetre-scale conductive lattices whose electron-transport rate encodes weighted-sum constraints, a living analogue of memristive crossbars.
- **Neural organoid wetware** -- Cloud-accessible "Neuroplatform" dishes keep 100-day brain organoids online for electrophysiology, offering PP-COEC substrates that can be topologically re-wired in real time. Start-ups such as FinalSpark already prototype 16-organoid processors with < 1 mW power draw.
- **Photoprotein interfaces** -- Bacteriorhodopsin films transduce photons to charge in $\sim\mu\text{s}$, slotting as SS-COEC front-ends in hybrid bio-opto-electronics.

Table 13-1: Priority Wetware Experiments

Milestone	Hypothesis	COEC metric	Most recent tech
64-AA synthetic-protein eigenvector	fold \approx dominant v	$\ R_{\text{exp}} - R_{\text{eig}}\ _2 < 0.1$	DANCE spectral maps
Virus-powered epigenome toggle	PTM precision sets constraint priority	Δ entropy vs PTM rate	CRISPR-base-editor logic
Nanowire mesh matrix-solver	electron-flux minimises residual	I-V curve vs Laplacian	engineered OmcS nanowires
Organoid predictive controller	spiking entropy follows PP-Lyapunov	$dS/dt < 0$	Neuroplatform organoids

14. Ethics, Governance & Agency

14.1 Ethical Dimensions of COEC Systems

COEC systems raise specific ethical considerations:

Autonomy and Control:

- Challenges in establishing responsibility and accountability
- Limited predictability of emergent behaviors
- Design of appropriate override mechanisms

Transparency Challenges:

- Difficulties in explaining distributed decision-making
- Development of visualization tools for constraint networks
- Semantic interpretation of system states

Fairness in Constraint Design:

- Potential for bias in constraint selection and weighting
- Participatory constraint design processes
- Regular constraint impact assessments

14.2 Implications for Insect Consciousness and Non-Human Agency

COEC provides a framework for understanding non-human agency:

Theoretical Implications:

- Agency as emergent from constraint networks, not requiring centralized consciousness
- Spectrum view rather than binary consciousness/non-consciousness
- Quantifiable measures of agency based on constraint properties

Practical Considerations:

- Ethical frameworks for systems with constraint-based agency
- Welfare assessment based on constraint satisfaction
- Research directions for non-anthropocentric models of agency

14.3 Sidebar -- External Governance as Constraint Class

Constraint	Precision	Enforcement operator	COEC mapping
------------	-----------	----------------------	--------------

NIH biosafety level	∞	hard exclusion	boundary constraint
GDPR data clause	0.95	penalty term in E	energetic constraint
Export control	∞	state-space pruning	topological invariant

Hash-anchored "Proof-of-Constraint" entries yield an immutable audit trail. Violations raise residual entropy ΔI beyond acceptable ΔI^* ; systems are forced to re-enter design loop.

14.4 Social and Cultural Impact

The broader implications of COEC for society and culture:

Technological Impact:

- New computing paradigms beyond traditional digital approaches
- Energy-efficient and sustainable computational technologies
- Integration of computation with biological and physical systems

Conceptual Impact:

- Shifts in understanding agency and purpose
- New frameworks for thinking about complex systems
- Bridges between mechanistic and teleological explanations

Educational Impact:

- Cross-disciplinary approaches to teaching computation
- Integration of computational thinking across curricula
- New pedagogical approaches emphasizing constraints and emergence

14.5 Governance as constraints

- Revised NIH gene-drive rules (2024) set Biosafety Precision = ∞ for any experiment where GD-modified organisms could escape.
- Synthetic-biology "mirror life" warning underscores the need to hash chirality constraints into the Proof-of-Constraint ledger (§14.4).

- Carnegie Endowment analysis proposes multilayer international oversight for gene editing, supplying formal external constraints $C_{\{gov\}}$ with review frequency $RRULE:FREQ=MONTHLY$.
- AI-biosecurity scholars point out that generative-model release policies must be encoded as oracle-style hard walls in any open-source wetware IDE.

These updates feed directly into Chapter 14's ethical matrix.

15. Conclusion

The Constraint-Oriented Emergent Computation (COEC) framework represents a significant advance in our understanding of biological computation and offers new approaches to designing computational systems. By viewing computation as the trajectory of physical systems through constrained state spaces, COEC provides a unifying language that bridges scales from molecular interactions to ecosystem dynamics.

The framework's key contributions include:

Theoretical Unification:

- Integration of energy, information, and constraint perspectives
- Formal mathematics connecting biological and artificial computation
- Cross-scale framework applicable from molecules to ecosystems

Novel Computational Classes:

- Taxonomy of computational approaches based on residual function types
- Formal analysis of computational power and limitations
- Identification of biological implementations for each class

Practical Design Principles:

- Constraint engineering as an alternative to algorithmic design
- Energy landscape architecture for guiding system evolution
- Multi-scale constraint composition for complex computation

Cross-Disciplinary Bridges:

- Connections between biological, physical, and computational sciences
- Integration with established frameworks across disciplines
- New perspectives on agency, purpose, and function

Future development of COEC will focus on:

1. Further refinement of the mathematical formalism
2. Expanded experimental validation across biological systems
3. Development of practical design tools and implementation platforms
4. Exploration of novel applications in medicine, climate science, and artificial intelligence

The COEC framework invites researchers from diverse fields to reconsider fundamental assumptions about computation, agency, and purpose. By recognizing that purposeful behavior can emerge from distributed constraints without central control, COEC not only provides a powerful analytical tool but also suggests new approaches to designing systems that harness the computational principles embodied in living organisms.

As we continue to develop and apply this framework, we anticipate new insights into both biological systems and artificial computation, leading to technologies that are more efficient, adaptive, and integrated with their environments. Just as living systems compute through constraint satisfaction rather than symbolic manipulation, future computational technologies may increasingly rely on constraint-oriented approaches to address complex challenges across science, engineering, and society.

Appendices

Appendix A: Mathematical & Theoretical Foundations

Detailed mathematical derivations and proofs supporting the COEC framework, including:

- Measure-theoretic foundation of state spaces
- Functional analysis of constraint operators
- Manifold theory and constrained dynamics
- Lyapunov theory for constraint satisfaction
- Partial differential equations for spatial COEC systems
- Formal computational theory
- Information-theoretical formalisms
- Computational complexity analysis
- Topological data analysis
- Quantum COEC formalisms
- Hyperdimensional computing integration
- Algorithmic implementations
- Advanced cross-framework connections

Appendix B: Quick Guide / Reference

Constraint-Oriented Emergent Computation (COEC) — Quick Reference Guide

A concise map of the key ideas, formal tools, and application tracks that run through the thesis.

1. Core Premise

- **Computation = trajectory through constrained state space:** A physical/biological substrate S evolves under a constraint set C inside an energy-information landscape E ; the evolution operator Φ produces an output (residual) R while information structure I and precision weights P modulate the process.
- **Distributed agency:** Each constraint behaves like a micro-agent; apparent purpose emerges from their vector sum—no central controller needed.

2. Formal Ontology (the 7-tuple)

(S, C, E, Φ, R, I, P) with:

Symbol	Role	Typical instantiation
S	Substrate	protein chain, neural tissue, microbe colony
C	Constraints	bonds, graph topology, regulatory rules
E	Energy/info landscape	free-energy surface, metabolic cost field
Φ	Evolution law	ODE/PDE flow, Monte-Carlo sampler
R	Residual/output	fold geometry, oscillation, decision
I	Information structure	internal model, hierarchy, codebook

P Precision weights reliability of each constraint

3. Nine Canonical COEC Classes

Class	Essence	Bio exemplar	Eng. analogue
SS (Static-Structural)	ends in stable form	protein fold	self-assembly
DB (Dynamic-Behavioral)	outputs rhythm/pattern	circadian loop	oscillator circuits
DM (Distributed-Multiplicative)	non-local swarm compute	microbiome exchange	swarm robots
AP (Adaptive-Plastic)	self-edits constraints	synaptic plasticity	self-tuning nets
PP (Predictive-Probabilistic)	minimizes prediction error	visual cortex	anomaly filters
GCT	graph-metric guided	gene-reg nets	fairness routing
TDA	topology guided	morphogenesis	shape analytics
Cat	catalytic memory reuse	phase-separable enzymes	space-efficient algos

Sheaf

local-to-global
coherence

tissue morphogen

distributed
consistency

4. Computability Spectrum

- **Sub-Turing** → **Weak** → **Strong** → **Hyper-Turing** COEC, mirroring automata → PDA → TM → beyond-TM (parallel / quantum constraint satisfaction). Strong-Turing class proven *universal* via explicit UTM simulation.
- **Busy-Beaver bound**: exhaustive state search in COEC hits BB(n) limits; practical systems rely on heuristics or oracle-style constraints.

5. Mathematical Backbone

- **Measure-theoretic state spaces** and **constraint operators** with fixed-point sets.
- **Manifold & constrained flow** for continuous dynamics.
- **Lyapunov function** V_C guarantees convergence when $dV/dt < 0$.
- **PDE form** for spatial systems (e.g., morphogen gradients).

6. Information & Thermodynamics

- **Entropy lens**: $\Delta I = H(S) - H(S|C)$ quantifies computational work.
- **Entropy-minimization dynamics**: stable systems keep $dS/dt < 0$ and $d^2S/dt^2 \approx 0$; spikes trigger adaptive re-wiring.
- **Variational free-energy** formalism for PP-COEC; gradient descent on prediction error.

7. Dynamic Topology & ESSON

Entropy-Sensitive Self-Organizing Networks adjust graph edges via gradients of conditional entropy, PAC generalization, and sparsity—balancing expressiveness and parsimony.

8. Representation & Memory

Hypervector semantic embedding: high-dimensional binding \otimes and bundling \oplus enable fast associative recall; implemented with a Hypervector-Hash lattice for $O(\log n)$ lookup.

9. Category-theoretic Cohesion

- **COEC category**: objects = state spaces, morphisms = constraint-preserving maps.
- **Sheaf theory** stitches local constraints; non-trivial Čech H^1 flags global inconsistency.
- **Functorial dynamics** formalize time evolution.

10. Catalytic Space Complexity

Cat-COEC exploits *clean* + *catalytic* memory: same problem solvable with quadratically less dedicated space ($CSPACE \subseteq SPACE^2$).

11. Experimental & Design Playbook

- Protein folding precision-swap assay (SS-COEC test).
- Bacterial chemotaxis prediction-violation microfluidics (PP-COEC).
- Synthetic three-species metabolite ring (DM-COEC).

12. Expanded Themes & Practical Hooks

- **Constraint genesis** (environmental induction, composition, evolutionary fixation) and conflict resolution hierarchies.
- **Energy-information landscape sculpting** for design.
- Visual aids, executive summaries, ethics section, and an acronym glossary to boost accessibility for mixed audiences.

Summary:

Constraint-Oriented Emergent Computation (COEC) views every biological or physical system as a computer whose "program" is the mesh of constraints that shape its walk through state space: a substrate S evolves within a seven-element grammar (S, C, E, Φ, R, I, P) , where energetic topography, information flow and constraint precision jointly steer it toward a residual output R . From this vantage, protein folds, circadian clocks, swarms and adaptive brains are simply different COEC classes—static-structural, dynamic-behavioural, distributed-multiplicative, adaptive-plastic, predictive-probabilistic, graph-topological, topological-data, catalytic and sheaf-coherent—distinguished by the kind of residual they produce and the way constraints interact. The framework then maps those classes onto a computability ladder that runs from sub-Turing pattern recognizers up to hyper-Turing or oracle-aided systems, with Busy-Beaver bounds marking where exhaustive search gives way to heuristic or quantum shortcuts. Information-theoretic metrics (ΔI , entropy flow) and Lyapunov-style energy functions prove that

constraint satisfaction is both thermodynamically and mathematically grounded, while design heuristics—constraint engineering, energy-landscape sculpting, multi-scale composition and proof-of-constraint auditing—translate the theory into lab protocols and hardware blueprints. In one sweep, COEC thus unifies distributed agency, physics, information theory and category-theoretic coherence into a single language for analyzing—and deliberately building—computers made of molecules, tissues, robots or silicon.

Appendix C: Prioritized Open Questions and Solution Pathways

A systematically organized list of key research questions and potential approaches to addressing them, including:

Foundational Questions:

- Emergent constraint formation mechanisms
- Formal limits of distributed computation

Theoretical Extensions:

- Quantum integration pathways
- Relationships between COEC classes

Practical Applications:

- Implementation methods for complex constraints
- Validation metrics for experimental settings

Appendix D: Expanded Key Concepts

D.1 Constraint Genesis and Evolution

Constraints in COEC systems arise through diverse mechanisms and evolve over time. This section explores the fundamental processes that govern how constraints emerge, stabilize, and adapt across different system types.

D.1.1 Origins of Constraints

Constraints can originate through several distinct pathways:

Environmental Induction: External conditions often serve as the primary catalysts for constraint formation. For example:

- Temperature gradients induce thermodynamic constraints on reaction rates
- Resource limitations create metabolic constraints in cellular systems
- Physical barriers impose spatial constraints on diffusion processes

The mathematical formalization of environmentally-induced constraints follows:

$$c_{env}(\omega) = f_{transfer}(E_{external} \rightarrow S)$$

where $f_{transfer}$ is the transfer function mapping external energy landscapes to internal system constraints.

Evolutionary Fixation: Selection pressures can convert transient constraints into persistent ones through genetic encoding:

- Successful constraint patterns become encoded in regulatory networks
- Beneficial constraint hierarchies stabilize across generations
- Constraint precision evolves to match environmental reliability

This process can be modeled as:

$$P(c_i \text{ becomes fixed}) \propto \exp\left(\frac{\Delta F_{selection}}{k_B T_{selection}}\right)$$

where $\Delta F_{selection}$ represents the fitness advantage conferred by the constraint.

Compositional Emergence: Complex constraints can self-organize from simpler components:

- Pairwise molecular interactions give rise to higher-order assembly rules
- Simple feedback loops combine to create oscillatory timing constraints
- Local communication rules generate global synchronization constraints

Mathematically, compositional constraints can be expressed as:

$$c_{composite} = \bigotimes_{i \in I} c_i \oplus \bigotimes_{j \in J} c_j \oplus \dots$$

where \bigotimes and \oplus represent constraint binding and bundling operations respectively.

Regulatory Scaffolding: Existing constraints often guide the formation of new ones:

- Early developmental constraints provide a framework for later refinement
- Hierarchical constraint systems emerge, with higher-level constraints governing lower-level ones
- "Scaffold constraints" may be temporary, disappearing once dependent constraints stabilize

This process follows:

$$c_{new}(t+1) = g(c_{scaffold}(t), S(t))$$

where g is the constraint generation function dependent on existing scaffold constraints.

D.1.2 Constraint Evolution Dynamics

Once established, constraints undergo continuous refinement through several mechanisms:

Precision Drift: The precision weighting p_i of constraints evolves based on their predictive success:

$$\frac{dp_i}{dt} = \alpha \cdot \text{effectiveness}(c_i) - \beta \cdot \text{cost}(c_i)$$

where $\text{effectiveness}(c_i)$ measures the constraint's contribution to system success and $\text{cost}(c_i)$ represents the metabolic or computational cost of maintaining the constraint.

The adaptive variance of precision follows:

$$\text{Var}(p_i) \propto \frac{1}{\text{environmental stability}}$$

In stable environments, precision tends toward high values with low variance. In fluctuating environments, precision remains moderate with higher variance.

Constraint Composition: Simple constraints combine to form more complex constraints through several operations:

- Sequential composition: $c_A \rightarrow c_B$ (constraint A followed by B)
- Parallel composition: $c_A \parallel c_B$ (constraints A and B operating simultaneously)
- Conditional composition: $c_A \rightarrow c_B$ (constraint B activated by the satisfaction of A)

Each composition operation creates a new constraint with unique properties:

$$\begin{aligned} c_{A \rightarrow B}(\omega) &= c_B(c_A(\omega)) \\ c_{A \parallel B}(\omega) &= \min(c_A(\omega), c_B(\omega)) \\ c_{A \rightarrow B}(\omega) &= \begin{cases} c_B(\omega) & \text{if } c_A(\omega) > \theta_A \\ 1 & \text{otherwise} \end{cases} \end{aligned}$$

Constraint Relaxation: Under specific conditions, constraints may be temporarily suspended:

- Stress response systems can temporarily relax homeostatic constraints
- Developmental checkpoints may be bypassed during regeneration
- Learning systems can relax prior constraints during exploration phases

The relaxation function can be modeled as:

$$c'_i(\omega) = c_i(\omega) \cdot (1 - r_i(S, E))$$

where $r_i(S, E)$ is the relaxation factor that depends on system state and environmental conditions.

Constraint Conflict Resolution: When constraints conflict, systems employ various resolution strategies:

- Hierarchical resolution: Higher-priority constraints override lower ones
- Compromise resolution: Partial satisfaction of multiple constraints
- Temporal resolution: Alternating between conflicting constraints over time

This can be formalized as a constrained optimization problem:

$$\max_{\omega \in \Omega_S} \sum_i p_i \cdot c_i(\omega) \text{ subject to } \min_i c_i(\omega) \geq \theta_{\min}$$

where θ_{\min} is the minimum acceptable satisfaction threshold.

D.1.3 Constraint Lifecycle Phases

Constraints typically follow a lifecycle with distinct phases:

1. **Exploration Phase:** Initial, low-precision constraints emerge and compete
 - Characterized by high variance in constraint satisfaction
 - Rapid turnover of constraints through birth/death processes
 - Weak coupling between constraints
2. **Consolidation Phase:** Successful constraints stabilize and increase in precision
 - Constraint networks form with defined topological properties
 - Precision weighting differentiates critical from auxiliary constraints
 - Inter-constraint dependencies strengthen
3. **Exploitation Phase:** Established constraint networks drive efficient computation
 - High-precision primary constraints with low-precision modulatory constraints
 - Optimized energy-information trade-offs
 - Minimal constraint violations under normal conditions
4. **Adaptation Phase:** Constraints adjust in response to changing conditions
 - Precision re-weighting based on environmental feedback
 - Strategic constraint relaxation in novel scenarios
 - Birth of new constraints to address emerging challenges

The transition between phases can be detected by measuring the second derivative of system entropy:

$$\frac{d^2S}{dt^2} > \theta \implies \text{Phase transition}$$

This lifecycle pattern is observed across diverse COEC implementations, from molecular self-assembly to neural learning systems, reflecting a universal principle of constraint-based computation.

D.2 Energy-Information Landscape Interaction

The interplay between energy landscapes and information constraints forms the fundamental substrate upon which COEC systems operate. This section explores the rich mathematical and conceptual foundations of these interactions.

D.2.1 Unified Energy-Information Formalism

COEC systems navigate landscapes that combine physical energy with informational constraints. The unified landscape is expressed as:

$$E_{\text{total}}(\omega) = E_{\text{physical}}(\omega) + E_{\text{information}}(\omega)$$

Where:

- $E_{\text{physical}}(\omega)$ represents the physical energy of state ω (e.g., chemical potential, kinetic energy)
- $E_{\text{information}}(\omega)$ represents the informational cost of state ω

The informational component can be further decomposed:

$$E_{\text{information}}(\omega) = k_B T \cdot \text{DKL}(p(\omega) \parallel p_{\text{ref}}(\omega)) + \sum_i p_i \cdot [1 - c_i(\omega)]$$

Where:

- $\text{DKL}(p(\omega) \parallel p_{\text{ref}}(\omega))$ is the Kullback-Leibler divergence from a reference distribution
- $\sum_i p_i \cdot [1 - c_i(\omega)]$ is the weighted constraint violation term

This formulation unifies thermodynamic and information-theoretic perspectives, showing how constraint satisfaction can be viewed as an energy minimization process.

D.2.2 Landscape Topography Classification

Energy-information landscapes can be classified based on their topographical features:

Rugged Landscapes: Characterized by many local minima separated by high barriers

- Mathematical signature: High Lyapunov exponents and low autocorrelation length
- Common in problems with many competing constraints
- Exploration requires techniques like simulated annealing or constraint relaxation
- Examples: Protein folding, complex optimization problems

Funneled Landscapes: Possess a global structure that guides states toward specific attractors

- Mathematical signature: Negative correlation between energy and distance to target state
- Enable robust computation despite local variations
- Typical of evolved biological systems with clear functional targets
- Examples: Well-designed enzyme active sites, trained neural networks

Neutral Networks: Contain connected regions of states with similar energy/constraint satisfaction

- Mathematical signature: High percolation across iso-energy surfaces
- Enable exploration without penalty, supporting evolvability
- Critical for adaptive systems that must maintain function while exploring alternatives
- Examples: RNA secondary structures, regulatory network motifs

Multiscale Landscapes: Exhibit different features at different resolution scales

- Mathematical signature: Scale-dependent fractal dimension
- Combine small-scale roughness with large-scale directionality
- Support hierarchical constraint satisfaction
- Examples: Developmental processes, hierarchical learning systems

The formal characterization of landscape type can be quantified through several metrics:

$$\text{Ruggedness}(E) = \frac{1}{N} \sum_{i=1}^N |\nabla^2 E(\omega_i)|$$

$$\text{Funneling}(E, \omega_{\text{target}}) = -\text{Corr}(E(\omega), \text{dist}(\omega, \omega_{\text{target}}))$$

$$\text{Neutrality}(E, \theta) = \frac{|\{\omega, \omega' \mid ||\omega - \omega'|| = 1 \text{ and } |E(\omega) - E(\omega')| < \theta\}|}{|\{\omega, \omega' \mid ||\omega - \omega'|| = 1\}|}$$

D.2.3 Entropic Forces and Information Channels

COEC systems harness entropic forces to drive computation. Entropic forces arise from statistical tendencies rather than direct physical interactions:

$$F_{\text{entropic}} = T \nabla S$$

where ∇S is the gradient of entropy in state space.

These forces drive the system toward states with higher multiplicity (more microstates), which often coincide with constraint satisfaction. The interplay between entropic forces and constraints creates information channels through state space.

An information channel in COEC is formalized as:

$$\mathcal{C}(S \rightarrow S') = \{\omega(t) \mid t \in [t_0, t_1], \omega(t_0) \in S, \omega(t_1) \in S'\}$$

where S and S' are regions of state space.

The capacity of such a channel is:

$$\text{Cap}(\mathcal{C}) = \max_{p(\omega)} I(\Omega_{t_0}; \Omega_{t_1})$$

where $I(\Omega_{t_0}; \Omega_{t_1})$ is the mutual information between initial and final states.

Information channels with high capacity represent reliable computational pathways, enabling precise transitions between specific regions of state space.

D.2.4 Landscape Manipulation Techniques

Energy-information landscapes can be deliberately shaped to guide computation:

Potential Warping: Modifying the underlying energy function

- Mathematical form: $E'(\omega) = E(\omega) + \Delta E(\omega)$
- Implementation: External fields, chemical modifications, engineered binding sites
- Examples: Optical tweezers, chemical chaperones, allosteric modulators

Entropic Barriers: Creating statistical bottlenecks in state space

- Mathematical form: $\Delta S_{\text{barrier}}(\omega) = -k_B \log(\Omega_{\text{accessible}}(\omega))$
- Implementation: Excluded volume effects, molecular crowding, confinement
- Examples: Molecular nanopores, cellular compartmentalization, physical confinement

Precision Modulation: Dynamically adjusting constraint importance

- Mathematical form: $p_i(t) = p_i(0) \cdot f(t, S, E)$
- Implementation: Regulatory feedback, competitive binding, attention mechanisms
- Examples: Enzymatic regulation, neural attention, immune system tuning

Constraint Coupling: Creating dependencies between constraints

- Mathematical form: $c_{\text{coupled}}(\omega) = g(c_1(\omega), c_2(\omega), \dots, c_n(\omega))$
- Implementation: Allosteric interactions, cooperative binding, regulatory cascades
- Examples: Hemoglobin oxygen binding, transcriptional regulation, neural gating

These techniques can be combined to create sophisticated computational pathways in COEC systems, enabling complex information processing with minimal energy expenditure.

D.2.5 Thermodynamic Efficiency of Constraint Satisfaction

The efficiency of COEC systems can be analyzed through the lens of thermodynamic efficiency:

$$\eta_{\text{COEC}} = \frac{\Delta I(S, C)}{\Delta E_{\text{dissipated}}}$$

where $\Delta I(S, C)$ is the information gain from constraint application and $\Delta E_{\text{dissipated}}$ is the energy dissipated during the process.

This efficiency measure quantifies how effectively a system converts energy into computational work. The theoretical maximum efficiency is bounded by:

$$\eta_{\text{max}} = \frac{k_B T \ln(2)}{E_{\text{bit}}}$$

where E_{bit} is the minimum energy required to process one bit of information.

Biological COEC systems often approach this theoretical maximum through sophisticated energy-information landscape engineering. For example:

- Molecular motors achieve efficiency near 100% for specific transitions
- Neural signaling balances energy expenditure with information transmission
- Cellular decision-making operates near the thermodynamic limit for binary decisions

The study of these efficiency principles provides insights for designing artificial COEC systems with minimal energy requirements.

D.3 Dynamic Topology and Node Lifecycle

The topology of constraint networks in COEC systems is not static but evolves dynamically as the system computes. This section explores the mathematics and mechanisms of topological evolution.

D.3.1 Topological Representations of Constraint Networks

Constraint networks can be formally represented using several mathematical structures:

Graph Representation: A weighted directed graph $G = (V, E, W)$ where:

- Vertices V represent constraints
- Edges E represent interactions between constraints
- Edge weights W represent interaction strengths

The adjacency matrix A of this graph captures the direct influences between constraints:

$$A_{ij} = \begin{cases} w_{ij} & \text{if constraints } i \text{ and } j \text{ interact} \\ 0 & \text{otherwise} \end{cases}$$

Simplicial Complex Representation: A combinatorial structure $\mathcal{K} = (V, \Sigma)$ where:

- Vertices V represent constraints
- Simplices Σ represent higher-order interactions
- A k -simplex $\sigma \in \Sigma$ represents $(k+1)$ mutually interacting constraints

This representation captures the rich higher-order relationships between constraints that cannot be expressed by pairwise interactions alone. The structure is analyzed through simplicial homology:

$$H_n(\mathcal{K}) = \text{ker}(\partial_n) / \text{im}(\partial_{n+1})$$

where ∂_n is the boundary operator that maps n -simplices to $(n-1)$ -simplices.

Sheaf Representation: A sheaf \mathcal{F} on a topological space X where:

- The base space X represents the system's state space
- Stalks \mathcal{F}_x represent local constraint values at state x
- Restriction maps $\rho_{U,V}$ ensure consistency between regions

This representation excels at modeling how local constraints must be compatible on overlapping regions of state space, capturing the global consistency requirements of distributed constraint systems.

D.3.2 Topological Metrics and Phase Transitions

The topology of constraint networks can be characterized using several key metrics:

Degree Distribution: The distribution $P(k)$ of the number of connections per constraint:

- Scale-free networks: $P(k) \sim k^{-\gamma}$ (power law)
- Random networks: $P(k) \sim e^{-\lambda k}$ (exponential)
- Small-world networks: Combination of local clustering and short path lengths

Spectral Properties: The eigenvalue spectrum of the graph Laplacian $L = D - A$:

- The second smallest eigenvalue λ_2 (algebraic connectivity) measures network robustness
- The spectral gap $\lambda_2 - \lambda_1$ indicates how well-connected the network is
- The spectral density $\rho(\lambda)$ characterizes global network properties

Persistent Homology: Multiscale topological features tracked through filtration:

- Birth and death of topological features (connected components, loops, voids)
- Persistence diagrams visualizing topological significance
- Bottleneck or Wasserstein distance measuring topological similarity

Topological phase transitions in constraint networks are often associated with computational phase transitions. These transitions occur when network structure undergoes qualitative changes, leading to different computational behaviors:

$$P(\text{Order parameter}) = \begin{cases} 0 & \text{if } \text{control parameter} < \text{critical value} \\ > 0 & \text{if } \text{control parameter} > \text{critical value} \end{cases}$$

Examples include:

- Percolation transitions when constraint networks reach connectivity thresholds
- Synchronization transitions when coupled oscillatory constraints align
- Computational phase transitions between chaotic, critical, and ordered regimes

D.3.3 Node Birth-Death Dynamics

Constraint nodes in COEC networks undergo birth and death processes governed by utility metrics:

Birth Process: New constraints emerge when they would increase overall computational efficiency:

$$P(\text{birth of } c_{\text{new}}) \propto \exp\left(\frac{\Delta U_{\text{expected}}(c_{\text{new}})}{T_{\text{effective}}}\right)$$

where $\Delta U_{\text{expected}}(c_{\text{new}})$ is the expected utility gain and $T_{\text{effective}}$ is an effective temperature controlling exploration.

The expected utility is estimated through:

$$\Delta U_{\text{expected}}(c_{\text{new}}) = \frac{\partial \Delta I(S, C \cup \{c_{\text{new}}\})}{\partial c_{\text{new}}} - \kappa_{\text{new}}$$

where κ_{new} is the cost of implementing and maintaining the new constraint.

Death Process: Existing constraints are eliminated when their utility falls below a threshold:

$$P(\text{death of } c_i) \propto \exp\left(-\frac{\Delta U_i}{T_{\text{effective}}}\right)$$

where ΔU_i is the current utility of constraint c_i .

The birth and death processes interact to create an evolutionary dynamic in the constraint topology. This dynamic implements a form of natural selection at the constraint level, optimizing the network topology for computational efficiency.

D.3.4 Adaptive Rewiring and Plasticity

Beyond node birth and death, constraint networks undergo continuous rewiring that modifies their connection patterns:

Hebbian-Like Rewiring: Connection strengths increase between co-active constraints:

$$\frac{dw_{ij}}{dt} = \eta \cdot c_i(\omega) \cdot c_j(\omega) - \delta \cdot w_{ij}$$

where η is the learning rate and δ is a decay term.

Homeostatic Plasticity: Total incoming influence is regulated to maintain stability:

$$\frac{dw_{ij}}{dt} = \eta \cdot (w_{\text{target}} - \sum_k w_{ik}) \cdot w_{ij}$$

where w_{target} is the target total input weight.

Topological Reinforcement: Connections that participate in successful computational pathways are strengthened:

$$\frac{dw_{ij}}{dt} = \eta \cdot \frac{\partial R}{\partial w_{ij}} - \delta \cdot w_{ij}$$

where $\frac{\partial R}{\partial w_{ij}}$ is the sensitivity of the residual function to connection weight.

Constraint Diffusion: Constraints can spread influence to neighboring nodes:

$$c_i^{\text{new}}(\omega) = \alpha \cdot c_i(\omega) + (1-\alpha) \cdot \sum_{j \in N(i)} \frac{w_{ij}}{\sum_{k \in N(i)} w_{ki}} \cdot c_j(\omega)$$

where $N(i)$ is the set of nodes neighboring constraint i .

These adaptive mechanisms allow COEC networks to continuously refine their topology in response to computational demands, creating efficiently structured constraint systems tailored to specific tasks.

D.3.5 Topological Modularity and Hierarchy

Constraint networks often develop modular and hierarchical structures that enhance computational efficiency:

Modularity: Grouping of constraints into densely connected clusters with sparse inter-cluster connections:

$$Q = \frac{1}{2m} \sum_{ij} \left[A_{ij} - \frac{k_i k_j}{2m} \right] \delta(c_i, c_j)$$

where Q is the modularity score, k_i is the degree of node i , m is the total number of edges, and $\delta(c_i, c_j)$ equals 1 if nodes i and j are in the same community and 0 otherwise.

High modularity enables:

- Parallel processing of distinct computational subtasks
- Containment of errors within modules
- Reuse of modular constraint patterns

Hierarchy: Organization of constraints into multiple levels of influence:

$$H = 1 - \frac{\sum_i \sum_{j \in N(i)} |l_i - l_j|}{k_i \cdot L}$$

where H is the hierarchy score, l_i is the hierarchical level of node i , k_i is its degree, L is the total number of levels, and n is the number of nodes.

Hierarchical structures provide:

- Multiscale constraint processing
- Top-down control of computational priorities
- Efficient information aggregation and distribution

The evolution toward modular and hierarchical topologies is a hallmark of advanced COEC systems, enabling them to manage complexity through structured constraint organization.

D.4 Predictive Processing Mechanisms

Predictive processing represents a fundamental computational strategy in many COEC systems, particularly those that must interact with dynamic environments. This section explores the mathematical foundations and implementations of predictive processing within the COEC framework.

D.4.1 Generative Models as Constraint Systems

In PP-COEC systems, internal generative models function as constraint sets that predict future states:

$$M: S(t) \mapsto \hat{S}(t+\Delta t)$$

where $\hat{S}(t+\Delta t)$ is the predicted future state.

These models can be formalized as constraint sets:

$$C_M = \{c_1, c_2, \dots, c_n\}$$

where each constraint c_i captures a specific aspect of the prediction.

The precision-weighted satisfaction of these constraints determines the model's prediction:

$$\hat{S}(t+\Delta t) = \arg\max_{S' \in \Omega_S} \sum_{i=1}^n p_i \cdot c_i(S')$$

This formulation allows generative models to be seamlessly integrated into the broader COEC framework, where prediction becomes a special case of constraint satisfaction.

D.4.2 Prediction Error Minimization

The core computational principle in PP-COEC systems is the minimization of prediction error:

$$\delta(t) = S_{\text{actual}}(t) - S_{\text{predicted}}(t)$$

This error drives system evolution through several mechanisms:

State Update: The system state is adjusted to reduce prediction error:

$$\Delta S(t) = -\alpha \cdot \delta(t)$$

where α is the learning rate for state updates.

Model Update: The generative model itself is revised based on persistent errors:

$$\Delta C_M = -\beta \cdot \nabla_{C_M} \mathbb{E}[\delta^2]$$

where β is the learning rate for model updates and $\nabla_{C_M} \mathbb{E}[\delta^2]$ is the gradient of expected squared prediction error with respect to the model constraints.

Precision Weighting: The reliability of different prediction components is continuously adjusted:

$$\Delta p_i = \gamma \cdot \left(\frac{1}{|\delta_i|^2 + \epsilon} - p_i \right)$$

where γ is the learning rate for precision updates, δ_i is the prediction error component associated with constraint c_i , and ϵ is a small constant preventing division by zero.

This multilevel error minimization process implements a sophisticated form of Bayesian inference, where the system continually refines both its beliefs (state) and its belief-formation mechanisms (constraints) based on evidence.

D.4.3 Hierarchical Predictive Processing

PP-COEC systems typically organize predictive constraints into hierarchical structures:

$$C_M = \{C_1, C_2, \dots, C_L\}$$

where C_l represents the constraint set at level l of the hierarchy.

Each level makes predictions about the level below it:

$$\hat{S}_l(t) = f_l(S_{l+1}(t))$$

where f_l is the prediction function for level l .

Prediction errors propagate upward through the hierarchy:

$$\delta_l(t) = S_l(t) - \hat{S}_l(t) \quad S_{l+1}(t+1) = S_{l+1}(t) - \alpha_{l+1} \cdot \delta_l(t)$$

The precision assigned to prediction errors varies across hierarchical levels:

$$p_l = g(l, \text{context})$$

where g is a function that determines precision based on hierarchical level and context.

This hierarchical organization enables:

- Simultaneous prediction at multiple temporal and spatial scales
- Abstract, high-level constraints influencing concrete, low-level predictions
- Context-sensitive processing through top-down modulation of precision

D.4.4 Active Inference and Control

PP-COEC systems extend beyond passive prediction to active control through the principle of active inference:

$$\text{Action} = \arg\min_a \mathbb{E}[\text{Free Energy}(S(t+1) | a)]$$

where $\text{Free Energy}(S) = \mathbb{E}_q[\log q(S) - \log p(S, O)]$ is the variational free energy, $q(S)$ is the approximate posterior distribution over states, $p(S, O)$ is the generative model linking states and observations, and a represents potential actions.

This can be reformulated in terms of constraint satisfaction:

$$\text{Action} = \arg\max_a \mathbb{E}[\sum_i p_i \cdot c_i(S(t+1) | a)]$$

Actions are thus selected to maximize expected future constraint satisfaction, implementing a form of control as constrained optimization.

This approach unifies perception and action within a single framework:

- Perception: Updating internal state to satisfy constraints
- Action: Modifying external state to satisfy the same constraints

The implementation of active inference involves:

1. Maintaining a set of policy constraints that predict outcomes of different actions
2. Evaluating expected free energy (constraint satisfaction) under different policies
3. Selecting actions that minimize expected free energy
4. Updating policy constraints based on observed outcomes

D.4.5 Attention as Precision Allocation

Attention mechanisms in PP-COEC systems can be formalized as dynamic precision allocation:

$$p_i(t) = p_{\text{base},i} \cdot \text{Attention}(i, t)$$

where $p_{\text{base},i}$ is the baseline precision for constraint c_i and $\text{Attention}(i, t)$ is the attentional modulation at time t .

Attention is allocated based on several factors:

Uncertainty Reduction: Directing attention to reduce uncertainty in critical predictions:

$$\text{Attention}(i, t) \propto \frac{\partial \mathbb{H}[S]}{\partial p_i}$$

where $\mathbb{H}[S]$ is the entropy of the system state.

Goal Relevance: Enhancing precision for constraints relevant to current goals:

$$\text{Attention}(i, t) \propto \text{Relevance}(c_i, \text{Goals}(t))$$

Surprise Detection: Increasing attention to unexpected prediction errors:

$$\text{Attention}(i, t) \propto \exp\left(\frac{|\delta_i(t)|^2}{2\sigma^2}\right)$$

where σ^2 is the expected variance of prediction errors.

By dynamically modulating precision, attention mechanisms effectively create a flexible "spotlight" that highlights the most informationally valuable constraints at each moment, optimizing computational resources and enhancing performance in complex, changing environments.

D.5 Semantic Representation and Drift Detection

COEC systems require sophisticated mechanisms for representing, storing, and monitoring the semantic content of their computations. This section explores the mathematical foundations of these semantic mechanisms.

D.5.1 Hyperdimensional Semantic Spaces

Semantic information in COEC systems is represented in high-dimensional vector spaces:

$$\mathcal{H} \subset \mathbb{R}^d \text{ where } d \gg 1000$$

These high-dimensional spaces possess several advantageous properties:

Quasi-Orthogonality: Random vectors are approximately orthogonal:

$\|\mathbb{E}[\cos(v_1, v_2)]\| \approx 0$ for random $v_1, v_2 \in \mathcal{H}$

Superposition: Multiple vectors can be combined through addition:

$$v_{\text{bundle}} = v_1 + v_2 + \dots + v_n$$

Binding: Vectors can be bound together through various operations:

$$v_{\text{bind}} = v_1 \otimes v_2$$

where \otimes may represent element-wise multiplication, circular convolution, or other binding operations.

These properties enable rich compositional semantics within a fixed-dimensional representation, supporting complex symbolic operations in a distributed format.

D.5.2 Semantic Encoding of Constraints

Constraints are encoded as hypervectors through various mapping functions:

Direct Encoding: Mapping constraint parameters directly to vector components:

$$h_{\text{direct}}(c_i) = \Phi_{\text{direct}}(\text{parameters}(c_i))$$

Functional Encoding: Representing constraints by their effect on state space:

$$h_{\text{functional}}(c_i) = \Phi_{\text{functional}}(\{c_i(\omega) \mid \omega \in \Omega_{\text{sample}}\})$$

Compositional Encoding: Building constraint representations from component parts:

$$h_{\text{compositional}}(c_i) = \bigoplus_{j=1}^m (h_{\text{role},j} \otimes h_{\text{filler},j})$$

where \oplus represents bundling, \otimes represents binding, and $h_{\text{role},j}$ and $h_{\text{filler},j}$ represent roles and values in the constraint's structure.

The choice of encoding method depends on the specific COEC system and application, with many systems employing multiple encoding strategies for different aspects of their constraint sets.

D.5.3 Semantic Memory Systems

COEC systems store and retrieve semantic information through specialized memory systems:

Hypervector-Hash Lattice: A multidimensional hashing structure for efficient retrieval:

$$\text{HashLattice}(h) = \lfloor \text{LSH}(h) \rfloor$$

where LSH is a locality-sensitive hashing function.

This structure enables $O(\log n)$ retrieval time for nearest-neighbor queries, supporting efficient associative memory even for very large semantic spaces.

Semantic Tensor Networks: Storing relationship information in tensor structures:

$$\mathcal{T}(i,j,k,\dots) = \text{strength of relationship } (i,j,k,\dots)$$

These tensors can be factorized for efficiency and generalization:

$$\mathcal{T} \approx \sum_{r=1}^R a^{(1)}_r \otimes a^{(2)}_r \otimes \dots \otimes a^{(n)}_r$$

where $a^{(i)}_r$ are factor vectors and R is the tensor rank.

Episodic Traces: Temporal sequences of semantic states:

$$E = (h_1, h_2, \dots, h_T)$$

These traces can be compressed into single episodic vectors:

$$h_E = \sum_{t=1}^T \text{position}(t) \otimes h_t$$

where $\text{position}(t)$ is a position encoding vector.

The integration of these memory systems enables COEC systems to build, maintain, and utilize complex semantic knowledge structures that guide their computational processes.

D.5.4 Semantic Drift Detection

Monitoring semantic drift is crucial for maintaining computational integrity in COEC systems:

Hypercore References: Stable reference points in semantic space:

$$H_{\text{core}} = \{h_{\text{core},1}, h_{\text{core},2}, \dots, h_{\text{core},k}\}$$

These core vectors represent fundamental semantic anchors against which current states can be compared.

Drift Metrics: Quantitative measures of semantic change:

$$\text{Drift}(S_t) = 1 - \cos(h(S_t), h_{\text{core}})$$

where $h(S_t)$ is the semantic embedding of the current state and h_{core} is the most relevant semantic hypercore.

Drift Thresholds: Critical values triggering corrective actions:

$$\text{if } \text{Drift}(S_t) > \theta_{\text{drift}} \text{ then activate drift correction}$$

where θ_{drift} is the maximum acceptable drift before intervention.

Drift Correction Mechanisms: Processes to realign semantic content:

$$h_{\text{corrected}}(S_t) = \alpha \cdot h(S_t) + (1-\alpha) \cdot h_{\text{core}}$$

where α controls the balance between current state and core reference.

These drift detection mechanisms provide COEC systems with semantic stability while allowing appropriate adaptation, preventing both excessive rigidity and uncontrolled drift.

D.5.5 Phase Mapping and Semantic Classification

COEC systems classify their states into semantic phase profiles that capture high-level behavioral and functional characteristics:

Phase Vector: A multidimensional characterization of system state:

$$\Psi(S) = (p_1, p_2, \dots, p_k)$$

where each dimension p_i represents a specific phase property.

Phase Mapping Functions: Transforms from state to phase profile:

$$p_i = f_i(S, C, E, I, P)$$

These functions extract semantic essence from complex system states, mapping high-dimensional configurations to interpretable characteristics.

Phase Transitions: Qualitative shifts in behavior:

$$\text{if } \|\Psi(S_{t+1}) - \Psi(S_t)\| > \theta_{\text{phase}} \text{ then phase transition has occurred}$$

where θ_{phase} is the threshold for recognizing a phase transition.

Phase Libraries: Collections of prototypical phase profiles:

$$\mathcal{L} = \{\Psi_1, \Psi_2, \dots, \Psi_m\}$$

These libraries enable rapid classification of current system behavior by comparing observed phase profiles to known patterns.

Phase mapping provides a crucial semantic interface between COEC systems and human understanding, translating the complex internal dynamics of constraint satisfaction into interpretable behavioral and functional categories.

Appendix E: Acronym Glossary

Acronym	Full Term	Brief Description
COEC	Constraint-Oriented Emergent Computation	Framework describing computation as the trajectory of systems through constrained state spaces
SS-COEC	Static-Structural COEC	Systems producing stable structural outputs (e.g., protein folding)
DB-COEC	Dynamic-Behavioral COEC	Systems producing stable temporal patterns (e.g., circadian rhythms)
DM-COEC	Distributed-Multiplicative COEC	Systems where computation emerges from multiple interacting subsystems
AP-COEC	Adaptive-Plastic COEC	Systems that modify their own constraints over time
PP-COEC	Predictive-Probabilistic COEC	Systems using internal models to anticipate future states
GCT-COEC	Graph-Constrained Topology COEC	Systems using graph metrics to guide computation
TDA-COEC	Topological Data Analysis COEC	Systems using topological features to guide computation
Cat-COEC	Catalytic COEC	Systems using transient external memory for computation

PAC	Probably Approximately Correct	Learning theory framework for generalization guarantees
ESSON	Entropy-Sensitive Self-Organizing Network	Network topology evolving according to entropy gradients
FDSC	Fluid Distributed Constraint Sharing / Folding-Driven Structural Computation	Privacy-preserving distributed computation / Topological data analysis via folding
BB	Busy Beaver	Function defining computational limits for exhaustive search
VSA	Vector Symbolic Architecture	Framework for operating on distributed representations using high-dimensional vectors
IIT	Integrated Information Theory	Theory linking consciousness to information integration in complex systems
FEP	Free Energy Principle	Biological systems act to minimize variational free energy
FBA	Flux Balance Analysis	Constraint-based modeling of metabolic networks
LSH	Locality-Sensitive Hashing	Technique for dimensionality reduction preserving similarity relationships

TDA	Topological Data Analysis	Mathematical approach to analyzing data using topological features
NAS	Neural Architecture Search	Automated discovery of optimal neural network architectures
ReRAM	Resistive Random-Access Memory	Non-volatile memory technology used in neuromorphic hardware
CIM	Compute-in-Memory	Approach that performs computations where data is stored
POVM	Positive Operator-Valued Measure	Quantum measurement formalism used in Quantum-COEC
CDL	Constraint Definition Language	Formal language for defining constraints across domains

16. Road-Test Plan (24 months)

Milestone	Hypothesis	COEC metric	Method
64-dim protein eigen-solver	fold \rightarrow eigenvector	$\ R_{\text{exp}} - R_{\text{eigen}}\ _2 < 0.1$	CD, LC-MS
Virus-protein-neuron tri-culture	residual propagation cross-tiers	mutual-info ≥ 0.5 bit	Ca-imaging + qPCR

128×128 memristor
PDE chip

poly vs exp scaling

error <5 % at $N=10^4$

on-chip ADC
logs

Proof-of-Constraint field
test

ledger immutability

tamper prob $<10^{-12}$

Mina-style
zk-snarks

Each experiment closes the loop from constraint theory → substrate → residual, reinforcing COEC falsifiability.